

2023

DS 7333 | Quantifying the World

CASE STUDY 2 | DIABETES & HOSPITAL READMISSION

JOEY HERNANDEZ
DANIEL CHANG

Introduction

Background

Diabetes, a chronic metabolic disorder characterized by elevated blood sugar levels, is a significant global health concern. The management of diabetes often involves hospitalization, especially for individuals experiencing acute complications or undergoing significant treatment adjustments. One critical aspect of diabetes care is the prevention of readmission into a hospital within a short period after the initial discharge. Readmissions can indicate various issues, such as inadequate treatment during the initial hospitalization, complications arising post-discharge, or a lack of effective outpatient care.

Objective and Scope

The primary objective of this report is to develop a predictive model to anticipate whether a patient with diabetes will be readmitted to a hospital within 30 days following their initial discharge. Achieving this objective can offer several advantages, including improved patient care, reduced healthcare costs, and enhanced resource allocation within healthcare facilities.

To accomplish our goal, we will employ logistic regression, a widely used statistical technique for binary classification. Logistic regression is well-suited for this task as it allows us to model the probability of readmission based on a set of relevant predictor variables. By analyzing these predictors, we can gain insights into the factors contributing to readmission risk among diabetic patients. Additionally, we will utilize imputation techniques to handle missing data, ensuring that our analysis is robust and representative of the patient population.

Data Source

This case study will utilize two datasets, “diabetic_data” and “IDs_mapping”. The data was provided to us in the Case 2 Study Module and is in the form of two separate csv files. When combined the data contains 101766 observations and 53 features.

Data Inspection

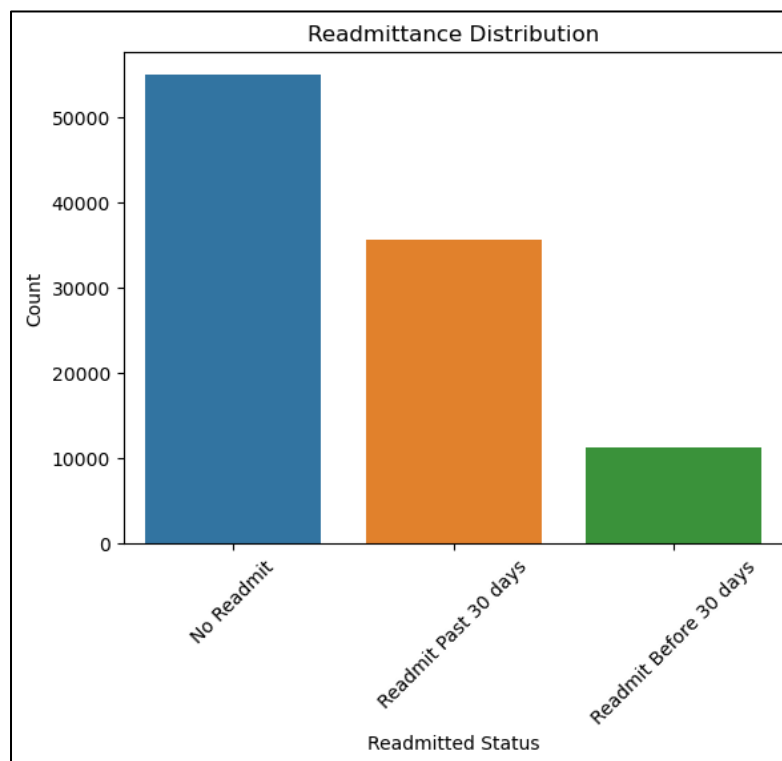
Before creating any models or analysis with the data our first step was to inspect our data to better understand data types (such as “int”, “cat”, “object”, etc.), distributions of values, identification of missing values, duplicated data, and outliers. This step is vital in understanding how we should approach any types of transformations or adjustments to the modeling and analysis process of our data.

Target Variable Inspection

To gain insight into the distribution of our classification target variable, we performed a thorough examination of the data. The target variable in this analysis “readmitted” represents the status of no readmission, readmission within 30 days or greater than 30 days of the initial hospitalization for patients. It is a critical factor in our predictive model, as it helps determine whether a patient falls into one of two categories: "no readmittance" or "less than 30 days readmittance." Which is the objective classification problem for our case study.

Initially, we visualized the distribution of the target variable in its original form. This initial examination revealed a substantial class imbalance between the three categories. Most instances were labeled as "no readmittance," while a considerably smaller portion represented "less than 30 days readmittance", and “greater than 30 days readmittance”. This significant imbalance posed a challenge for our predictive modeling, as models may tend to be biased toward the majority class, potentially leading to suboptimal performance.

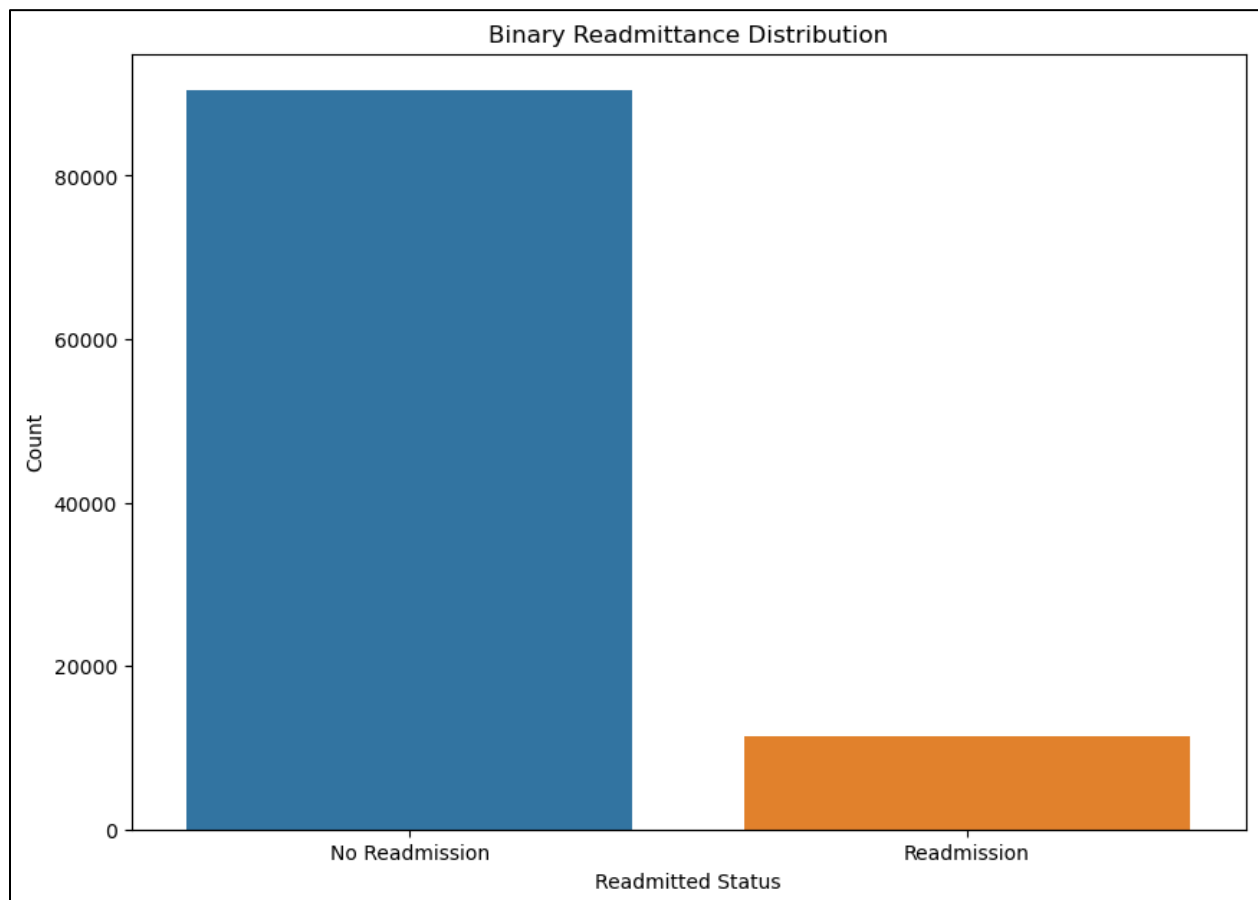
Figure 1: Count Plot Illustrating the Distribution of the Target Variable across the classes “No” “>30,” “<30”



Description: The count plot indicates that out of all the readmitted status. Most were not readmitted.

The first amendment to our data and thus the scope of our project began with removal of the 'greater than 30 days readmittance'. The objective of our data as provided to us by the client specifically mentioned the desire to predict patients that would be readmitted to the hospital within 30 days, thus this change in the data allows us to remove information from our data that may otherwise bias our prediction of what type of readmittance is occurring.

Figure 2: Count Plot Illustrating the new Distribution of the Target Variable across the binary values of “No Readmittance” and “Readmittance”, where Readmittance is only occurrences where the patient was readmitted <30 days.



Description: The count plot indicates that out of all the readmitted statuses (after converting them to binary). Our “No Readmission” count is far higher than the “readmission” count.

Patient Data Privacy And Information Inclusion/Exclusion

Respecting patient data privacy is of paramount importance in healthcare research, and this study is no exception. We recognize that certain demographic and clinical variables, such as race, gender, and age, can be invaluable in identifying potential contributors to the prediction of patient readmission. However,

it is essential to emphasize that in our data analysis and modeling process, stringent measures have been implemented to safeguard patient privacy.

All patient identifiers, including but not limited to names, addresses, and specific identification numbers, have been rigorously excluded from our dataset. Additionally, any data attributes that could be used to directly infer individual patient identities or sensitive personal information have been carefully removed or anonymized to ensure the utmost protection of patient privacy.

The information used in our predictive model is limited to objective and de-identified data points relevant to the healthcare context. This means that our analysis is focused solely on variables that contribute to the accurate prediction of patient readmission risk, without compromising the confidentiality and privacy of individual patients. Our commitment to data privacy aligns with the highest ethical standards and legal requirements, ensuring that the insights derived from this study are both valuable and ethically sound.

Missing Data

In the process of preparing and cleaning the dataset for our analysis, we encountered various missing data points across different attributes. Our approach to handling these missing values is guided by thorough investigations into the nature and potential implications of the missing data.

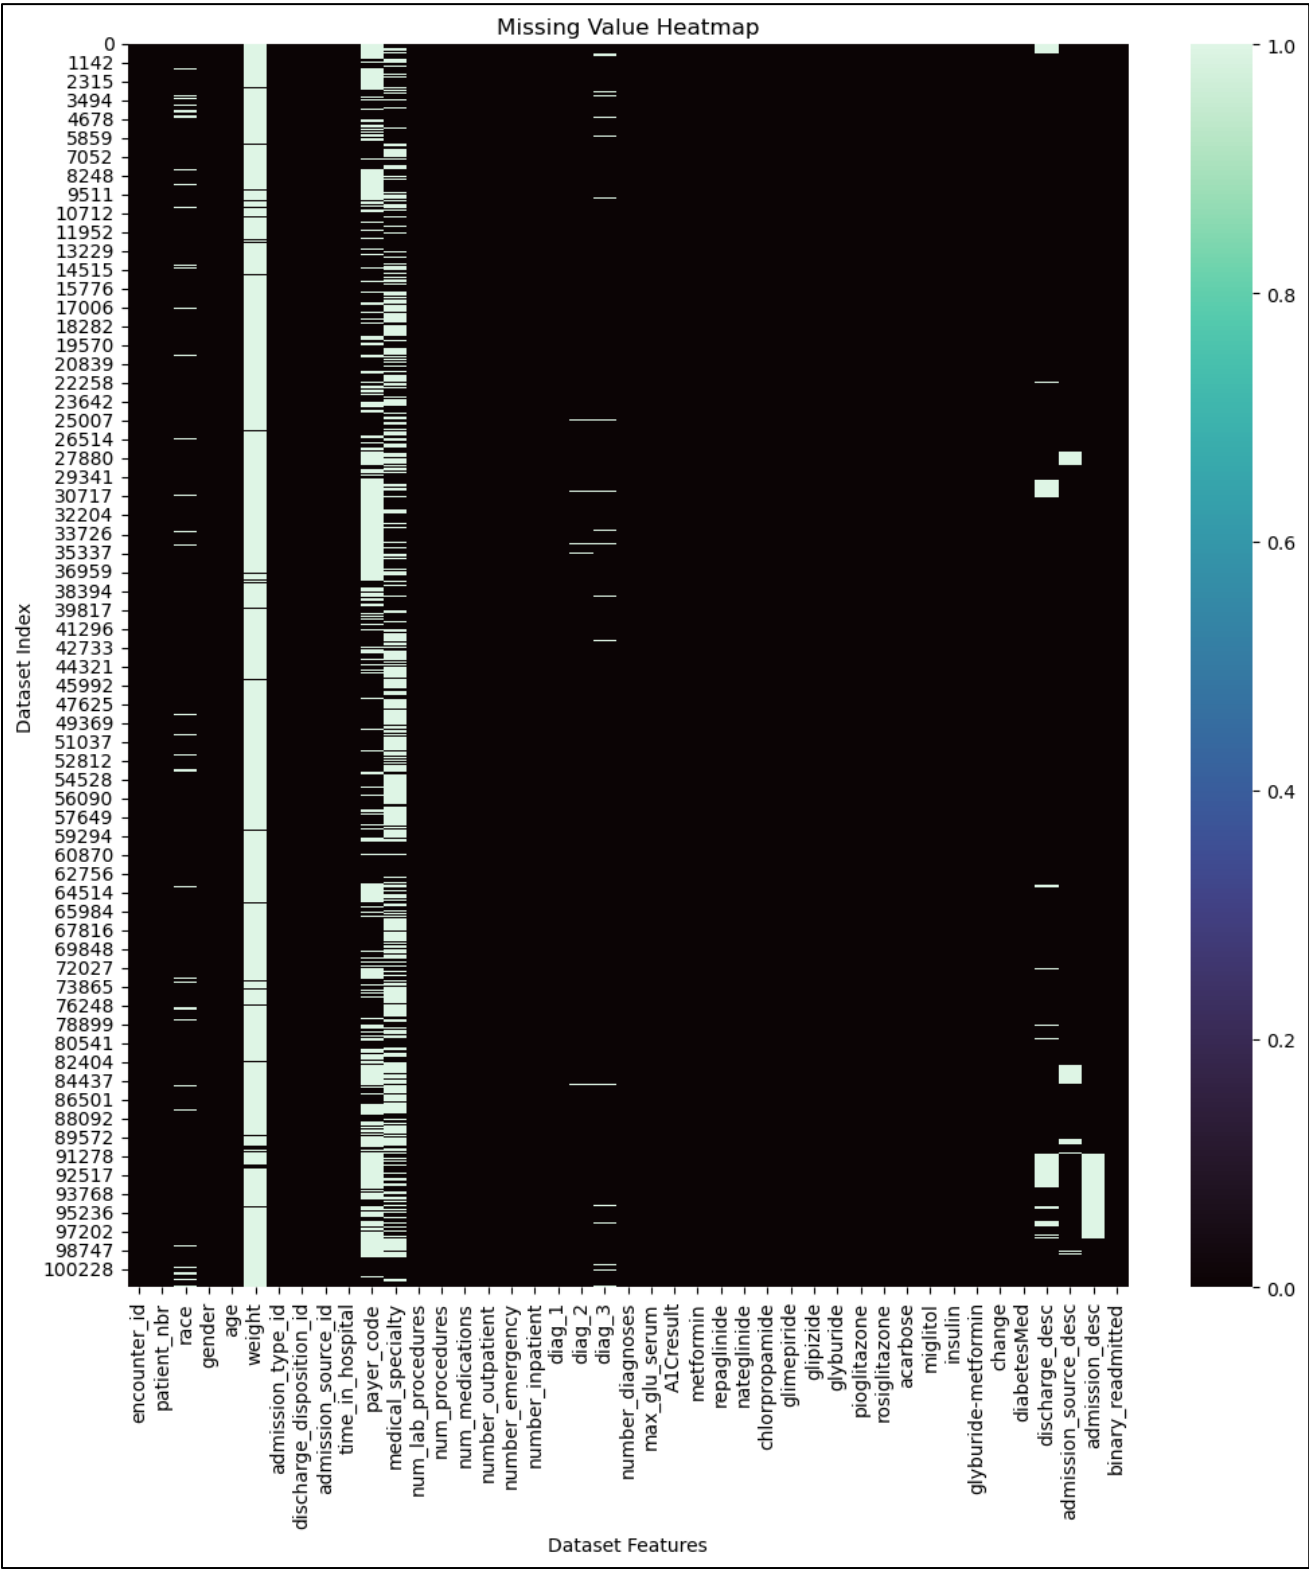
1. **Race:** In the 'race' attribute, we identified missing values denoted by "?". After further investigation, including visual analysis and cross-tabulations, we determined that these missing values were Missing Completely at Random (MCAR). As a result, we chose to proceed with imputation by replacing the missing values with the most frequent category. This approach allows us to maintain the integrity of the dataset while addressing the missing data issue.
2. **Weight:** The 'weight' attribute exhibited a significant percentage of missing values, exceeding 90%. Similar to the 'race' attribute, we investigated the nature of this missing data and confirmed that it was MCAR. However, due to the substantial extent of missing values and the limited potential usefulness of this feature, we made the decision to remove the 'weight' attribute from our dataset.
3. **Payer Code:** Missing values represented by "?" in the 'payer_code' attribute were also identified as MCAR upon thorough investigation. In this case, instead of imputation or removal, we chose to re-label the missing values and include them as a separate category within the 'payer_code' feature. This approach ensures that we retain valuable information while handling the missing data appropriately.
4. **Medical Specialty:** For the 'medical_specialty' attribute, there was no discernible relationship between the values in our dataset and the missingness of medical specialty values. To address this, we proceeded with inputting the missing values by assigning them a new category labeled as "other." This imputation strategy helps preserve the overall structure of the data while accounting for the missing information.

5. **Diagnosis Codes (diag_1, diag_2, diag_3):** While these attributes did contain missing values, it is worth noting that some of these gaps may be attributed to patients not having a specific diagnosis to report. Given the inconsequential amount of missing data in the diagnosis codes, we opted for a conservative approach by removing the null values from these attributes.

Our data preprocessing efforts regarding missing data aim to ensure that the resulting dataset is as informative and representative as possible while mitigating the potential bias introduced by the missing values. These carefully considered strategies allow us to maintain the integrity of the data and facilitate meaningful analyses for our predictive model of patient readmission.

Figure 3: A heatmap representing the distribution of missing values across the dataset. Darker areas indicate the absence of data points, while lighter regions denote complete information.

Understanding the pattern of missingness is crucial for effective data preprocessing and predictive modeling.



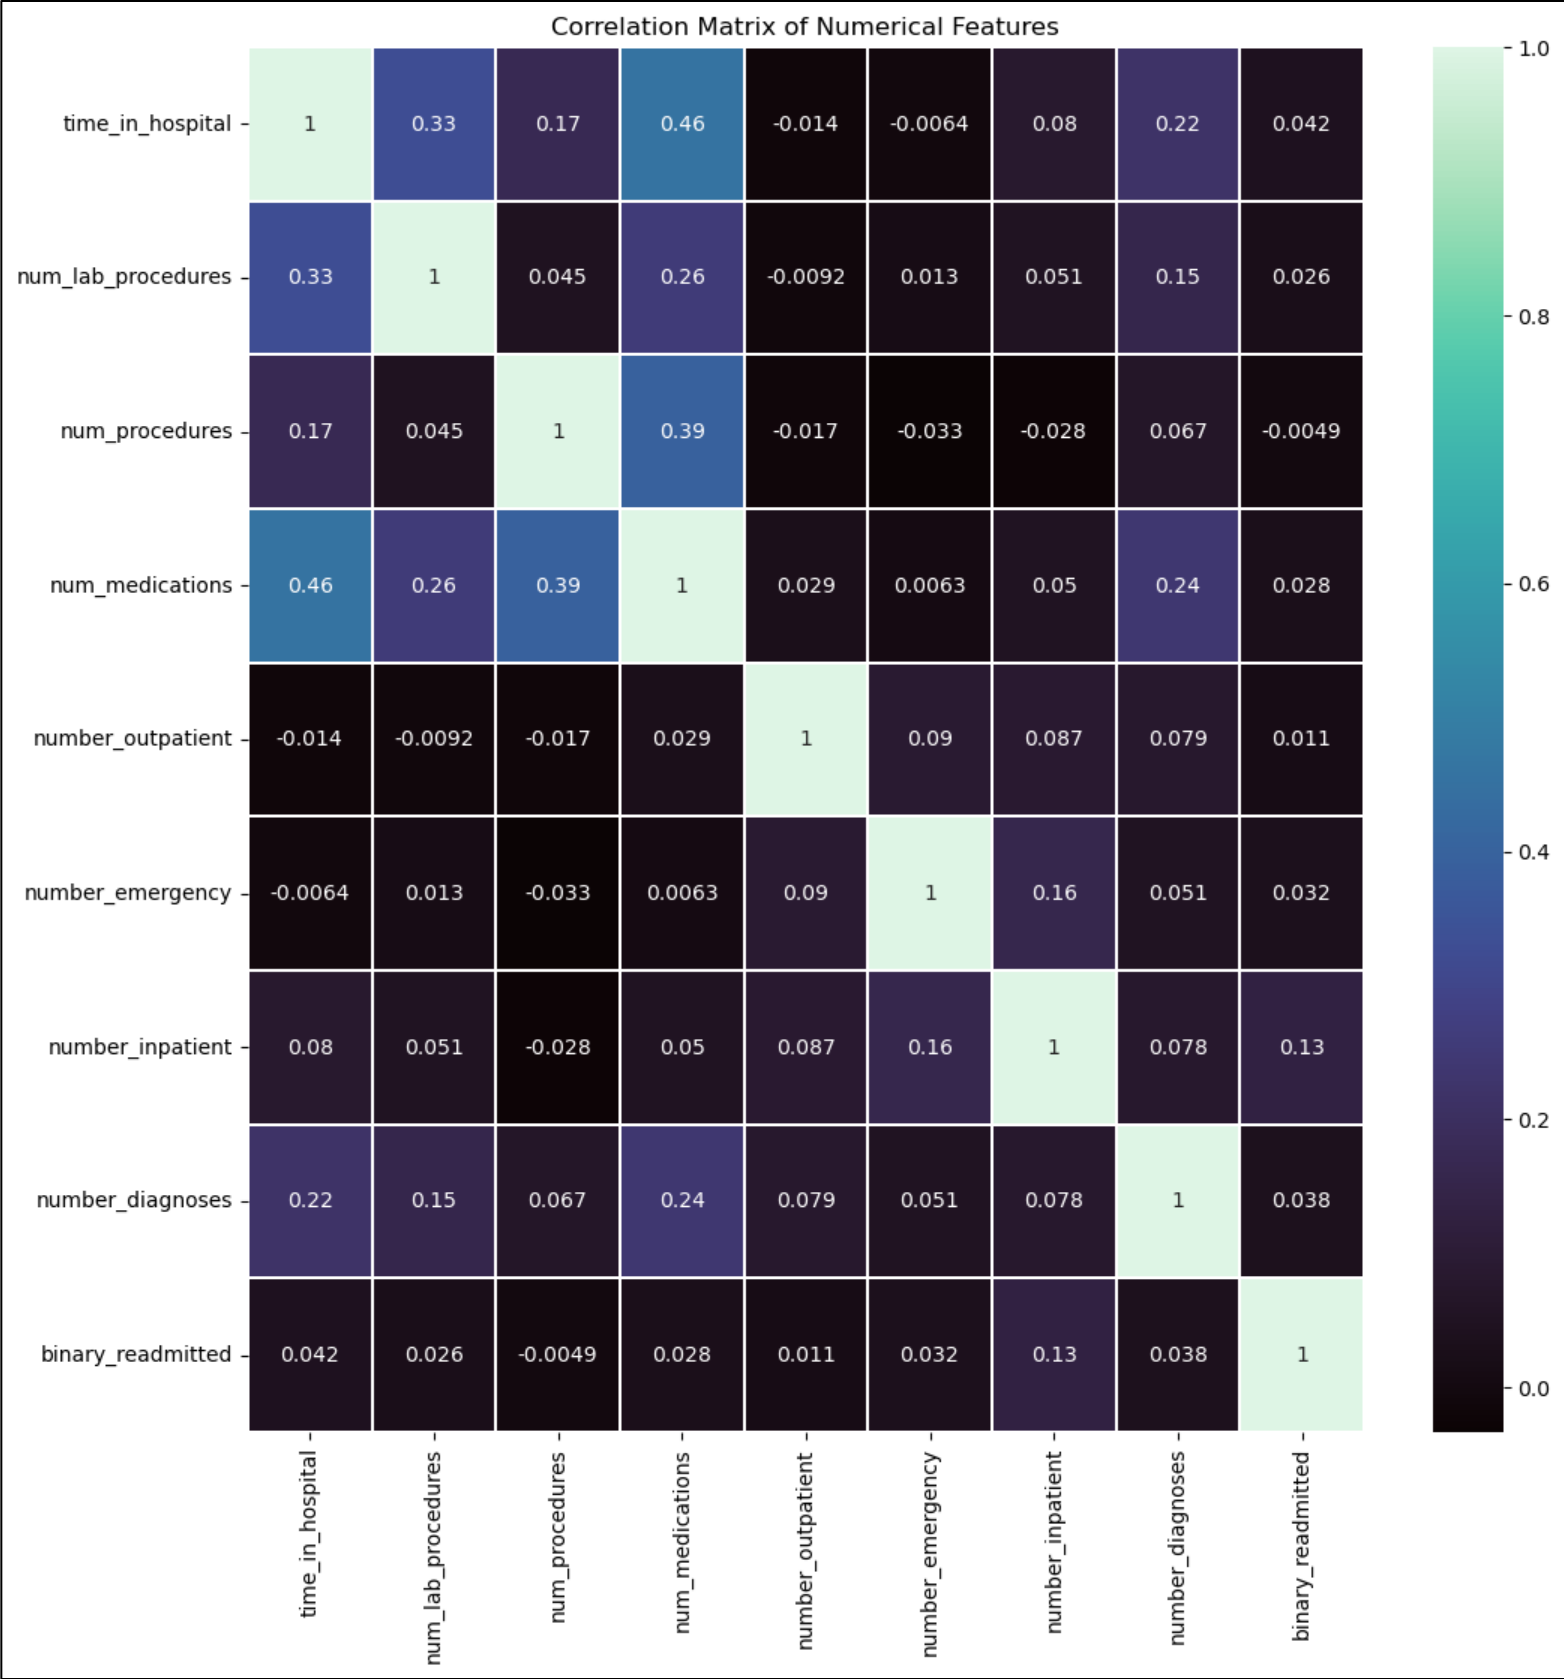
Description: Here is the heatmap of missing values. The lighter color represents a missing value, while the darker color represents a non-missing value. Right away, we can see that the weight column has the most number of missing values.

Correlation Plot (Original Target Variable)

Prioritizing the preprocessing steps for the target variable in the step prior to this was important because it allowed us to assess integrity and accuracy of subsequent analyses. By addressing the target variable's distribution and more narrow scope, the resulting correlation values can be better trusted to either accurately reflect the underlying relationships between variables or understand what limitations may arise from the less than desirable target variable distribution. Failure to preprocess and identify data discrepancies in the target variable could lead to misinterpretations, as correlations might be influenced by skewedness, outliers, or nonlinearities within the target data.

Performing a correlation heatmap provides a visually informative representation of the relationships between variables within a dataset. By illustrating the strength and direction of linear associations, the heatmap becomes an indispensable tool for uncovering patterns and dependencies that might not be immediately apparent from individual variable analyses. Each cell in the heatmap corresponds to a pair of variables, with the color gradient indicating the magnitude of correlation. This enables the rapid identification of high and low correlation values, highlighting potential areas of interest for further investigation.

Figure 4 (next page): A heatmap illustrating the correlation among integer-type features in the dataset. The color intensity reflects the strength and direction of associations between variables. Analyzing feature correlations is essential for identifying potential patterns and dependencies that can inform our predictive modeling efforts.



Modeling

Lasso (L1 Regularization) Logistic Regression Model

To build an effective predictive model for patient readmission within 30 days of initial hospitalization, we began with a Lasso Logistic Regression approach. Lasso (Least Absolute Shrinkage and Selection Operator) is a regularization technique that helps prevent overfitting by penalizing the absolute values of the regression coefficients. In our analysis, we utilized Lasso Logistic Regression as a starting point due to its ability to perform feature selection by driving some coefficient estimates to zero, thereby simplifying the model and enhancing its interpretability.

Our model-building process included several crucial steps:

- 1. Grid Search for Optimal Alpha Value:** We performed a grid search to determine the best alpha value for the Lasso Logistic Regression model. Alpha controls the strength of the L1 regularization penalty, and finding the optimal value is essential for achieving the right balance between model complexity and performance.
- 2. Train-Test Split:** To assess the model's performance effectively, we divided the dataset into training and testing subsets. The training data were used to train the model, while the testing data served as an independent dataset for evaluating its generalization performance.
- 3. Feature Scaling and Dummitizing:** Proper scaling of the features is crucial for logistic regression models. We standardized the numerical features to have zero mean and unit variance to ensure that all features contributed equally to the model. Additionally, we employed one-hot encoding (dummitizing) for categorical variables to convert them into a format suitable for the logistic regression model.
- 4. Applying Threshold to Classification:** In binary classification problems like ours, a probability threshold is applied to determine the predicted class labels. We experimented with various threshold values to optimize the trade-off between sensitivity and specificity, tailoring the model to prioritize either minimizing false positives or false negatives, depending on the clinical context.

By implementing these steps, we aimed to develop a robust Lasso Logistic Regression model capable of predicting patient readmission within 30 days. This model not only provides predictive accuracy but also offers interpretability, allowing us to identify the most influential factors contributing to the likelihood of readmission. In the subsequent sections, we will detail the results of our model evaluation and discuss its implications for improving patient care and healthcare resource allocation.

Results:

Optimal Alpha (Lambda): The grid search procedure identified the best alpha value for Lasso regularization as 0.0100. This parameter controls the strength of regularization, striking a balance between model complexity and performance.

Test Set Accuracy: The Lasso Logistic Regression model achieved an accuracy of 0.6104 on the test dataset. This metric reflects the proportion of correctly classified instances, indicating that the model correctly predicted the readmission outcomes for approximately 61.04% of the test cases.

User-Defined Threshold: In binary classification, the choice of threshold for converting predicted probabilities into class labels can significantly impact the model's performance. In this case, a user-defined threshold of 0.65 was applied to the model's predicted probabilities to determine the classification labels.

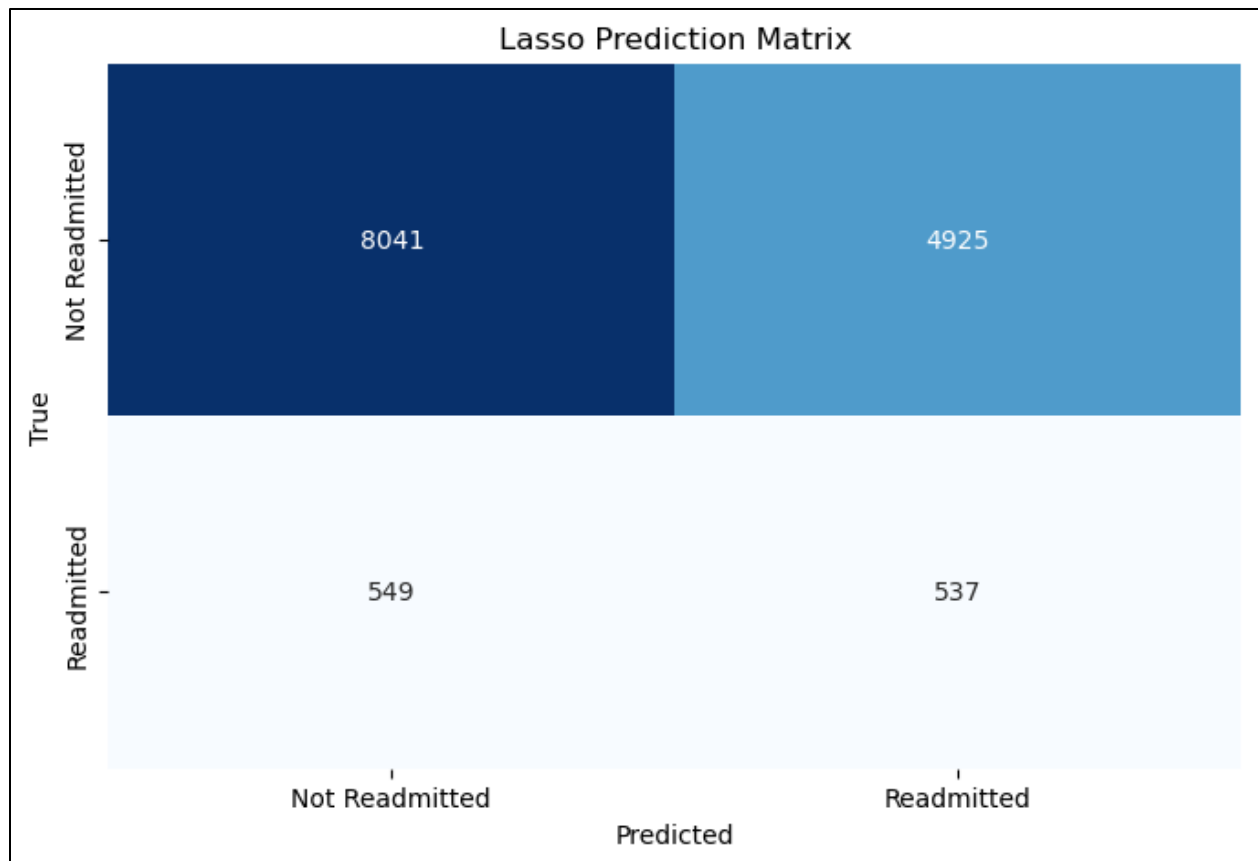
Classification Report with Custom Threshold:

- **Precision:** The precision for class "0" (no readmittance) was notably high at 0.94. This indicates that when the model predicted "no readmittance," it was accurate 94% of the time. Intuitively this makes sense given that the majority class of our data was "no readmittance" so it would be expected that the model would be bias to predicting "no readmittance".
- **Recall:** The recall for class "0" was 0.62, indicating that the model captured 62% of the actual "no readmittance" cases. For class "1" (less than 30 days readmittance), the recall was 0.49, suggesting that the model identified 49% of the actual cases.
- **F1-Score:** The F1-score, which balances precision and recall, was 0.75 for class "0" and 0.16 for class "1."
- **Support:** The support represents the number of instances in each class within the test dataset. In this case, class "0" had significantly more instances (12,966) than class "1" (1,086).

Macro and Weighted Averages: The macro-average F1-score was 0.46, reflecting the overall balance between precision and recall across both classes. The weighted-average F1-score, which accounts for class imbalances, was 0.70. These averages provide a comprehensive view of the model's performance across all classes.

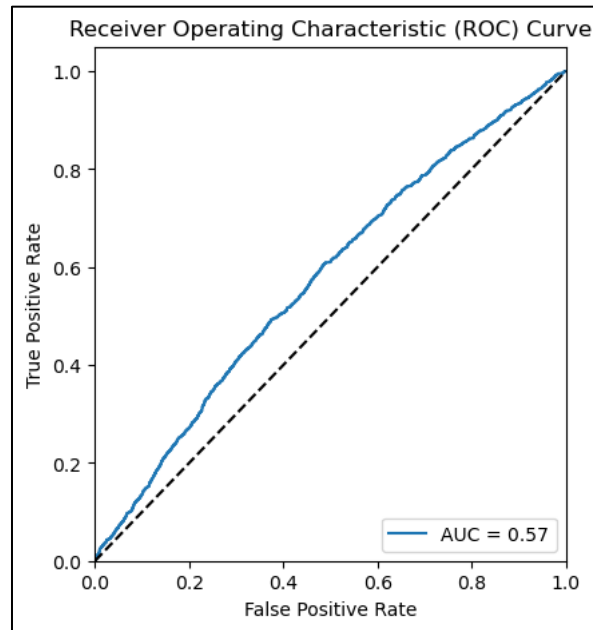
In summary, the Lasso Logistic Regression model, with an optimized alpha value and a user-defined threshold, demonstrated reasonable accuracy in predicting patient readmission within 30 days. It exhibited strong precision for the "no readmittance" class, indicating a high degree of confidence in negative predictions. However, there is room for improvement in recall and F1-score, particularly for the "less than 30 days readmittance" class, which represents cases where early readmission is crucial to identify. These results provide a baseline for our predictive modeling efforts, and further analysis and model refinements will be discussed in subsequent sections.

Figure 5 : A heatmap of a confusion matrix.



Description: This is a confusion matrix that was generated from our LASSO Logistic Regression.

Figure 6 : Receiver Operating Characteristics(ROC) Curve.



Description: This is a ROC Curve Plot with the False Positive Rate on the x-axis and True Positive Rate on the y-axis.

Figure 7 : Classification Report of LASSO Logistic Regression

Classification Report				
	Precision	Recall	F1- Score	Support
Not Readmitted	0.94	0.62	0.75	12966
Readmitted	0.10	0.49	0.16	1086
Accuracy			0.61	14052
Macro Average		0.56	0.46	14052
Weighted Average		0.61	0.70	14052

Description: Our classification report shows that our accuracy is 0.61 and our Macro Average F-1 Score is 0.46.

Lasso Logistic Regression Model – Under Sampling

For our second Lasso Logistic Regression model, we followed a similar procedure to the previous model but introduced a key modification: under sampling. Under sampling is a technique used to address class imbalance, which is particularly relevant in our binary classification problem due to the significant disproportion between the "no readmittance" and "less than 30 days readmittance" classes.

The steps in building this model included:

1. Grid Search for Optimal Alpha Value: We conducted a grid search to identify the optimal alpha value for Lasso Logistic Regression, maintaining the importance of regularization in feature selection and model simplification.

2. Train-Test Split: We again divided the dataset into training and testing subsets, ensuring an independent evaluation of model performance.

3. Feature Scaling and Dummitizing: Consistent with the previous model, we standardized numerical features and applied one-hot encoding to categorical variables for compatibility with logistic regression.

4. Under sampling: To mitigate the effects of class imbalance, we employed under sampling. This technique involved randomly selecting a subset of the majority class (in this case, "no readmittance") to balance the class distribution. By reducing the number of instances in the majority class, we aimed to ensure that the model did not disproportionately favor this class during training.

The introduction of under sampling in this Lasso Logistic Regression model allowed us to address the class imbalance issue more effectively. It enabled the model to learn from a more balanced dataset, potentially leading to better generalization and predictive performance, especially for the minority class ("less than 30 days readmittance"). In the subsequent sections, we will present the results of this model, including its evaluation and impact on predictive accuracy for patient readmission within 30 days.

Results

Optimal Alpha (Lambda): The grid search determined the best alpha value for Lasso regularization to be 0.0100, consistent with our first model.

Test Set Accuracy: The Lasso Logistic Regression model with under-sampling achieved an accuracy of 0.6600 on the test dataset. This metric indicates that the model correctly predicted the readmission outcomes for approximately 66% of the test cases.

User-Defined Threshold: To classify predictions into classes, a user-defined threshold of 0.98 was applied to the model's predicted probabilities.

Classification Report with Custom Threshold:

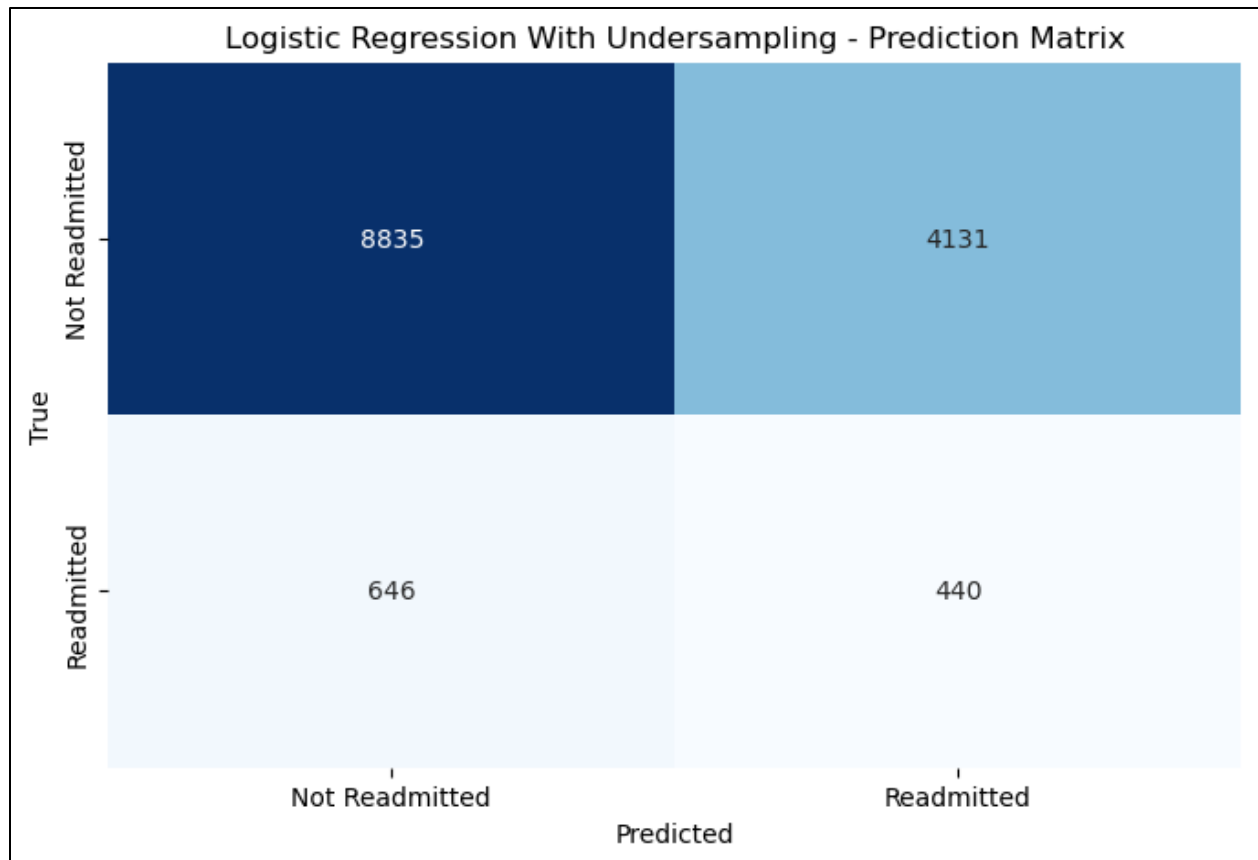
- **Precision:** The precision for class "0" (no readmittance) remained high at 0.93, indicating that when the model predicted "no readmittance," it was accurate 93% of the time. For class "1" (less than 30 days readmittance), the precision was 0.10, reflecting that the model's positive predictions for this class had a lower precision rate.

- **Recall:** The recall for class "0" was 0.68, suggesting that the model captured 68% of the actual "no readmittance" cases. The recall for class "1" was 0.41, indicating that the model identified 41% of the actual cases of "less than 30 days readmittance."
- **F1-Score:** The F1-score for class "0" improved to 0.79, demonstrating a better balance between precision and recall compared to our first model. However, the F1-score for class "1" remained at 0.16.
- **Support:** As in the previous model, class "0" had a substantially larger number of instances (12,966) compared to class "1" (1,086).

Macro and Weighted Averages: The macro-average F1-score was 0.47, reflecting the overall balance between precision and recall across both classes. The weighted-average F1-score, accounting for class imbalances, improved to 0.74.

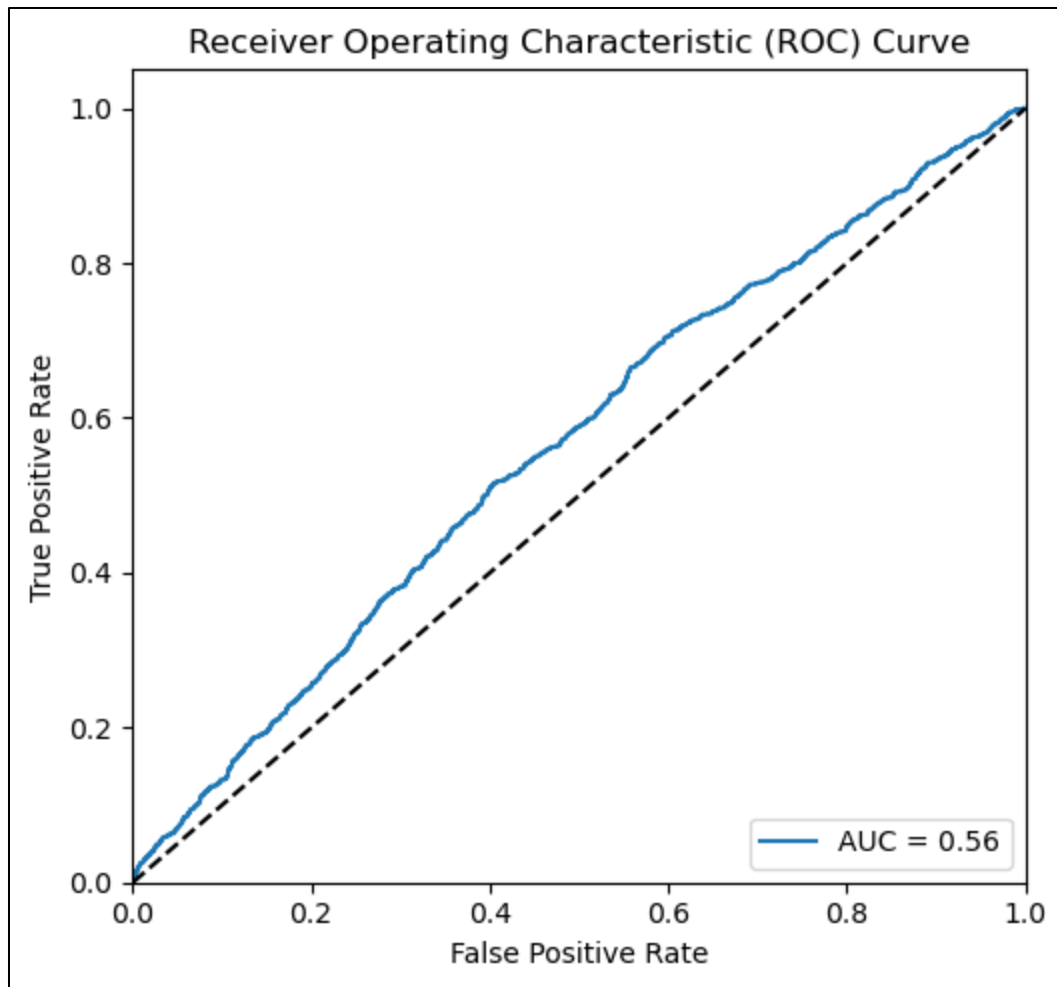
In summary, the Lasso Logistic Regression model with undersampling showed an improvement in accuracy compared to our first model. While precision remained high for class "0," the model's ability to identify cases of "less than 30 days readmittance" improved, as reflected in the higher recall for this class and a slightly improved F1-score. The undersampling technique helped address class imbalance, leading to a more balanced model. However, there is still room for further enhancement, especially in improving the model's ability to identify cases of early readmission. These results provide valuable insights for our ongoing analysis and model refinement efforts.

Figure 8 : *Confusion Matrix of Logistic Regression with UnderSampling.*



Description: Here is the confusion matrix of our Logistic Regression model with UnderSampling. It seems somewhat similar to our LASSO model.

Figure 9 : Receiver Operating Characteristics(ROC) Curve of Logistic Regression with UnderSampling.



Description: This is a ROC Curve Plot with the False Positive Rate on the x-axis and True Positive Rate on the y-axis.

Figure 10 : Classification Report of Logistic Regression with UnderSampling.

Classification Report				
	Precision	Recall	F1- Score	Support
Not Readmitted	0.93	0.68	0.79	12966
Readmitted	0.10	0.41	0.16	1086
Accuracy			0.66	14052
Macro Average	0.51	0.54	0.47	14052
Weighted Average	0.87	0.66	0.74	14052

Description: Classification Report for Logistic Regression with UnderSampling. From the report, we can see that our model has an improved accuracy of 0.66.

Random Forest

In addition to Lasso Logistic Regression, we explored the application of a Random Forest classifier as our third predictive model for patient readmission within 30 days of initial hospitalization. The Random Forest algorithm is an ensemble learning method that leverages the collective decision-making of multiple decision trees. This approach often yields robust and accurate predictions, making it a valuable tool for complex classification tasks like ours.

Our strategy for implementing the Random Forest model consisted of the following steps:

1. Hyperparameter Tuning: To optimize the performance of the Random Forest model, we conducted hyperparameter tuning. This involved exploring various settings, such as the number of trees in the forest and the maximum depth of each tree. We aimed to identify the combination of hyperparameters that would result in the most effective model.

2. Train-Test Split: Similar to our previous models, we divided the dataset into training and testing subsets. This separation allowed us to train the Random Forest model on one portion of the data and evaluate its performance on an independent dataset, ensuring that the model generalizes well to new, unseen instances.

3. Feature Importance Analysis: One of the advantages of Random Forest is its ability to provide insight into feature importance. We conducted a feature importance analysis to identify which variables had the most significant influence on predicting patient readmission. This analysis can be valuable for healthcare providers and decision-makers to prioritize interventions and allocate resources effectively.

4. Model Evaluation: We rigorously assessed the performance of the Random Forest model using a range of evaluation metrics, including accuracy, precision, recall, F1-score, and ROC-AUC. These metrics allowed us to gauge the model's predictive accuracy and its ability to balance sensitivity and specificity in predicting readmission outcomes.

5. Class Weight: In contrast to the under sampling technique utilized in the logistic regression model Random Forest Classifier has a parameter "class_weight" where we can set class weight equal to 'balanced'. This is a more convenient option compared to under/oversampling because it doesn't require modifying the training dataset's size. Instead, it automatically calculates class weights inversely proportional to the class frequencies in the training data. In other words, it assigns higher weights to the minority class and lower weights to the majority class.

By incorporating the Random Forest model into our analysis, we aimed to leverage its ensemble capabilities to capture complex relationships in the data and provide a more accurate and robust prediction of patient readmission within 30 days. The results of our Random Forest

model will be discussed in subsequent sections, shedding light on its effectiveness in enhancing patient care and healthcare resource allocation.

Results:

Our third model employed the Random Forest classifier with class weights set to 'balanced' to address the class imbalance issue. Here are the results obtained from this model:

Test Set Accuracy: The Random Forest model with balanced class weights achieved an accuracy of 0.6050 on the test dataset. This metric indicates that the model correctly predicted the readmission outcomes for approximately 60.50% of the test cases.

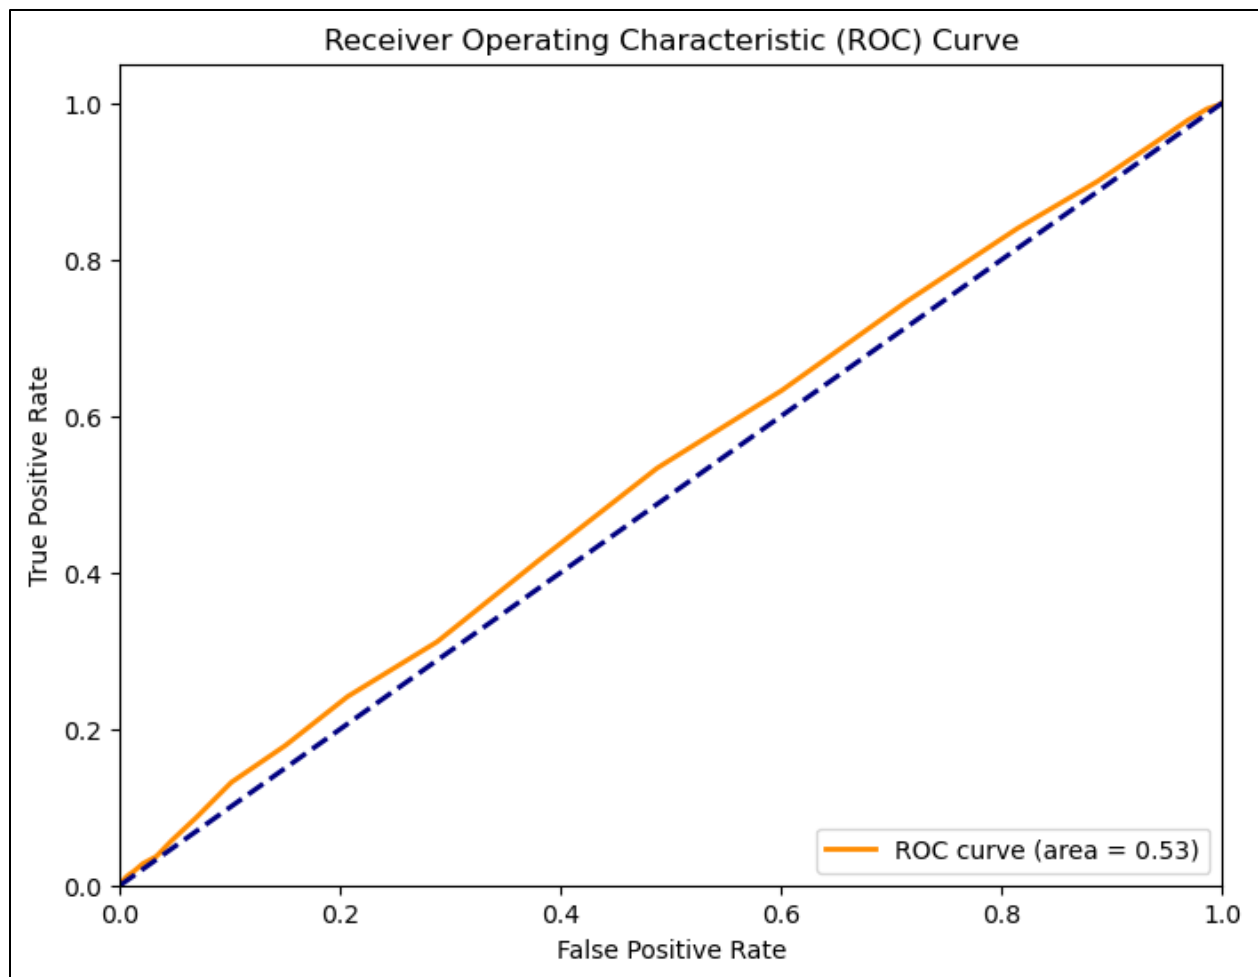
Classification Report:

- **Precision:** The precision for class "0" (no readmittance) remained high at 0.93, indicating that when the model predicted "no readmittance," it was accurate 93% of the time. However, the precision for class "1" (less than 30 days readmittance) was lower at 0.08, reflecting a lower precision rate for positive predictions.
- **Recall:** The recall for class "0" was 0.62, suggesting that the model captured 62% of the actual "no readmittance" cases. For class "1," the recall was 0.41, indicating that the model identified 41% of the actual cases of "less than 30 days readmittance."
- **F1-Score:** The F1-score for class "0" was 0.74, demonstrating a good balance between precision and recall. However, the F1-score for class "1" remained low at 0.14.
- **Support:** As in the previous models, class "0" had a significantly larger number of instances (12,966) compared to class "1" (1,086).

Macro and Weighted Averages: The macro-average F1-score was 0.44, reflecting the overall balance between precision and recall across both classes. The weighted-average F1-score, accounting for class imbalances, was 0.70.

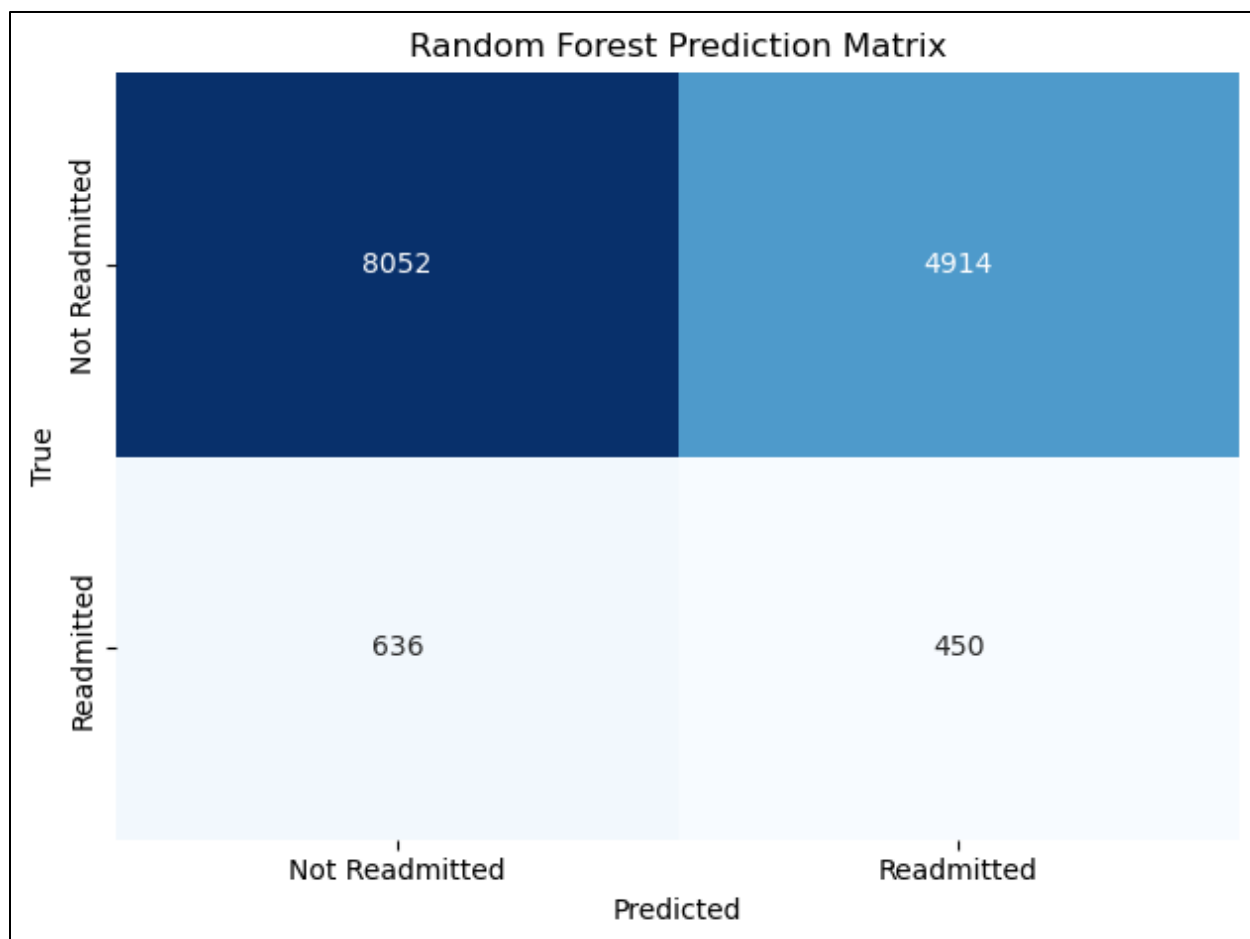
In summary, the Random Forest Classifier with balanced class weights exhibited a moderate accuracy in predicting patient readmission within 30 days. The model demonstrated high precision for class "0" (no readmittance), but the ability to identify cases of "less than 30 days readmittance" remained limited, as reflected in the lower recall and F1-score for class "1." While balanced class weights helped mitigate class imbalance, further model refinement may be needed to improve its performance in identifying early readmission cases. These results provide insights for ongoing analysis and potential model enhancements.

Figure 11(Next Page) : Receiver Operating Characteristics(ROC) Curve of Random Forest.



Description: This is a ROC Curve Plot with the False Positive Rate on the x-axis and True Positive Rate on the y-axis.

Figure 12(Next Page) :Confusion Matrix of Random Forest.



Description: Here is the confusion matrix of our Random Forest. It seems somewhat similar to our LASSO model.

Figure 13 : Classification Report of Random Forest.

Classification Report				
	Precision	Recall	F1- Score	Support
Not Readmitted	0.93	0.62	0.74	12966
Readmitted	0.08	0.41	0.14	1086
Accuracy			0.61	14052
Macro Average	0.51	0.52	0.44	14052
Weighted Average	0.86	0.61	0.70	14052

Description: Classification Report for Random Forest. From the report, we can see that our model has an improved accuracy of 0.61.

Conclusion

In this comprehensive analysis, we delved into the critical task of predicting patient readmission within 30 days of their initial hospitalization, a significant concern in the management of diabetes and other chronic conditions. Our study revolved around the evaluation of three distinct predictive models: Lasso Logistic Regression, Lasso Logistic Regression with Under sampling, and Random Forest Classifier with Balanced Class Weights. Each model brought its unique strengths and considerations to the table, contributing to our understanding of this complex problem.

The Lasso Logistic Regression model served as a foundational step, showcasing the importance of regularization in feature selection and model simplification. Despite reasonable accuracy, this initial model revealed the challenges of striking a balance between precision and recall, particularly for cases of early readmission. Feature engineering and custom thresholding were essential steps to fine-tune its performance.

The introduction of under sampling in our second model marked an effective attempt to mitigate class imbalance, providing a more balanced dataset for training. This adjustment yielded improvements in overall accuracy and recall for cases of early readmission, although there is still room for enhancement in precision and F1-score.

Our third model, the Random Forest Classifier with balanced class weights, showcased the robustness of ensemble learning in addressing complex classification tasks. While achieving a moderate level of accuracy, this model upheld high precision for class "0" predictions but struggled to effectively identify cases of "less than 30 days readmittance." Balancing class weights helped, yet further optimization is needed to bridge the gap between precision and recall for class "1."

In conclusion, the prediction of patient readmission within a 30-day window is a multifaceted challenge with significant clinical implications. Our models provide a foundation for understanding this problem, and their results shed light on the intricacies of balancing accuracy, precision, and recall, especially concerning early readmission cases.

Recommendations

Based on our analysis and findings, we propose several recommendations to enhance the predictive modeling of patient readmission within 30 days of initial hospitalization, with a focus on improving model performance and the practical utility of the predictions:

1. Feature Engineering and Selection:

- **Feature Engineering:** Continue to explore feature engineering techniques to create meaningful and predictive variables. Consider interactions between variables, time-based features, and domain-specific variables that may impact readmission risk.
- **Dimensionality Reduction:** Given the dimensionality of the dataset, consider employing dimensionality reduction techniques, or feature selection algorithms, to reduce noise and enhance model interpretability.

2. Dual Threshold Model:

- Develop a dual-threshold model that assigns probabilities or likelihood scores to the likelihood of readmission. This model can provide more nuanced insights by distinguishing between patients with a higher risk of early readmission and those with a lower risk. Implementing two thresholds allows healthcare providers to prioritize interventions more effectively.

3. Handling Class Imbalance:

- Continue experimenting with techniques to address class imbalance. Undersampling and oversampling methods, such as Synthetic Minority Over-sampling Technique (SMOTE), may further improve the model's ability to identify early readmission cases.

4. SME Collaboration:

- Collaborate closely with healthcare professionals to clinically validate the model's predictions. Incorporate domain expertise to refine the model's features, thresholds, and evaluation metrics to align better with the practical needs of healthcare providers.

5. Interpretability and Explainability:

- Prioritize the development of model interpretability and explainability techniques. Transparent models are essential for healthcare practitioners to understand and trust model predictions.

6. Data Quality and Collection:

- Invest in improving data quality by addressing missing data issues and conducting regular data audits. Collect additional relevant variables that might not be part of the current dataset but could have significant predictive power.

By implementing these recommendations, healthcare providers and data scientists can work together to develop more accurate, reliable, and clinically relevant models for predicting patient readmission. These efforts can ultimately contribute to improved patient care, better allocation of healthcare resources, and enhanced outcomes in the management of chronic conditions such as diabetes.

Appendix

Import Pacakges

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

Import Data

Diabetic data (main set)

```
In [2]: pd.set_option("display.max_columns", None)

diabetic_data = pd.read_csv('diabetic_data.csv')
diabetic_data.head()
```

```
Out[2]:
```

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	discharge_disposition_id	a
0	2278392	8222157	Caucasian	Female	[0-10)	?	6		25
1	149190	55629189	Caucasian	Female	[10-20)	?	1		1
2	64410	86047875	AfricanAmerican	Female	[20-30)	?	1		1
3	500364	82442376	Caucasian	Male	[30-40)	?	1		1
4	16680	42519267	Caucasian	Male	[40-50)	?	1		1

ID dataset (3 of them)

```
In [3]: id_data = pd.read_csv('IDs_mapping.csv', header = None)
id_data.head()
```

```
Out[3]:
```

	0	1
0	admission_type_id	description
1	1	Emergency
2	2	Urgent
3	3	Elective
4	4	Newborn

```
In [4]: # dataset one - admission_type
empty_rows = id_data[id_data.isna().all(axis=1)].index.tolist()

id_admission_type = id_data.iloc[:empty_rows[0]]
id_admission_type.columns = id_admission_type.iloc[0] # get the column to be first row
id_admission_type = id_admission_type.drop(id_admission_type.index[0]) # remove the first row
id_admission_type = id_admission_type.reset_index(drop=True) # reset index

id_admission_type.head(20)
```

```
Out[4]:
```

	admission_type_id	description
0	1	Emergency
1	2	Urgent
2	3	Elective
3	4	Newborn
4	5	Not Available
5	6	NaN
6	7	Trauma Center
7	8	Not Mapped

```
In [5]: # dataset two - discharge id
id_discharge = id_data.iloc[empty_rows[0]+1 : empty_rows[1]]
id_discharge.columns = id_discharge.iloc[0]
id_discharge = id_discharge.drop(id_discharge.index[0])
id_discharge = id_discharge.reset_index(drop=True)

id_discharge.head(31)
```

```
Out[5]:
```

	10	discharge_disposition_id	description
0	1		Discharged to home
1	2		Discharged/transferred to another short term h...
2	3		Discharged/transferred to SNF
3	4		Discharged/transferred to ICF
4	5		Discharged/transferred to another type of inpa...
5	6		Discharged/transferred to home with home healt...
6	7		Left AMA
7	8		Discharged/transferred to home under care of H...
8	9		Admitted as an inpatient to this hospital
9	10		Neonate discharged to another hospital for neo...
10	11		Expired
11	12		Still patient or expected to return for outpat...
12	13		Hospice / home
13	14		Hospice / medical facility
14	15		Discharged/transferred within this institution...

15	16	Discharged/transferred/referred another instit...
16	17	Discharged/transferred/referred to this instit...
17	18	NaN
18	19	Expired at home. Medicaid only, hospice.
19	20	Expired in a medical facility. Medicaid only, ...
20	21	Expired, place unknown. Medicaid only, hospice.
21	22	Discharged/transferred to another rehab fac in...
22	23	Discharged/transferred to a long term care hos...
23	24	Discharged/transferred to a nursing facility c...
24	25	Not Mapped
25	26	Unknown/Invalid
26	30	Discharged/transferred to another Type of Heal...
27	27	Discharged/transferred to a federal health car...
28	28	Discharged/transferred/referred to a psychiatr...
29	29	Discharged/transferred to a Critical Access Ho...

```
In [6]: # dataset three - admission source id
id_admission_source = id_data.iloc[empty_rows[1]+1 : ]
id_admission_source.columns = id_admission_source.iloc[0]
id_admission_source = id_admission_source.drop(id_admission_source.index[0])
id_admission_source = id_admission_source.reset_index(drop = True)
id_admission_source.head(26)
```

```
Out[6]:
```

42	admission_source_id	description
0	1	Physician Referral
1	2	Clinic Referral
2	3	HMO Referral
3	4	Transfer from a hospital
4	5	Transfer from a Skilled Nursing Facility (SNF)
5	6	Transfer from another health care facility
6	7	Emergency Room
7	8	Court/Law Enforcement
8	9	Not Available
9	10	Transfer from critial access hospital
10	11	Normal Delivery
11	12	Premature Delivery
12	13	Sick Baby
13	14	Extramural Birth
14	15	Not Available
15	17	NaN
16	18	Transfer From Another Home Health Agency

17	19	Readmission to Same Home Health Agency
18	20	Not Mapped
19	21	Unknown/Invalid
20	22	Transfer from hospital inpt/same fac reslt in...
21	23	Born inside this hospital
22	24	Born outside this hospital
23	25	Transfer from Ambulatory Surgery Center
24	26	Transfer from Hospice

Merge Datasets

```
In [7]: id_admission_type['admission_type_id'] = id_admission_type['admission_type_id'].astype(int)
id_discharge['discharge_disposition_id'] = id_discharge['discharge_disposition_id'].astype(int)
id_admission_source['admission_source_id'] = id_admission_source['admission_source_id'].astype(int)
```

```
In [8]: df = pd.merge(diabetic_data, id_admission_type, how='inner', on='admission_type_id')
df = pd.merge(df, id_discharge, how='inner', on='discharge_disposition_id') # description
df = pd.merge(df, id_admission_source, how='inner', on='admission_source_id') # description
```

```
In [9]: df.rename(columns={'description': 'admission_desc',
                           'description_x': 'discharge_desc',
                           'description_y': 'admission_source_desc'}, inplace=True)
```

```
In [10]: df.head()
```

```
Out[10]:
```

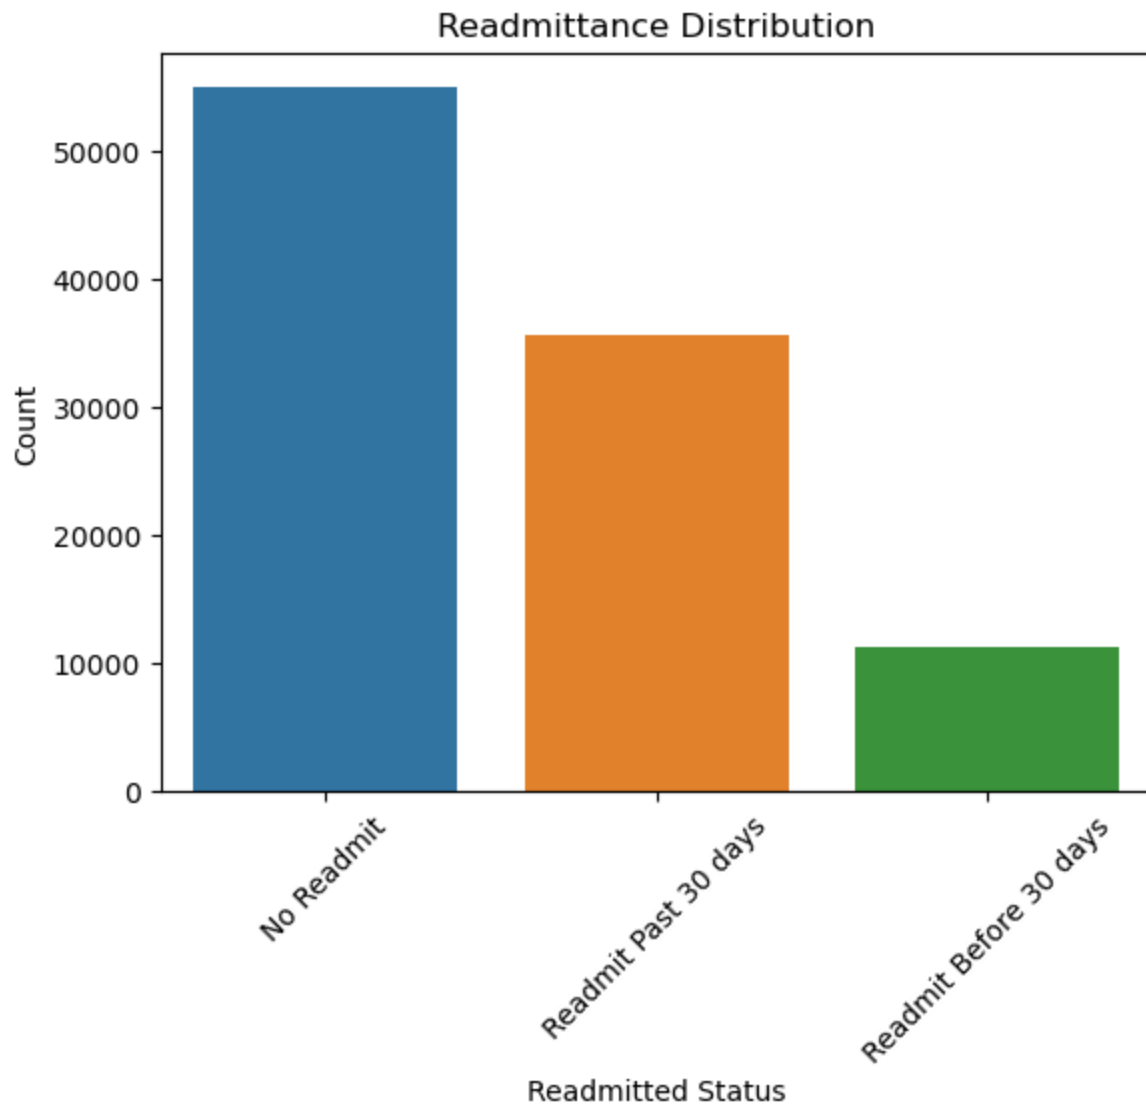
	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	discharge_disposition_id	admission_source_desc
0	2278392	8222157	Caucasian	Female	[0-10)	?	6	25	
1	1968528	720936	Caucasian	Female	[70-80)	?	6	25	
2	2223336	558360	AfricanAmerican	Female	[60-70)	?	6	25	
3	2298006	2519748	Caucasian	Male	[60-70)	?	6	25	
4	2356308	608841	AfricanAmerican	Female	[50-60)	?	6	25	

Target Variable Inspection

```
In [11]: custom_x = ['No Readmit', 'Readmit Past 30 days', 'Readmit Before 30 days']

ax = sns.countplot(df, x='readmitted')
ax.set_xticklabels(custom_x)
plt.xticks(rotation=45)
```

```
plt.title('Readmittance Distribution')
plt.xlabel('Readmitted Status')
plt.ylabel('Count')
plt.show()
```



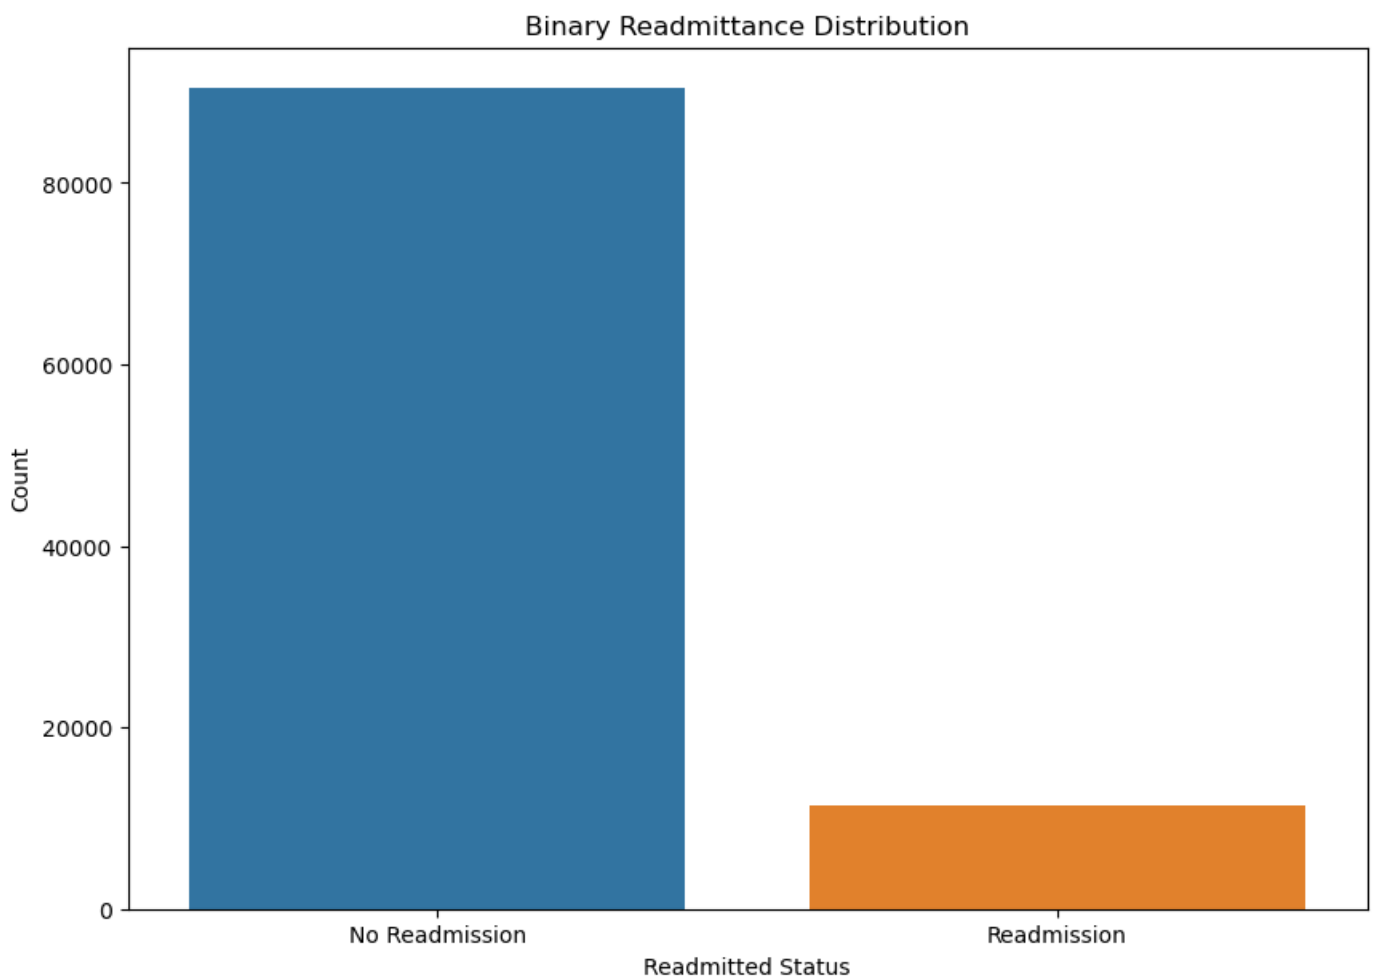
Binary Target Variable

- Because we are looking to classify readmittance within 30 days, we will make this a binary type of target instead of multiclass.
- This will look like "readmitted_30? = Yes or No"

```
In [12]: # df['binary_readmitted'] = df['readmitted'].apply(lambda x: 'yes' if x != 'NO' else 'no')
df['binary_readmitted'] = df['readmitted'].apply(lambda x: 1 if x == '<30' else 0)
```

```
In [13]: custom_x = ['No Readmission', 'Readmission']

plt.figure(figsize=(10,7))
ax = sns.countplot(df, x = 'binary_readmitted')
ax.set_xticklabels(custom_x)
plt.title('Binary Readmittance Distribution')
plt.xlabel('Readmitted Status')
plt.ylabel('Count')
plt.show()
```



Missing Values

Examining Value Counts

```
In [14]: value_counts_dict = {}
for column in df.columns:
    value_counts_dict[column] = df[column].value_counts()

for column, i in value_counts_dict.items():
    print(f'Column: {column}')
    print(i)
    print('\n')
```

```
Column: encounter_id
2278392      1
217273374    1
217188990    1
217165278    1
217146714    1
..
48648786     1
48582588     1
48525978     1
48513912     1
370756376    1
Name: encounter_id, Length: 101766, dtype: int64
```

```

Column: patient_nbr
88785891      40
43140906      28
1660293       23
23199021      23
88227540      23
..
112244895     1
5899338       1
108997866     1
3452139       1
87775947      1
Name: patient_nbr, Length: 71518, dtype: int64

```

```

Column: race
Caucasian      76099
AfricanAmerican 19210
?              2273
Hispanic       2037
Other          1506
Asian          641
Name: race, dtype: int64

```

```

Column: gender
Female      54708
Male       47055
Unknown/Invalid 3
Name: gender, dtype: int64

```

```

Column: age
[70-80)  26068
[60-70)  22483
[50-60)  17256
[80-90)  17197
[40-50)   9685
[30-40)   3775
[90-100)  2793
[20-30)   1657
[10-20)    691
[0-10)    161
Name: age, dtype: int64

```

```

Column: weight
?      98569
[75-100)  1336
[50-75)   897
[100-125)  625
[125-150)  145
[25-50)    97
[0-25)     48
[150-175)  35
[175-200)  11
>200       3
Name: weight, dtype: int64

```

```

Column: admission_type_id
1    53990
3    18869
2    18480

```

```
6      5291
5      4785
8       320
7       21
4       10
Name: admission_type_id, dtype: int64
```

Column: discharge_disposition_id

```
1      60234
3      13954
6      12902
18     3691
2      2128
22     1993
11     1642
5      1184
25     989
4      815
7      623
23     412
13     399
14     372
28     139
8      108
15     63
24     48
9      21
17     14
16     11
19     8
10     6
27     5
12     3
20     2
Name: discharge_disposition_id, dtype: int64
```

Column: admission_source_id

```
7      57494
1      29565
17     6781
4      3187
6      2264
2      1104
5      855
3      187
20     161
9      125
8      16
22     12
10     8
14     2
11     2
25     2
13     1
Name: admission_source_id, dtype: int64
```

Column: time_in_hospital

```
3      17756
2      17224
1      14208
4      13924
5      9966
6      7539
```



```
7      5859
8      4391
9      3002
10     2342
11     1855
12     1448
13     1210
14     1042
Name: time_in_hospital, dtype: int64
```

Column: payer_code

```
?      40256
MC     32439
HM     6274
SP     5007
BC     4655
MD     3532
CP     2533
UN     2448
CM     1937
OG     1033
PO      592
DM      549
CH      146
WC      135
OT       95
MP       79
SI       55
FR        1
Name: payer_code, dtype: int64
```

Column: medical_specialty

```
?                                49949
InternalMedicine                 14635
Emergency/Trauma                 7565
Family/GeneralPractice           7440
Cardiology                      5352
...
Surgery-PlasticwithinHeadandNeck 1
Neurophysiology                 1
Speech                          1
Psychiatry-Addictive            1
Pediatrics-InfectiousDiseases   1
Name: medical_specialty, Length: 73, dtype: int64
```

Column: num_lab_procedures

```
1      3208
43     2804
44     2496
45     2376
38     2213
...
118     1
121     1
126     1
132     1
129     1
Name: num_lab_procedures, Length: 118, dtype: int64
```

Column: num_procedures

```
0     46652
1     20742
```

```
2      12717
3      9443
6      4954
4      4180
5      3078
Name: num_procedures, dtype: int64
```

Column: num_medications

```
13      6086
12      6004
11      5795
15      5792
14      5707
```

...

```
75      2
70      2
79      1
81      1
74      1
```

Name: num_medications, Length: 75, dtype: int64

Column: number_outpatient

```
0      85027
1      8547
2      3594
3      2042
4      1099
5       533
6       303
7       155
8        98
9        83
10       57
11       42
13       31
12       30
14       28
15       20
16       15
17        8
21        7
20        7
22        5
18        5
27        3
24        3
19        3
23        2
25        2
26        2
29        2
36        2
33        2
35        2
42        1
34        1
40        1
38        1
37        1
28        1
39        1
```

Name: number_outpatient, dtype: int64

Column: number_emergency

0	90383
1	7677
2	2042
3	725
4	374
5	192
6	94
7	73
8	50
10	34
9	33
11	23
13	12
12	10
22	6
16	5
18	5
19	4
20	4
14	3
15	3
21	2
25	2
28	1
42	1
76	1
37	1
29	1
46	1
64	1
63	1
54	1
24	1

Name: number_emergency, dtype: int64

Column: number_inpatient

0	67630
1	19521
2	7566
3	3411
4	1622
5	812
6	480
7	268
8	151
9	111
10	61
11	49
12	34
13	20
14	10
15	9
16	6
19	2
21	1
18	1
17	1

Name: number_inpatient, dtype: int64

Column: diag_1

428	6862
414	6581
786	4016

410	3614
486	3508
...	
640	1
314	1
817	1
61	1
915	1

Name: diag_1, Length: 717, dtype: int64

Column: diag_2

276	6752
428	6662
250	6071
427	5036
401	3736
...	
256	1
752	1
316	1
963	1
917	1

Name: diag_2, Length: 749, dtype: int64

Column: diag_3

250	11555
401	8289
276	5175
428	4577
427	3955
...	
744	1
732	1
E945	1
122	1
193	1

Name: diag_3, Length: 790, dtype: int64

Column: number_diagnoses

9	49474
5	11393
8	10616
7	10393
6	10161
4	5537
3	2835
2	1023
1	219
16	45
10	17
13	16
11	11
15	10
12	9
14	7

Name: number_diagnoses, dtype: int64

Column: max_glu_serum

None	96420
Norm	2597
>200	1485
>300	1264

Name: max_glu_serum, dtype: int64

Column: A1Cresult

None 84748

>8 8216

Norm 4990

>7 3812

Name: A1Cresult, dtype: int64

Column: metformin

No 81778

Steady 18346

Up 1067

Down 575

Name: metformin, dtype: int64

Column: repaglinide

No 100227

Steady 1384

Up 110

Down 45

Name: repaglinide, dtype: int64

Column: nateglinide

No 101063

Steady 668

Up 24

Down 11

Name: nateglinide, dtype: int64

Column: chlorpropamide

No 101680

Steady 79

Up 6

Down 1

Name: chlorpropamide, dtype: int64

Column: glimepiride

No 96575

Steady 4670

Up 327

Down 194

Name: glimepiride, dtype: int64

Column: acetohexamide

No 101765

Steady 1

Name: acetohexamide, dtype: int64

Column: glipizide

No 89080

Steady 11356

Up 770

Down 560

Name: glipizide, dtype: int64

Column: glyburide

No 91116
Steady 9274
Up 812
Down 564
Name: glyburide, dtype: int64

Column: tolbutamide
No 101743
Steady 23
Name: tolbutamide, dtype: int64

Column: pioglitazone
No 94438
Steady 6976
Up 234
Down 118
Name: pioglitazone, dtype: int64

Column: rosiglitazone
No 95401
Steady 6100
Up 178
Down 87
Name: rosiglitazone, dtype: int64

Column: acarbose
No 101458
Steady 295
Up 10
Down 3
Name: acarbose, dtype: int64

Column: miglitol
No 101728
Steady 31
Down 5
Up 2
Name: miglitol, dtype: int64

Column: troglitazone
No 101763
Steady 3
Name: troglitazone, dtype: int64

Column: tolazamide
No 101727
Steady 38
Up 1
Name: tolazamide, dtype: int64

Column: examide
No 101766
Name: examide, dtype: int64

Column: citoglipton
No 101766
Name: citoglipton, dtype: int64

Column: insulin
No 47383
Steady 30849
Down 12218
Up 11316
Name: insulin, dtype: int64

Column: glyburide-metformin
No 101060
Steady 692
Up 8
Down 6
Name: glyburide-metformin, dtype: int64

Column: glipizide-metformin
No 101753
Steady 13
Name: glipizide-metformin, dtype: int64

Column: glimepiride-pioglitazone
No 101765
Steady 1
Name: glimepiride-pioglitazone, dtype: int64

Column: metformin-rosiglitazone
No 101764
Steady 2
Name: metformin-rosiglitazone, dtype: int64

Column: metformin-pioglitazone
No 101765
Steady 1
Name: metformin-pioglitazone, dtype: int64

Column: change
No 54755
Ch 47011
Name: change, dtype: int64

Column: diabetesMed
Yes 78363
No 23403
Name: diabetesMed, dtype: int64

Column: readmitted
NO 54864
>30 35545
<30 11357
Name: readmitted, dtype: int64

Column: discharge_desc
Emergency 53990
Elective 18869
Urgent 18480
Not Available 4785

Not Mapped 320
 Trauma Center 21
 Newborn 10
 Name: discharge_desc, dtype: int64

Column: admission_source_desc
 Discharged to home 60234
 Discharged/transferred to SNF 13954
 Discharged/transferred to home with home health service 12902
 Discharged/transferred to another short term hospital 2128
 Discharged/transferred to another rehab fac including rehab units of a hospital . 1993
 Expired 1642
 Discharged/transferred to another type of inpatient care institution 1184
 Not Mapped 989
 Discharged/transferred to ICF 815
 Left AMA 623
 Discharged/transferred to a long term care hospital. 412
 Hospice / home 399
 Hospice / medical facility 372
 Discharged/transferred/referred to a psychiatric hospital of psychiatric distinct part u
 nit of a hospital 139
 Discharged/transferred to home under care of Home IV provider 108
 Discharged/transferred within this institution to Medicare approved swing bed 63
 Discharged/transferred to a nursing facility certified under Medicaid but not certified
 under Medicare. 48
 Admitted as an inpatient to this hospital 21
 Discharged/transferred/referred to this institution for outpatient services 14
 Discharged/transferred/referred another institution for outpatient services 11
 Expired at home. Medicaid only, hospice. 8
 Neonate discharged to another hospital for neonatal aftercare 6
 Discharged/transferred to a federal health care facility. 5
 Still patient or expected to return for outpatient services 3
 Expired in a medical facility. Medicaid only, hospice. 2
 Name: admission_source_desc, dtype: int64

Column: admission_desc
 Emergency Room 57494
 Physician Referral 29565
 Transfer from a hospital 3187
 Transfer from another health care facility 2264
 Clinic Referral 1104

Transfer from a Skilled Nursing Facility (SNF)	855
HMO Referral	187
Not Mapped	161
Not Available	125
Court/Law Enforcement	16
Transfer from hospital inpt/same fac reslt in a sep claim	12
Transfer from critial access hospital	8
Extramural Birth	2
Normal Delivery	2
Transfer from Ambulatory Surgery Center	2
Sick Baby	1

Name: admission_desc, dtype: int64

Column: binary_readmitted

0	90409
1	11357

Name: binary_readmitted, dtype: int64

Features that may be redundant or contain redundant values

- citoglipton - single value
- examide - single value
- Gender - remove the unknown gender
- metformin-pioglitazone - binary
- metformin-rosiglitazone - binary
- glimepiride-pioglitazone - binary
- glipizide-metformin - binary
- tolazamide - binary
- troglitazone - binary
- tolbutamide - binary
- acetohexamide - binary

```
In [15]: df = df[df['gender']!='Unknown/Invalid'] # taking out the 'unknown/invalid' gender obser

# dropping observations where the values were insignificant
df.drop(['citoglipton','examide','metformin-pioglitazone',
        'metformin-rosiglitazone','glimepiride-pioglitazone',
        'glipizide-metformin','tolazamide','troglitazone',
        'tolbutamide','acetohexamide','readmitted'], axis = 1, inplace = True)

In [16]: # some patients have been back multiple times. We are trying to determine readmittance,
# when patients have come back by keeping the first return visit rather than every singl
df = df.drop_duplicates(subset = 'patient_nbr', keep = 'first')

print(df.shape) # looking at how the shape has been reduced thus far

(71515, 43)
```

Hidden N/A Values

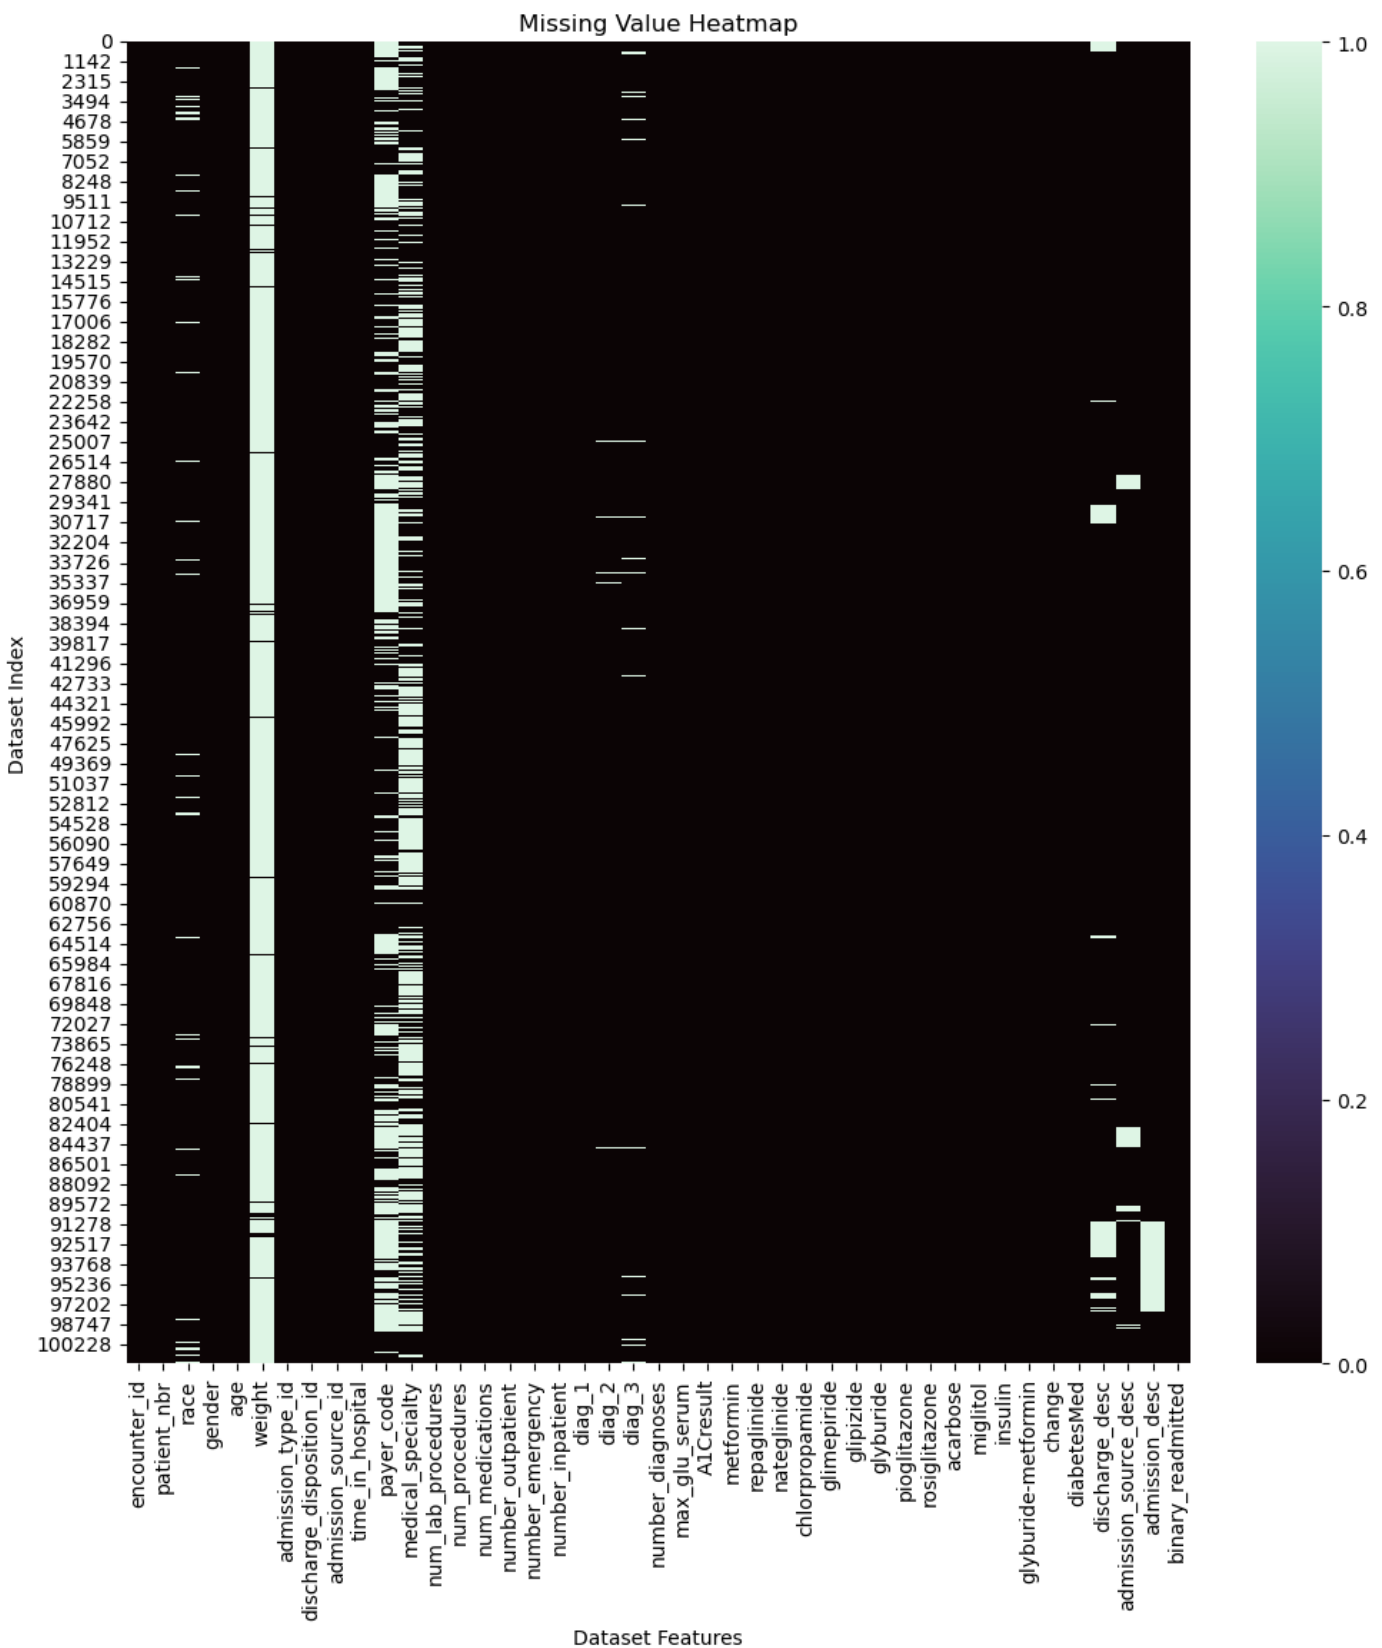
- **race:** contain ?
 - Upon further investigation (visually and with cross tabulations) we have determined this data to be MCAR (Missing Completely at Random)

- We will proceed with imputation of the most frequent category
- **weight:** > 90% missing values
 - Upon further investigation (visually and with cross tabulations) we have determined this data to be MCAR (Missing Completely at Random)
 - We will proceed with the removal of this feature.
- **payer_code:** contain ?
 - Upon Further Investigation (visually and with cross tabulations) we have determined that this data is MCAR.
 - We will not impute these values, or remove them as there is a significant amount missing, instead we will re-label the missing values and include them as a feature in our data.
- **medical_specialty:** contain ?
 - There is seemingly no relationship between the values in our data and the missingness of our medical specialty values.
 - We will proceed with imputing the missing values as their own value of "other".
- **diag-1,2,3:** all contain '?'
 - There is an inconsequential amount of missing data in the Diagnosis 1,2,3. We will proceed with removing the null values.

```
In [17]: # we are replacing '?' with NaN so that we can visually see where the missing data is ha
df['race'].replace('?', np.nan, inplace = True)
df['weight'].replace('?', np.nan, inplace = True)
df['payer_code'].replace('?', np.nan, inplace = True)
df['medical_specialty'].replace('?', np.nan, inplace = True)
df['diag_1'].replace('?', np.nan, inplace = True)
df['diag_2'].replace('?', np.nan, inplace = True)
df['diag_3'].replace('?', np.nan, inplace = True)
```

Heatmap for Missing Values

```
In [18]: plt.figure(figsize=(12,12))
sns.heatmap(df.isnull(), cbar = True, cmap = 'mako')
plt.title('Missing Value Heatmap')
plt.xlabel('Dataset Features')
plt.ylabel('Dataset Index')
plt.show()
```



Discharge Desc. & Admission Desc.

- The heatmap and cross tabulations show that there is some relationship between the missing values in Discharge Desc. & Admission Desc.
- This relationship somewhat makes sense intuitively because if there is missing information regarding Admission, it would make sense that Discharge missing information could be related.
- Ultimately this relationship is only a partial correlation and isn't quite considered MNAR (missing not at random). We will proceed with filling the NaN values for these as "other".

```
In [19]: cont_table = pd.crosstab(df['discharge_desc'].isna(), df['admission_desc'].isna())
cont_table
```

```
Out[19]: admission_desc    False    True
```

```
discharge_desc
```

```
False    64622    2307
```

```
True      2004    2582
```

```
In [20]: df[(df['discharge_desc'] == 'NAN') & (df['admission_type_id'] == 6)]
```

```
Out[20]: encounter_id  patient_nbr  race  gender  age  weight  admission_type_id  discharge_disposition_id  admission_so
```

Payer Code Vs Medical Specialty

- Looking at the heatmap above it appears that there are some instances where payer code and medical specialty flip between missing observations.
- to further explore if maybe the missingness can be determined by the missingness of the other we will create a contingency table.

```
In [21]: cont_table = pd.crosstab(df['payer_code'].isnull(), df['medical_specialty'].isnull())
cont_table
```

```
Out[21]: medical_specialty    False    True
```

```
payer_code
```

```
False    19288    21474
```

```
True      18202    12551
```

From this table, we can observe that when 'payer_code' is not missing (False), 'medical_specialty' can be either not missing (False) or missing (True). Similarly, when 'payer_code' is missing (True), 'medical_specialty' can also be either not missing (False) or missing (True).

There doesn't seem to be a strong relationship indicating that when one feature is missing, the other is not missing, or vice versa. In other words, the missingness of 'payer_code' and 'medical_specialty' does not appear to be dependent on each other.

Define Percent Missing

```
In [22]: def percent_missing(df):
percent_nan = 100*df.isnull().sum()/len(df)
percent_nan = percent_nan[percent_nan > 0].sort_values()
return(percent_nan)
percent_nan = percent_missing(df)
print(round(percent_nan,2))
```

```
diag_1                0.02
```

```
diag_2                0.41
```

```
diag_3                1.68
```

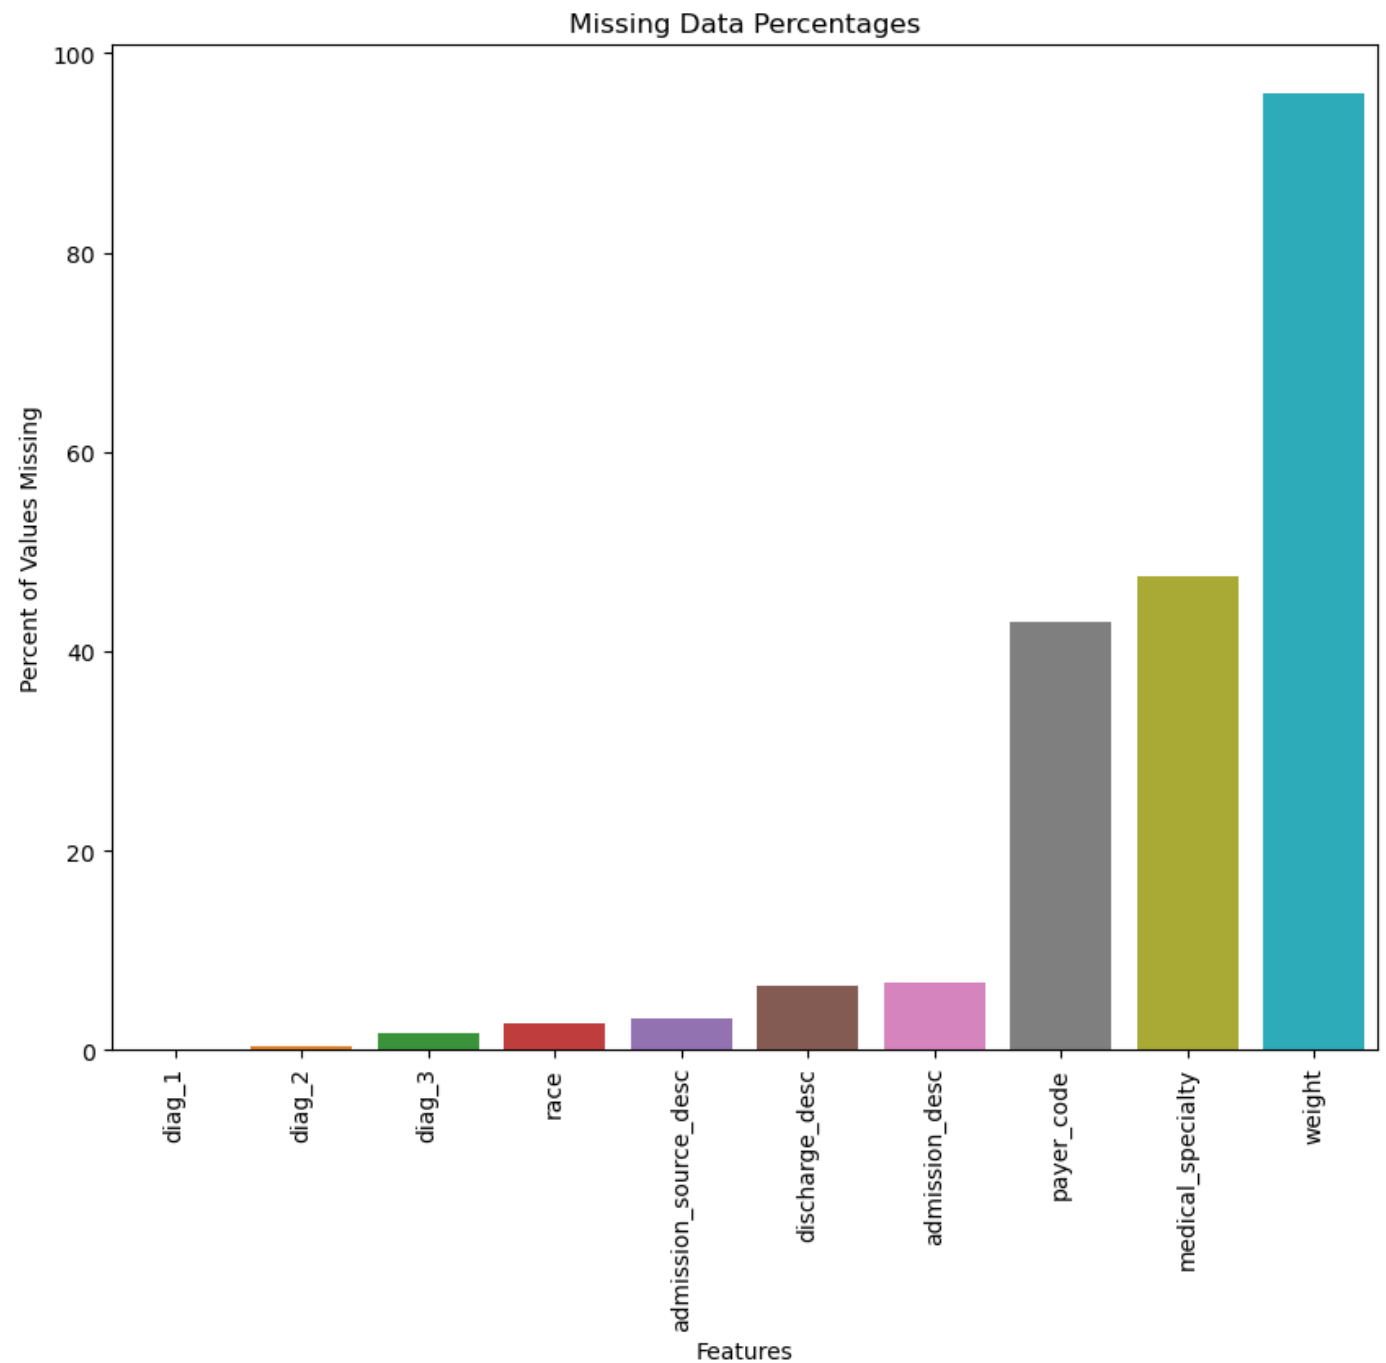
```
race                  2.67
```

```
admission_source_desc 3.22
```

```
discharge_desc      6.41
admission_desc      6.84
payer_code          43.00
medical_specialty    47.58
weight              96.02
dtype: float64
```

Percent Missing Bar Plot

```
In [23]: plt.figure(figsize = (10,8))
sns.barplot(x= percent_nan.index, y = percent_nan)
plt.title('Missing Data Percentages')
plt.xlabel('Features')
plt.ylabel('Percent of Values Missing')
plt.xticks(rotation = 90)
plt.show()
```



```
In [24]: # filling NaN with "NaN" so that we can see in our data where these values are occurring.
# this gives us a means of representing the missing data which is useful in investigatio
```

```

#impute race missing values with mode
mode_race = df['race'].mode()[0]
df['race'].fillna(mode_race, inplace = True)

# drop the weight column
df.drop('weight', inplace = True, axis = 1)

# payer_code impute with it's own level to maintain information in data.
df['payer_code'].fillna('Not Specified', inplace = True)

df['medical_specialty'].fillna('Other', inplace = True)

# fill NaN with 'other'
df['admission_desc'].fillna('Other', inplace = True)
df['admission_source_desc'].fillna('Other', inplace = True)
df['discharge_desc'].fillna('Other', inplace = True)

# remove NaN from Diag 1-3
df.dropna(subset=['diag_1', 'diag_2', 'diag_3'], inplace = True)

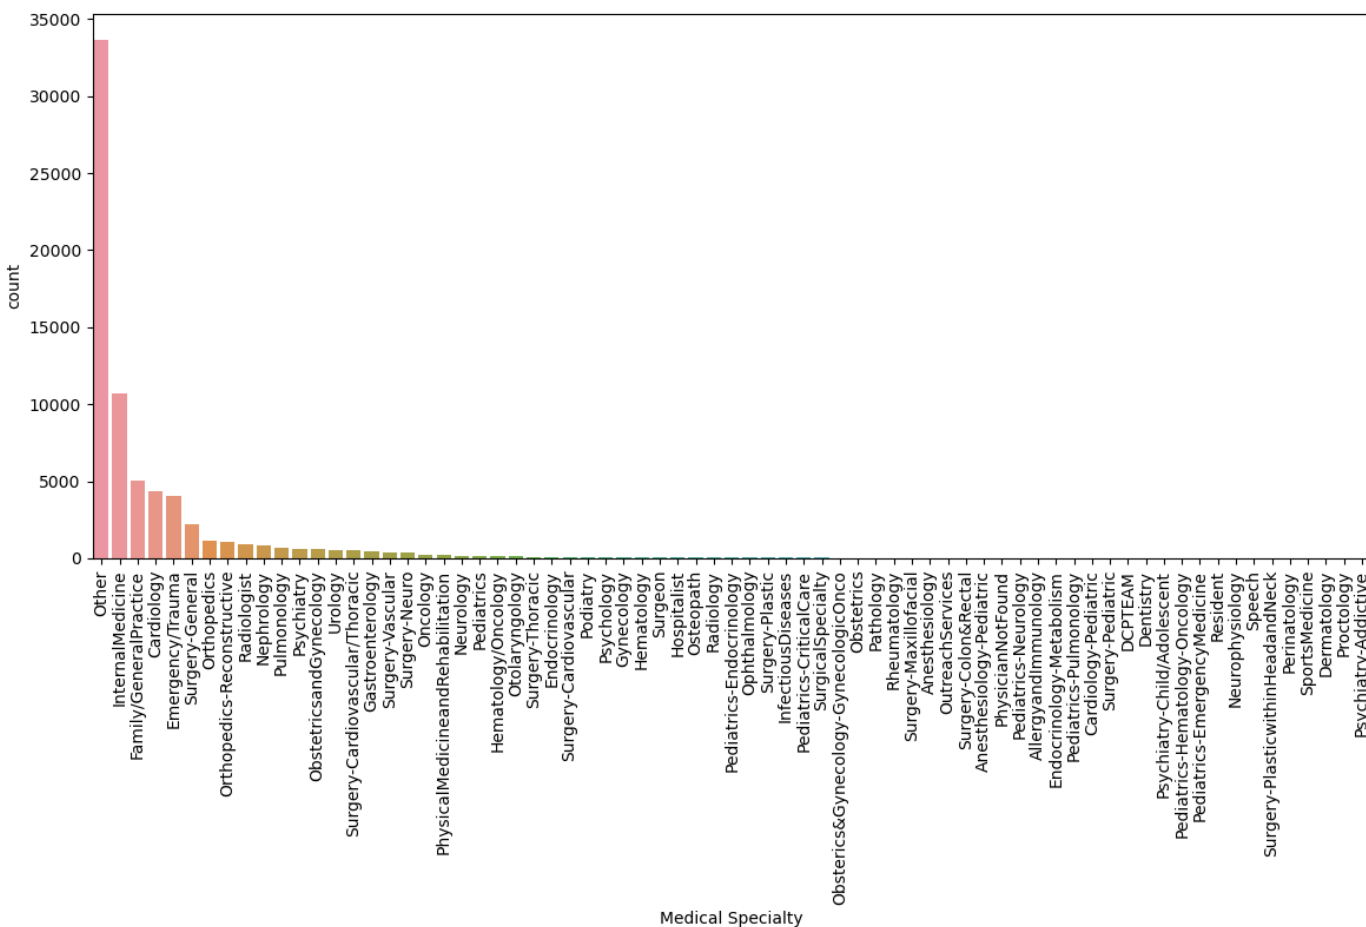
```

Medical Specialty - With use of "other" value to represent the NaN

```

In [25]: plt.figure(figsize=(14,6))
sns.countplot(df, x = 'medical_specialty', order = df['medical_specialty'].value_counts()
plt.xticks(rotation = 90)
plt.xlabel('Medical Specialty')
plt.show()

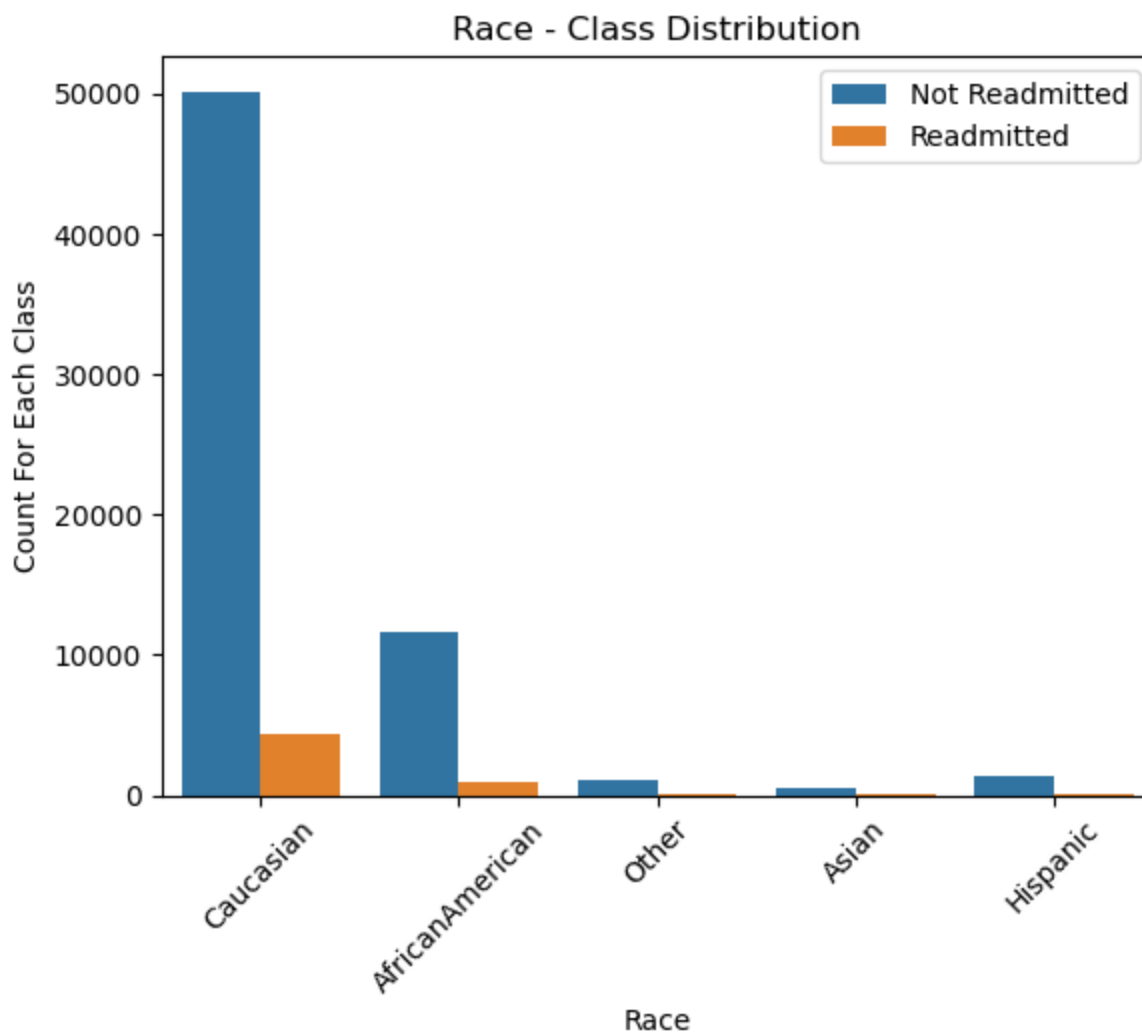
```



Race - With imputed Mode

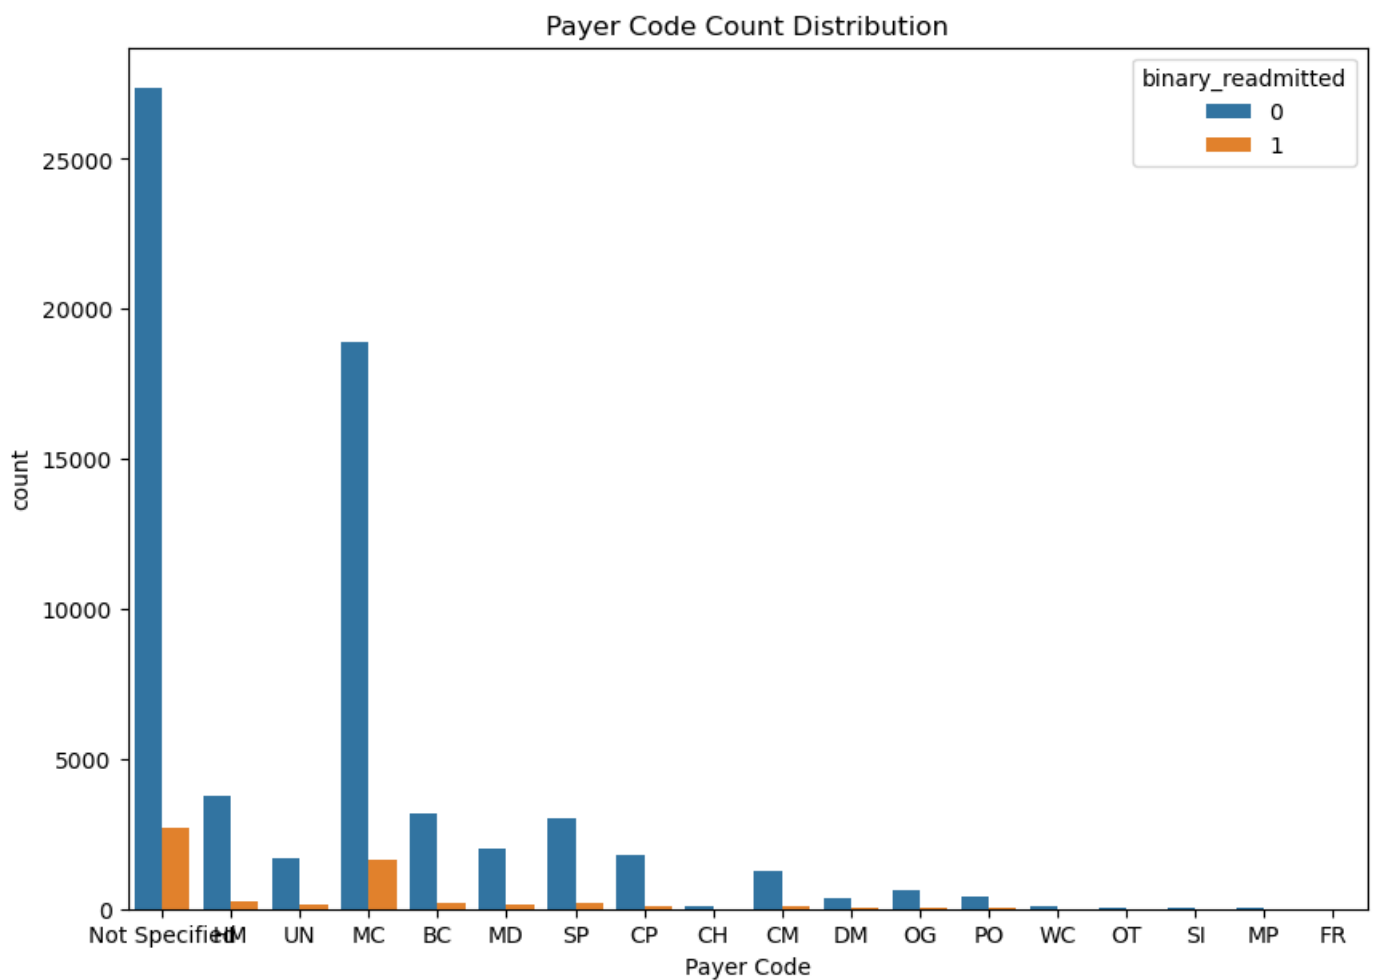
```
In [26]: legend_labels = ['Not Readmitted', 'Readmitted']

ax = sns.countplot(df, x = 'race', hue = 'binary_readmitted')
ax.legend(legend_labels)
plt.xticks(rotation = 45)
plt.title('Race - Class Distribution')
plt.xlabel('Race')
plt.ylabel('Count For Each Class')
plt.show()
```



Payer Code - Relabeld NaN as Not Specified

```
In [27]: plt.figure(figsize=(10,7))
sns.countplot(df, x = 'payer_code', hue='binary_readmitted')
plt.title('Payer Code Count Distribution')
plt.xlabel('Payer Code')
plt.show()
```



[Click Here For information on Diagnosis Codes](#)

```
In [28]: # diag_1
custom_categories = {
    'Diabetes': ['250', '250.02', '250.03', '250.04', '250.1', '250.11', '250.12', '250.
    'Heart Disease': ['410', '411', '412', '413', '414', '415', '416', '417', '418', '41
    'Respiratory': ['460', '461', '462', '463', '464', '465', '466', '470', '471', '472'
    'Injury': ['800', '801', '802', '803', '804', '805', '806', '807', '808', '809', '81
    'Other': ['E909', 'V25', 'V26', 'V43', 'V45', 'V51', 'V53', 'V54', 'V55', 'V56', 'V5

}

# Create a new column 'diag_1_new' and set it to 'Other' initially
df['diag_1_new'] = 'Other'

# Map the values in 'diag_1' to custom categories
for category, codes in custom_categories.items():
    df.loc[df['diag_1'].isin(codes), 'diag_1_new'] = category

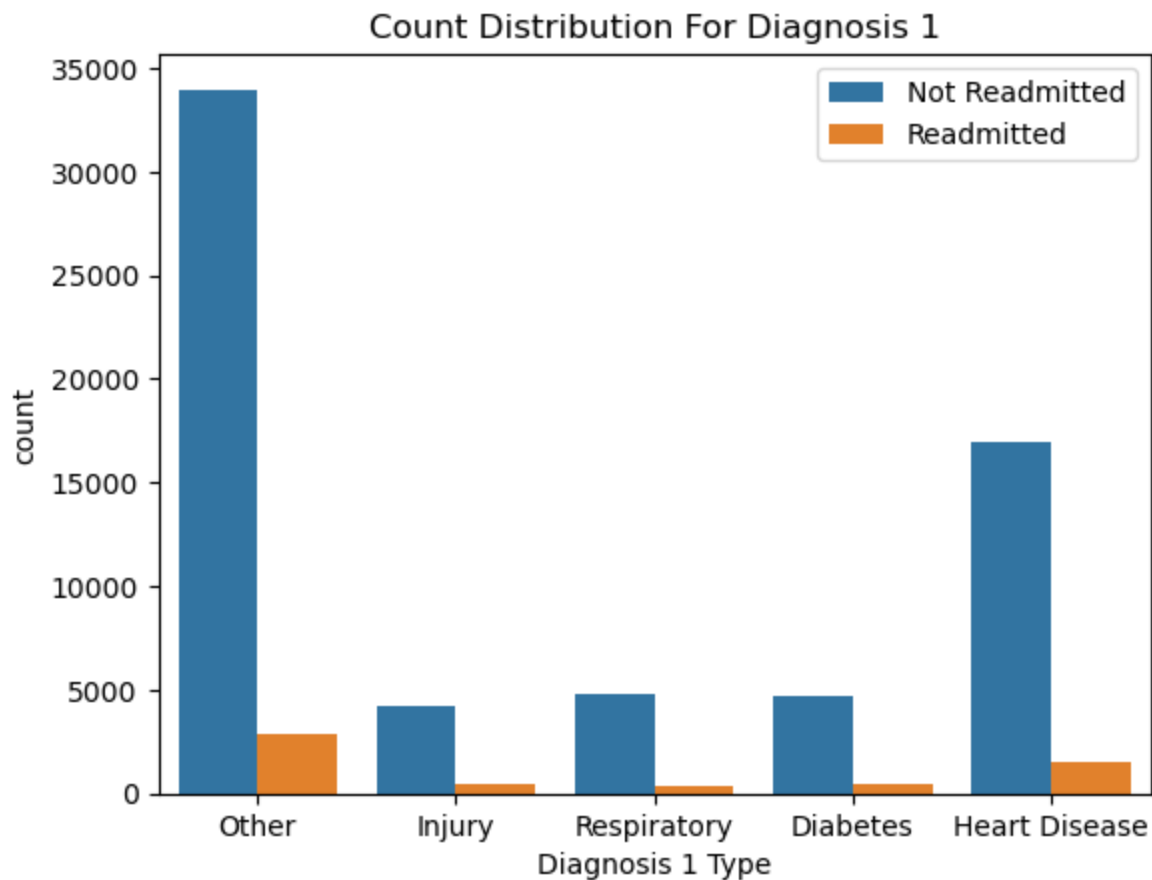
# Display the DataFrame with the new 'diag_1_new' feature
print(df[['diag_1', 'diag_1_new']])
```

	diag_1	diag_1_new
1	440	Other
2	997	Injury
3	486	Respiratory
4	250.03	Diabetes
5	414	Heart Disease
...

101752	998	Injury
101753	996	Injury
101754	V57	Other
101755	V57	Other
101762	V57	Other

[70256 rows x 2 columns]

```
In [29]: legend_label = ['Not Readmitted', 'Readmitted']
ax = sns.countplot(data = df, x = 'diag_1_new', hue = 'binary_readmitted')
ax.legend(legend_label)
plt.title('Count Distribution For Diagnosis 1')
plt.xlabel('Diagnosis 1 Type')
plt.show()
```



```
In [30]: # diag 2
custom_categories_diag_2 = {
    'Diabetes': ['250', '250.02', '250.03', '250.04', '250.1', '250.11', '250.12', '250.
    'Heart Disease': ['410', '411', '412', '413', '414', '415', '416', '417', '418', '41
    'Respiratory': ['460', '461', '462', '463', '464', '465', '466', '470', '471', '472'
    'Injury': ['800', '801', '802', '803', '804', '805', '806', '807', '808', '809', '81
    'Other': ['?', 'V', 'E']
}

# Create a new column 'diag_2_new' and set it to 'Other' initially
df['diag_2_new'] = 'Other'

# Map the values in 'diag_2' to custom categories
for category, codes in custom_categories_diag_2.items():
    df.loc[df['diag_2'].isin(codes), 'diag_2_new'] = category

# Display the DataFrame with the new 'diag_2_new' feature
print(df[['diag_2', 'diag_2_new']])
```

	diag_2	diag_2_new
1	413	Heart Disease
2	8	Other

```

3         250      Diabetes
4         401      Other
5         340      Other
...
101752     786      Other
101753    E878      Other
101754     788      Other
101755     348      Other
101762     599      Other

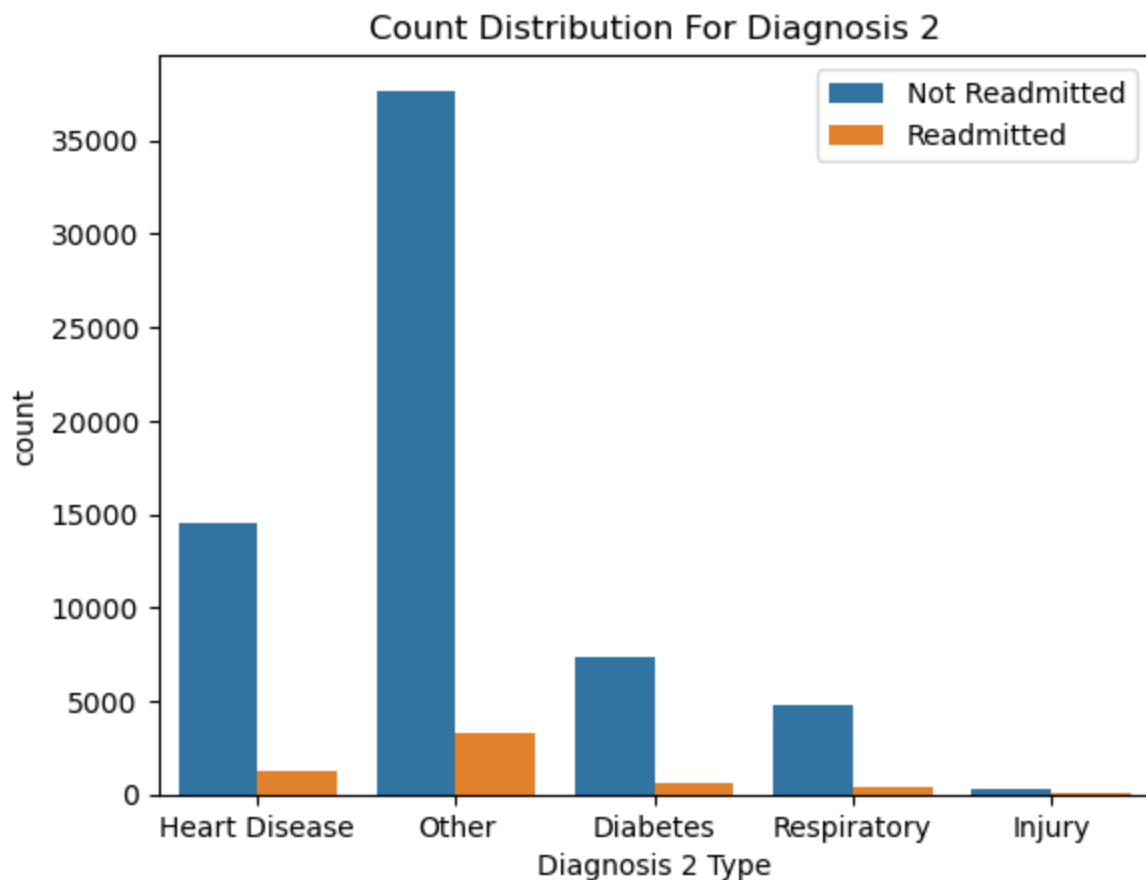
```

[70256 rows x 2 columns]

```

In [31]: legend_label = ['Not Readmitted', 'Readmitted']
ax = sns.countplot(data = df, x = 'diag_2_new', hue = 'binary_readmitted')
ax.legend(legend_label)
plt.title('Count Distribution For Diagnosis 2')
plt.xlabel('Diagnosis 2 Type')
plt.show()

```



```

In [32]: # diag 3
custom_categories_diag_3 = {
    'Diabetes': ['250', '250.02', '250.03', '250.04', '250.1', '250.11', '250.12', '250.
    'Heart Disease': ['410', '411', '412', '413', '414', '415', '416', '417', '418', '41
    'Respiratory': ['460', '461', '462', '463', '464', '465', '466', '470', '471', '472'
    'Injury': ['800', '801', '802', '803', '804', '805', '806', '807', '808', '809', '81
    'Other': ['NAN', 'E888', 'E878', 'V', 'V43', 'V70']
}

# Create a new column 'diag_3_new' and set it to 'Other' initially
df['diag_3_new'] = 'Other'

# Map the values in 'diag_3' to custom categories
for category, codes in custom_categories_diag_3.items():
    df.loc[df['diag_3'].isin(codes), 'diag_3_new'] = category

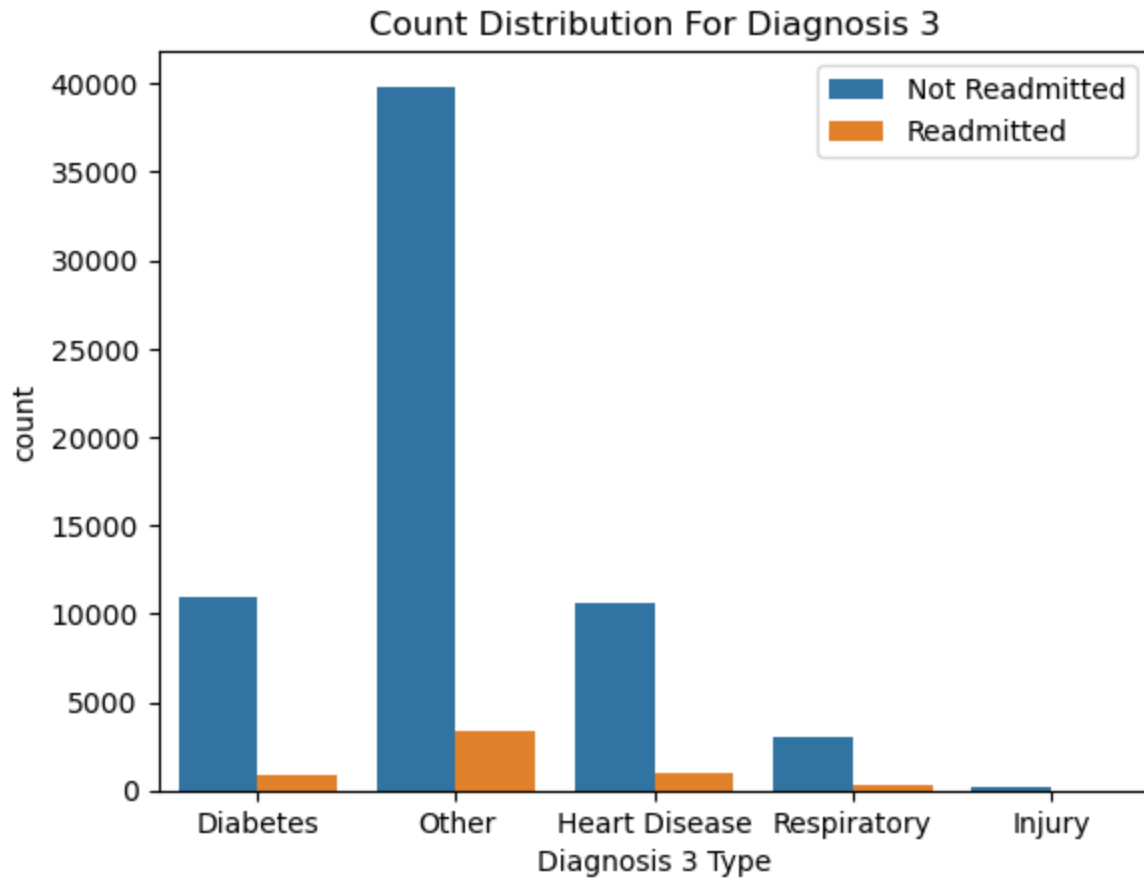
```

```
# Display the DataFrame with the new 'diag_3_new' feature
print(df[['diag_3', 'diag_3_new']])
```

	diag_3	diag_3_new
1	250.52	Diabetes
2	730	Other
3	427	Heart Disease
4	276	Other
5	401	Other
...
101752	584	Other
101753	401	Other
101754	349	Other
101755	V58	Other
101762	V43	Other

[70256 rows x 2 columns]

```
In [33]: legend_label = ['Not Readmitted', 'Readmitted']
ax = sns.countplot(data = df, x = 'diag_3_new', hue = 'binary_readmitted')
ax.legend(legend_label)
plt.title('Count Distribution For Diagnosis 3')
plt.xlabel('Diagnosis 3 Type')
plt.show()
```



Drop Columns no longer needed, or which pose redundant information

```
In [34]: df.drop(['encounter_id', 'patient_nbr', 'admission_type_id', 'discharge_disposition_id',
                  'admission_source_id', 'diag_1', 'diag_2', 'diag_3'], inplace = True, axis = 1)
```

Cross Tabulation Loop

```
In [35]: column_features = ['race', 'gender', 'age', 'time_in_hospital', 'payer_code',
    'medical_specialty', 'num_lab_procedures', 'num_procedures',
    'num_medications', 'number_outpatient', 'number_emergency',
    'number_inpatient', 'number_diagnoses', 'max_glu_serum', 'A1Cresult',
    'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide',
    'glimepiride', 'glipizide', 'glyburide', 'pioglitazone',
    'rosiglitazone', 'acarbose', 'miglitol', 'insulin',
    'glyburide-metformin', 'change', 'diabetesMed', 'discharge_desc',
    'admission_source_desc', 'admission_desc', 'binary_readmitted',
    'diag_1_new', 'diag_2_new', 'diag_3_new']

for feature in column_features:
    cross_tab = pd.crosstab(df['race'], df[feature])
    print(f"Cross-tabulation for race and {feature}:")
    print(cross_tab)
    print("\n")
```

Cross-tabulation for race and race:

race	AfricanAmerican	Asian	Caucasian	Hispanic	Other
AfricanAmerican	12625	0	0	0	0
Asian	0	489	0	0	0
Caucasian	0	0	54520	0	0
Hispanic	0	0	0	1467	0
Other	0	0	0	0	1155

Cross-tabulation for race and gender:

gender	Female	Male
AfricanAmerican	7685	4940
Asian	245	244
Caucasian	28086	26434
Hispanic	792	675
Other	583	572

Cross-tabulation for race and age:

age	[0-10)	[10-20)	[20-30)	[30-40)	[40-50)	[50-60)	[60-70)	\
AfricanAmerican	5	104	342	799	1813	2848	2820	
Asian	1	0	5	10	45	90	119	
Caucasian	54	216	601	1561	4444	8751	12201	
Hispanic	1	16	47	135	223	306	338	
Other	2	6	18	59	125	242	315	

age	[70-80)	[80-90)	[90-100)
AfricanAmerican	2481	1172	241
Asian	140	69	10
Caucasian	14916	10115	1661
Hispanic	280	110	11
Other	255	115	18

Cross-tabulation for race and time_in_hospital:

time_in_hospital	1	2	3	4	5	6	7	8	9	10	\
AfricanAmerican	1805	2225	2128	1692	1255	937	719	551	347	272	
Asian	103	102	78	62	39	33	16	22	10	6	
Caucasian	8127	9254	9697	7361	5233	3971	3087	2247	1543	1189	
Hispanic	230	305	266	181	138	104	63	55	41	29	
Other	211	199	207	152	86	61	59	57	32	31	

time_in_hospital	11	12	13	14
------------------	----	----	----	----

race				
AfricanAmerican	250	177	141	126
Asian	5	0	9	4
Caucasian	950	729	617	515
Hispanic	14	17	10	14
Other	20	16	14	10

Cross-tabulation for race and payer_code:

payer_code	BC	CH	CM	CP	DM	FR	HM	MC	MD	MP	\
race											
AfricanAmerican	616	22	333	265	124	0	598	2796	666	8	
Asian	14	1	7	5	3	0	25	83	31	0	
Caucasian	2653	93	975	1561	205	1	3272	17275	1266	33	
Hispanic	58	1	21	48	18	0	66	206	133	0	
Other	35	1	17	17	16	0	62	188	74	0	

payer_code	Not Specified	OG	OT	PO	SI	SP	UN	WC
race								
AfricanAmerican	6269	148	25	96	3	335	300	21
Asian	229	15	1	10	0	39	26	0
Caucasian	22249	382	39	324	29	2652	1423	88
Hispanic	742	41	1	9	1	85	32	5
Other	560	55	0	12	0	83	35	0

Cross-tabulation for race and medical_specialty:

medical_specialty	AllergyandImmunology	Anesthesiology	\
race			
AfricanAmerican		1	1
Asian		0	0
Caucasian		3	7
Hispanic		2	0
Other		0	1

medical_specialty	Anesthesiology-Pediatric	Cardiology	Cardiology-Pediatric	\
race				
AfricanAmerican		3	627	2
Asian		0	13	0
Caucasian		6	3604	3
Hispanic		0	48	0
Other		0	81	0

medical_specialty	DCPTEAM	Dentistry	Dermatology	Emergency/Trauma	\
race					
AfricanAmerican		0	2	0	412
Asian		0	0	0	87
Caucasian		4	2	1	3245
Hispanic		1	0	0	144
Other		0	0	0	161

medical_specialty	Endocrinology	Endocrinology-Metabolism	\
race			
AfricanAmerican		22	0
Asian		0	0
Caucasian		70	6
Hispanic		1	0
Other		0	0

medical_specialty	Family/GeneralPractice	Gastroenterology	Gynecology	\
race				
AfricanAmerican		1071	64	6
Asian		33	1	1
Caucasian		3699	369	38
Hispanic		175	4	2
Other		68	9	3

medical_specialty	Hematology	Hematology/Oncology	Hospitalist	\
race				
AfricanAmerican	3	15	2	
Asian	1	1	0	
Caucasian	39	114	38	
Hispanic	0	0	2	
Other	5	0	0	

medical_specialty	InfectiousDiseases	InternalMedicine	Nephrology	\
race				
AfricanAmerican	3	2837	332	
Asian	2	103	2	
Caucasian	24	7369	504	
Hispanic	1	252	8	
Other	0	141	7	

medical_specialty	Neurology	Neurophysiology	\
race			
AfricanAmerican	43	0	
Asian	1	0	
Caucasian	126	1	
Hispanic	3	0	
Other	6	0	

medical_specialty	Obsterics&Gynecology-GynecologicOnco	Obstetrics	\
race			
AfricanAmerican	6	3	
Asian	0	0	
Caucasian	12	12	
Hispanic	0	1	
Other	1	1	

medical_specialty	ObstetricsandGynecology	Oncology	Ophthalmology	\
race				
AfricanAmerican	156	61	6	
Asian	5	0	0	
Caucasian	396	169	26	
Hispanic	33	4	0	
Other	5	4	0	

medical_specialty	Orthopedics	Orthopedics-Reconstructive	Osteopath	Other	\
race					
AfricanAmerican	131	157	1	5382	
Asian	5	3	1	160	
Caucasian	980	836	35	26963	
Hispanic	17	17	0	633	
Other	8	9	0	530	

medical_specialty	Otolaryngology	OutreachServices	Pathology	Pediatrics	\
race					
AfricanAmerican	26	0	3	34	
Asian	1	0	0	1	
Caucasian	71	9	8	124	
Hispanic	4	0	0	2	
Other	2	0	1	3	

medical_specialty	Pediatrics-CriticalCare	Pediatrics-EmergencyMedicine	\
race			
AfricanAmerican	7	0	
Asian	1	0	
Caucasian	19	2	
Hispanic	0	0	
Other	2	0	

medical_specialty	Pediatrics-Endocrinology	Pediatrics-Hematology-Oncology	\
-------------------	--------------------------	--------------------------------	---

Asian	28	0	1
Caucasian	1777	9	307
Hispanic	43	0	5
Other	41	0	1

medical_specialty	Surgery-Pediatric	Surgery-Plastic	\
race			
AfricanAmerican	1	5	
Asian	0	0	
Caucasian	4	25	
Hispanic	0	0	
Other	0	1	

medical_specialty	Surgery-PlasticwithinHeadandNeck	Surgery-Thoracic	\
race			
AfricanAmerican		0	10
Asian		0	2
Caucasian		1	78
Hispanic		0	2
Other		0	5

medical_specialty	Surgery-Vascular	SurgicalSpecialty	Urology
race			
AfricanAmerican	144	6	106
Asian	0	0	1
Caucasian	246	18	411
Hispanic	4	1	15
Other	2	2	9

Cross-tabulation for race and num_lab_procedures:

num_lab_procedures	1	2	3	4	5	6	7	8	9	10	11	\
race												
AfricanAmerican	311	118	62	33	23	20	33	32	122	103	74	
Asian	27	3	3	2	1	0	3	3	8	6	0	
Caucasian	1787	618	415	241	178	169	213	202	537	451	374	
Hispanic	69	13	13	6	7	6	3	5	7	6	7	
Other	52	11	4	8	5	1	1	5	10	18	10	

num_lab_procedures	12	13	14	15	16	17	18	19	20	21	22	\
race												
AfricanAmerican	43	42	36	50	55	54	62	59	80	85	65	
Asian	1	1	6	0	1	3	6	5	5	6	8	
Caucasian	315	237	204	238	314	399	409	601	450	413	402	
Hispanic	4	8	6	5	8	8	6	14	22	15	7	
Other	5	8	7	3	4	2	6	18	14	11	13	

num_lab_procedures	23	24	25	26	27	28	29	30	31	32	33	\
race												
AfricanAmerican	70	65	58	143	105	107	196	177	166	147	178	
Asian	7	8	4	7	9	4	8	5	8	6	5	
Caucasian	491	436	614	620	468	598	739	727	731	745	712	
Hispanic	19	14	18	9	9	21	11	16	18	17	20	
Other	7	13	12	12	11	20	10	20	21	17	11	

num_lab_procedures	34	35	36	37	38	39	40	41	42	43	\
race											
AfricanAmerican	237	294	278	287	372	322	328	291	272	496	
Asian	9	15	11	8	6	6	11	8	3	15	
Caucasian	855	944	996	1091	1130	1079	1108	1073	1077	1295	
Hispanic	27	27	20	34	27	27	56	29	27	32	
Other	20	15	17	27	15	19	24	18	22	20	

num_lab_procedures	44	45	46	47	48	49	50	51	52	53	\
race											
AfricanAmerican	365	373	313	300	282	256	225	248	245	230	

Asian	14	9	11	12	7	8	7	10	5	11
Caucasian	1227	1175	1125	1077	1066	1036	1011	1008	940	941
Hispanic	27	20	23	28	33	35	26	19	23	26
Other	22	21	31	21	15	28	24	15	19	14

num_lab_procedures	54	55	56	57	58	59	60	61	62	63	64	\
race												
AfricanAmerican	208	229	195	197	189	184	190	205	184	167	128	
Asian	12	6	7	4	4	11	6	6	3	4	8	
Caucasian	1005	966	1001	941	891	877	821	855	783	753	701	
Hispanic	27	33	25	21	21	23	30	21	15	29	15	
Other	8	22	27	20	12	16	18	13	16	19	20	

num_lab_procedures	65	66	67	68	69	70	71	72	73	74	75	\
race												
AfricanAmerican	164	145	139	93	103	83	85	83	77	75	57	
Asian	7	9	4	4	4	1	4	4	4	3	2	
Caucasian	668	657	607	616	565	485	459	417	368	340	300	
Hispanic	22	19	19	18	13	15	11	12	6	11	10	
Other	12	6	9	11	8	17	14	13	11	5	3	

num_lab_procedures	76	77	78	79	80	81	82	83	84	85	86	\
race												
AfricanAmerican	40	57	36	37	37	30	27	29	15	13	13	
Asian	3	1	1	1	1	1	0	1	2	0	1	
Caucasian	276	231	209	178	158	162	115	134	92	93	70	
Hispanic	8	12	3	3	6	3	4	3	1	2	4	
Other	7	8	7	7	8	8	3	2	1	5	4	

num_lab_procedures	87	88	89	90	91	92	93	94	95	96	97	\
race												
AfricanAmerican	17	7	11	10	12	9	10	11	5	3	6	
Asian	1	0	0	0	0	0	0	0	1	0	2	
Caucasian	50	61	40	34	33	22	31	27	33	17	16	
Hispanic	1	3	3	2	1	1	1	1	2	1	0	
Other	5	2	2	1	2	0	1	1	0	1	1	

num_lab_procedures	98	99	100	101	102	103	104	105	106	107	108	\
race												
AfricanAmerican	4	2	3	2	2	1	1	0	0	1	1	
Asian	0	0	0	0	0	0	0	0	0	0	0	
Caucasian	15	5	7	7	4	4	0	6	5	0	2	
Hispanic	1	0	0	0	0	1	0	0	0	0	0	
Other	0	0	1	0	0	0	0	0	0	0	1	

num_lab_procedures	109	111	113	114	118	120	121	132				
race												
AfricanAmerican	0	1	2	0	0	0	1	1				
Asian	1	0	0	0	0	0	0	0				
Caucasian	1	0	0	2	1	1	0	0				
Hispanic	0	1	0	0	0	0	0	0				
Other	0	0	0	0	0	0	0	0				

Cross-tabulation for race and num_procedures:

num_procedures	0	1	2	3	4	5	6					
race												
AfricanAmerican	5725	2628	1656	1288	505	416	407					
Asian	211	98	75	60	20	7	18					
Caucasian	23380	10847	7130	5454	2476	1897	3336					
Hispanic	712	297	175	152	54	28	49					
Other	470	216	152	128	52	47	90					

Cross-tabulation for race and num_medications:

num_medications	1	2	3	4	5	6	7	8	9	10	11	\
-----------------	---	---	---	---	---	---	---	---	---	----	----	---

race						
AfricanAmerican	0	0	0	0	0	1
Asian	0	0	0	0	0	0
Caucasian	1	1	1	1	1	0
Hispanic	0	0	0	0	0	0
Other	0	0	0	0	0	0

Cross-tabulation for race and number_emergency:

number_emergency	0	1	2	3	4	5	6	7	8	9	10	11	14	\
race														
AfricanAmerican	11377	899	206	77	32	17	7	4	3	1	1	0	0	
Asian	462	21	5	0	0	0	0	1	0	0	0	0	0	
Caucasian	50442	3095	648	172	78	29	22	9	7	4	3	3	1	
Hispanic	1343	89	18	8	3	3	0	1	0	1	1	0	0	
Other	1063	72	13	5	2	0	0	0	0	0	0	0	0	

number_emergency	15	16	19	20	25	37	42	76
race								
AfricanAmerican	0	1	0	0	0	0	0	0
Asian	0	0	0	0	0	0	0	0
Caucasian	1	0	1	1	1	1	1	1
Hispanic	0	0	0	0	0	0	0	0
Other	0	0	0	0	0	0	0	0

Cross-tabulation for race and number_inpatient:

number_inpatient	0	1	2	3	4	5	6	7	8	9	10	11	\
race													
AfricanAmerican	10537	1418	425	124	61	28	18	6	1	3	3	0	
Asian	406	62	15	4	2	0	0	0	0	0	0	0	
Caucasian	44629	6792	1894	664	288	125	71	19	17	7	7	4	
Hispanic	1187	182	61	15	9	6	3	3	1	0	0	0	
Other	993	121	23	13	3	2	0	0	0	0	0	0	

number_inpatient	12	13	15
race			
AfricanAmerican	0	1	0
Asian	0	0	0
Caucasian	2	0	1
Hispanic	0	0	0
Other	0	0	0

Cross-tabulation for race and number_diagnoses:

number_diagnoses	3	4	5	6	7	8	9	10	11	12	13	\
race												
AfricanAmerican	561	1043	2013	1459	1394	1208	4942	1	0	0	1	
Asian	19	53	62	66	42	38	209	0	0	0	0	
Caucasian	1579	3057	6400	5759	5855	5946	25863	5	4	5	8	
Hispanic	85	147	180	242	139	146	524	1	0	0	1	
Other	69	79	136	145	133	121	472	0	0	0	0	

number_diagnoses	14	15	16
race			
AfricanAmerican	0	0	3
Asian	0	0	0
Caucasian	4	6	29
Hispanic	0	0	2
Other	0	0	0

Cross-tabulation for race and max_glu_serum:

max_glu_serum	>200	>300	None	Norm
race				
AfricanAmerican	87	70	12323	145

Asian	7	8	467	7
Caucasian	803	598	51621	1498
Hispanic	41	30	1346	50
Other	15	11	1104	25

Cross-tabulation for race and AlCresult:				
AlCresult	>7	>8	None	Norm
race				
AfricanAmerican	394	1259	10186	786
Asian	28	49	387	25
Caucasian	2267	4199	45314	2740
Hispanic	71	195	1102	99
Other	59	114	910	72

Cross-tabulation for race and metformin:				
metformin	Down	No	Steady	Up
race				
AfricanAmerican	68	10134	2271	152
Asian	4	386	96	3
Caucasian	338	43110	10438	634
Hispanic	14	1099	332	22
Other	6	911	227	11

Cross-tabulation for race and repaglinide:				
repaglinide	Down	No	Steady	Up
race				
AfricanAmerican	3	12513	102	7
Asian	1	476	12	0
Caucasian	25	53759	682	54
Hispanic	0	1438	24	5
Other	0	1134	19	2

Cross-tabulation for race and nateglinide:				
nateglinide	Down	No	Steady	Up
race				
AfricanAmerican	1	12513	109	2
Asian	1	479	9	0
Caucasian	6	54170	331	13
Hispanic	1	1456	10	0
Other	0	1142	13	0

Cross-tabulation for race and chlorpropamide:				
chlorpropamide	Down	No	Steady	Up
race				
AfricanAmerican	0	12620	4	1
Asian	0	489	0	0
Caucasian	1	54460	56	3
Hispanic	0	1466	1	0
Other	0	1154	1	0

Cross-tabulation for race and glimepiride:				
glimepiride	Down	No	Steady	Up
race				
AfricanAmerican	25	12033	527	40
Asian	0	468	20	1
Caucasian	107	51545	2679	189
Hispanic	0	1410	54	3
Other	4	1090	58	3

Cross-tabulation for race and glipizide:				
glipizide	Down	No	Steady	Up
race				
AfricanAmerican	59	11153	1317	96
Asian	0	419	63	7
Caucasian	296	47491	6307	426
Hispanic	8	1279	160	20
Other	4	982	155	14

Cross-tabulation for race and glyburide:				
glyburide	Down	No	Steady	Up
race				
AfricanAmerican	40	11509	979	97
Asian	2	438	47	2
Caucasian	349	48287	5397	487
Hispanic	7	1329	119	12
Other	2	1042	105	6

Cross-tabulation for race and pioglitazone:				
pioglitazone	Down	No	Steady	Up
race				
AfricanAmerican	11	11847	735	32
Asian	1	451	36	1
Caucasian	67	50306	4005	142
Hispanic	3	1348	111	5
Other	1	1089	63	2

Cross-tabulation for race and rosiglitazone:				
rosiglitazone	Down	No	Steady	Up
race				
AfricanAmerican	9	11875	714	27
Asian	0	466	22	1
Caucasian	65	50839	3517	99
Hispanic	1	1379	84	3
Other	0	1095	55	5

Cross-tabulation for race and acarbose:				
acarbose	Down	No	Steady	Up
race				
AfricanAmerican	0	12607	17	1
Asian	0	487	2	0
Caucasian	1	54352	159	8
Hispanic	0	1465	2	0
Other	0	1153	2	0

Cross-tabulation for race and miglitol:				
miglitol	Down	No	Steady	
race				
AfricanAmerican	0	12622	3	
Asian	0	489	0	
Caucasian	1	54504	15	
Hispanic	0	1467	0	
Other	0	1155	0	

Cross-tabulation for race and insulin:				
insulin	Down	No	Steady	Up
race				
AfricanAmerican	1462	5528	4288	1347
Asian	55	276	113	45
Caucasian	5512	27365	16462	5181

Hispanic	178	695	441	153
Other	187	514	293	161

Cross-tabulation for race and glyburide-metformin:

glyburide-metformin	Down	No	Steady	Up
race				
AfricanAmerican	1	12551	73	0
Asian	0	480	9	0
Caucasian	3	54149	361	7
Hispanic	0	1458	9	0
Other	0	1139	16	0

Cross-tabulation for race and change:

change	Ch	No
race		
AfricanAmerican	5646	6979
Asian	201	288
Caucasian	24298	30222
Hispanic	698	769
Other	577	578

Cross-tabulation for race and diabetesMed:

diabetesMed	No	Yes
race		
AfricanAmerican	3021	9604
Asian	132	357
Caucasian	13260	41260
Hispanic	371	1096
Other	253	902

Cross-tabulation for race and discharge_desc:

discharge_desc	Elective	Emergency	Newborn	Not Available	Not Mapped \
race					
AfricanAmerican	2195	7484	2	244	55
Asian	89	218	1	10	0
Caucasian	12088	25158	3	2561	228
Hispanic	255	722	2	66	2
Other	216	531	0	32	5

Cross-tabulation for race and discharge_desc:

discharge_desc	Other	Trauma Center	Urgent
race			
AfricanAmerican	538	2	2105
Asian	49	0	122
Caucasian	3612	18	10852
Hispanic	202	0	218
Other	99	0	272

Cross-tabulation for race and admission_source_desc:

admission_source_desc	Admitted as an inpatient to this hospital \
race	
AfricanAmerican	0
Asian	0
Caucasian	9
Hispanic	0
Other	0

Cross-tabulation for race and admission_source_desc:

admission_source_desc	Discharged to home	Discharged/transferred to ICF \
race		
AfricanAmerican	8637	108
Asian	359	1
Caucasian	33747	361

Hispanic	1098	0
Other	795	7
admission_source_desc Discharged/transferred to SNF \		
race		
AfricanAmerican	1129	
Asian	45	
Caucasian	7579	
Hispanic	109	
Other	94	
admission_source_desc Discharged/transferred to a federal health care facility. \		
race		
AfricanAmerican		2
Asian		0
Caucasian		1
Hispanic		0
Other		0
admission_source_desc Discharged/transferred to a long term care hospital. \		
race		
AfricanAmerican		76
Asian		0
Caucasian		156
Hispanic		1
Other		4
admission_source_desc Discharged/transferred to a nursing facility certified under Medi		
caid but not certified under Medicare. \		
race		
AfricanAmerican		5
Asian		0
Caucasian		19
Hispanic		0
Other		0
admission_source_desc Discharged/transferred to another rehab fac including rehab units		
of a hospital . \		
race		
AfricanAmerican		236
Asian		7
Caucasian		860
Hispanic		11
Other		19
admission_source_desc Discharged/transferred to another short term hospital \		
race		
AfricanAmerican		122
Asian		17
Caucasian		1098
Hispanic		24
Other		27
admission_source_desc Discharged/transferred to another type of inpatient care institut		

ion \	
race	
AfricanAmerican	140
Asian	6
Caucasian	618
Hispanic	23
Other	10
admission_source_desc	Discharged/transferred to home under care of Home IV provider \
race	
AfricanAmerican	7
Asian	0
Caucasian	50
Hispanic	0
Other	2
admission_source_desc	Discharged/transferred to home with home health service \
race	
AfricanAmerican	1178
Asian	35
Caucasian	6079
Hispanic	103
Other	107
admission_source_desc	Discharged/transferred within this institution to Medicare approv
ed swing bed \	
race	
AfricanAmerican	3
Asian	0
Caucasian	22
Hispanic	0
Other	0
admission_source_desc	Discharged/transferred/referred another institution for outpatien
t services \	
race	
AfricanAmerican	1
Asian	0
Caucasian	1
Hispanic	0
Other	0
admission_source_desc	Discharged/transferred/referred to a psychiatric hospital of psyc
hiatric distinct part unit of a hospital \	
race	
AfricanAmerican	16

Asian		0
Caucasian		44
Hispanic		3
Other		0

admission_source_desc	Discharged/transferred/referred to this institution for outpatient services \
race	
AfricanAmerican	2
Asian	0
Caucasian	4
Hispanic	0
Other	0

admission_source_desc	Expired	Expired at home. Medicaid only, hospice. \
race		
AfricanAmerican	211	0
Asian	6	1
Caucasian	892	5
Hispanic	12	0
Other	19	0

admission_source_desc	Expired in a medical facility. Medicaid only, hospice. \
race	
AfricanAmerican	0
Asian	0
Caucasian	1
Hispanic	0
Other	0

admission_source_desc	Hospice / home	Hospice / medical facility	Left AMA \
race			
AfricanAmerican	47	25	69
Asian	2	0	0
Caucasian	189	200	231
Hispanic	4	2	14
Other	7	4	13

admission_source_desc	Neonate discharged to another hospital for neonatal aftercare \
race	
AfricanAmerican	0
Asian	0
Caucasian	5
Hispanic	1
Other	0

admission_source_desc	Not Mapped	Other \
race		
AfricanAmerican	140	471
Asian	3	7
Caucasian	655	1692
Hispanic	5	57
Other	9	38

admission_source_desc	Still patient or expected to return for outpatient services
race	

AfricanAmerican	0
Asian	0
Caucasian	2
Hispanic	0
Other	0

Cross-tabulation for race and admission_desc:

admission_desc race	Court/Law Enforcement	Emergency Room	Extramural Birth \
AfricanAmerican	3	7191	2
Asian	0	305	0
Caucasian	4	26273	0
Hispanic	2	904	0
Other	0	677	0

admission_desc race	Not Available	Not Mapped	Physician Referral	Sick Baby \
AfricanAmerican	12	3	3874	0
Asian	0	1	153	0
Caucasian	63	132	19780	1
Hispanic	0	0	433	0
Other	0	0	337	0

admission_desc race	Transfer from Ambulatory Surgery Center \
AfricanAmerican	0
Asian	0
Caucasian	2
Hispanic	0
Other	0

admission_desc race	Transfer from a Skilled Nursing Facility (SNF) \
AfricanAmerican	119
Asian	3
Caucasian	355
Hispanic	3
Other	0

admission_desc race	Transfer from another health care facility \
AfricanAmerican	454
Asian	3
Caucasian	1091
Hispanic	6
Other	3

admission_desc race	Transfer from critical access hospital \
AfricanAmerican	0
Asian	0
Caucasian	6
Hispanic	0
Other	0

admission_desc race	Transfer from hospital inpt/same fac reslt in a sep claim \
AfricanAmerican	1
Asian	0
Caucasian	2
Hispanic	0
Other	0

admission_desc race	Clinic Referral	HMO Referral	Normal Delivery	Other \
------------------------	-----------------	--------------	-----------------	---------

AfricanAmerican	231	8	0	424
Asian	1	0	0	15
Caucasian	558	81	1	4203
Hispanic	6	0	0	94
Other	23	0	0	66

admission_desc	Transfer from a hospital
race	
AfricanAmerican	303
Asian	8
Caucasian	1968
Hispanic	19
Other	49

Cross-tabulation for race and binary_readmitted:

binary_readmitted	0	1
race		
AfricanAmerican	11652	973
Asian	456	33
Caucasian	50149	4371
Hispanic	1356	111
Other	1086	69

Cross-tabulation for race and diag_1_new:

diag_1_new	Diabetes	Heart Disease	Injury	Other	Respiratory
race					
AfricanAmerican	1498	2789	779	6787	772
Asian	19	124	31	288	27
Caucasian	3465	14913	3690	28288	4164
Hispanic	135	300	82	838	112
Other	83	311	88	589	84

Cross-tabulation for race and diag_2_new:

diag_2_new	Diabetes	Heart Disease	Injury	Other	Respiratory
race					
AfricanAmerican	1648	2199	42	8098	638
Asian	67	93	5	300	24
Caucasian	5894	13094	293	30883	4356
Hispanic	217	193	11	955	91
Other	145	228	5	701	76

Cross-tabulation for race and diag_3_new:

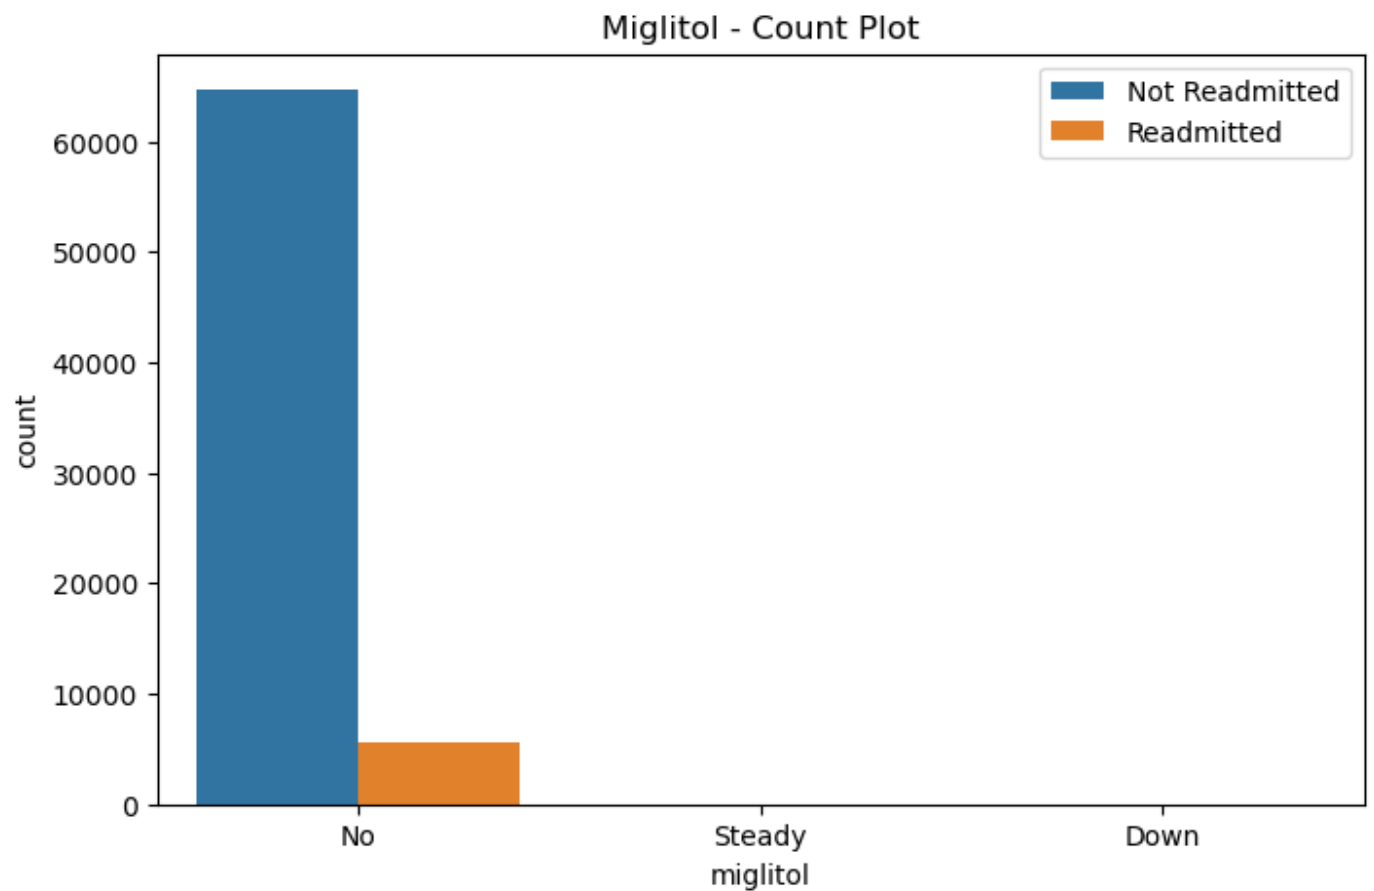
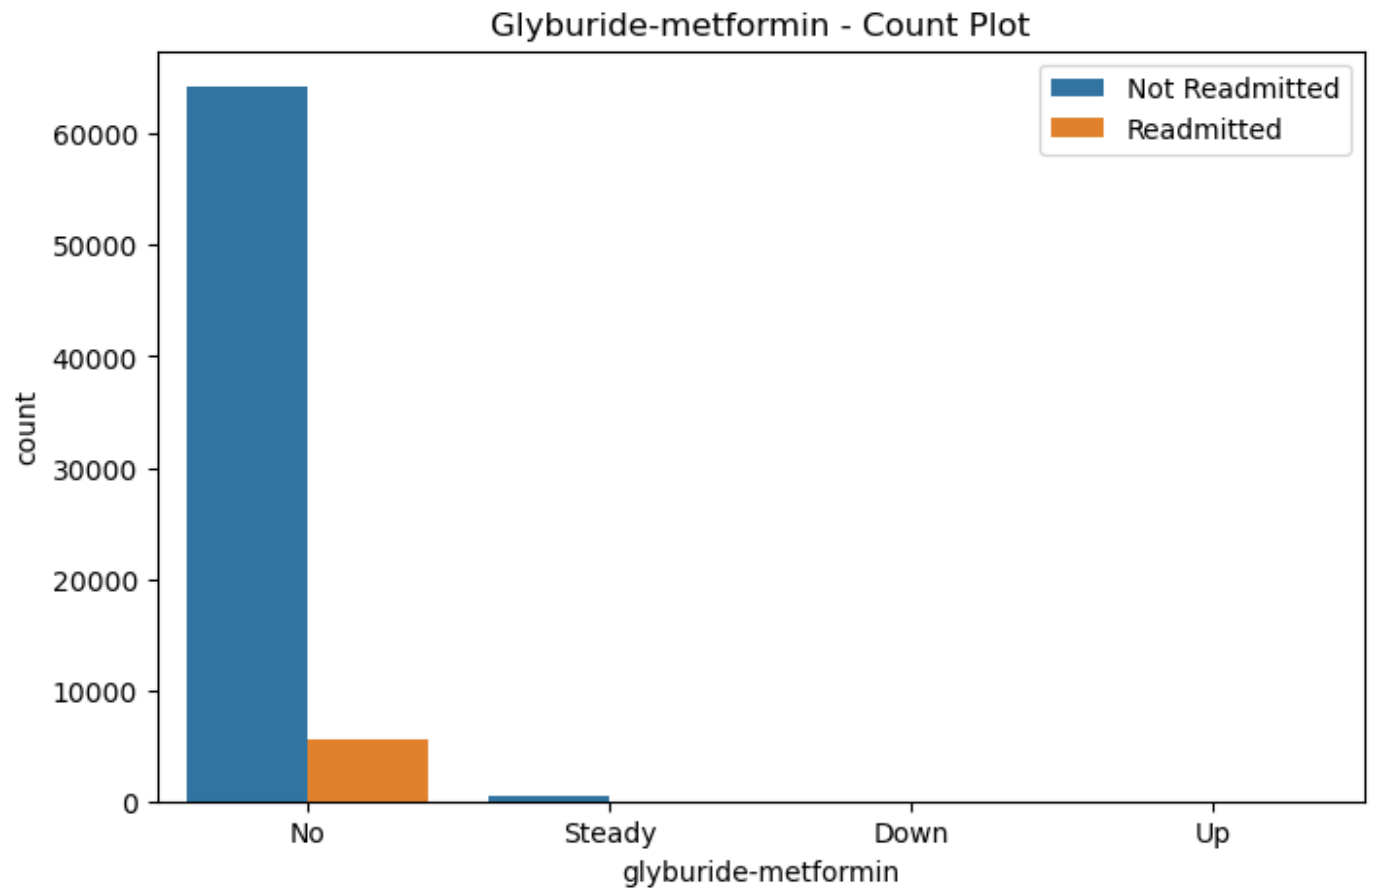
diag_3_new	Diabetes	Heart Disease	Injury	Other	Respiratory
race					
AfricanAmerican	2176	1618	24	8353	454
Asian	85	67	1	326	10
Caucasian	9102	9586	173	32878	2781
Hispanic	302	168	8	936	53
Other	219	143	5	746	42

```
In [36]: honed_cat_var = df[['glyburide-metformin','miglitol','acarbose',
                             'rosiglitazone','pioglitazone','glyburide','chlorpropamide',
                             'nateglinide','repaglinide', 'binary_readmitted']]
```

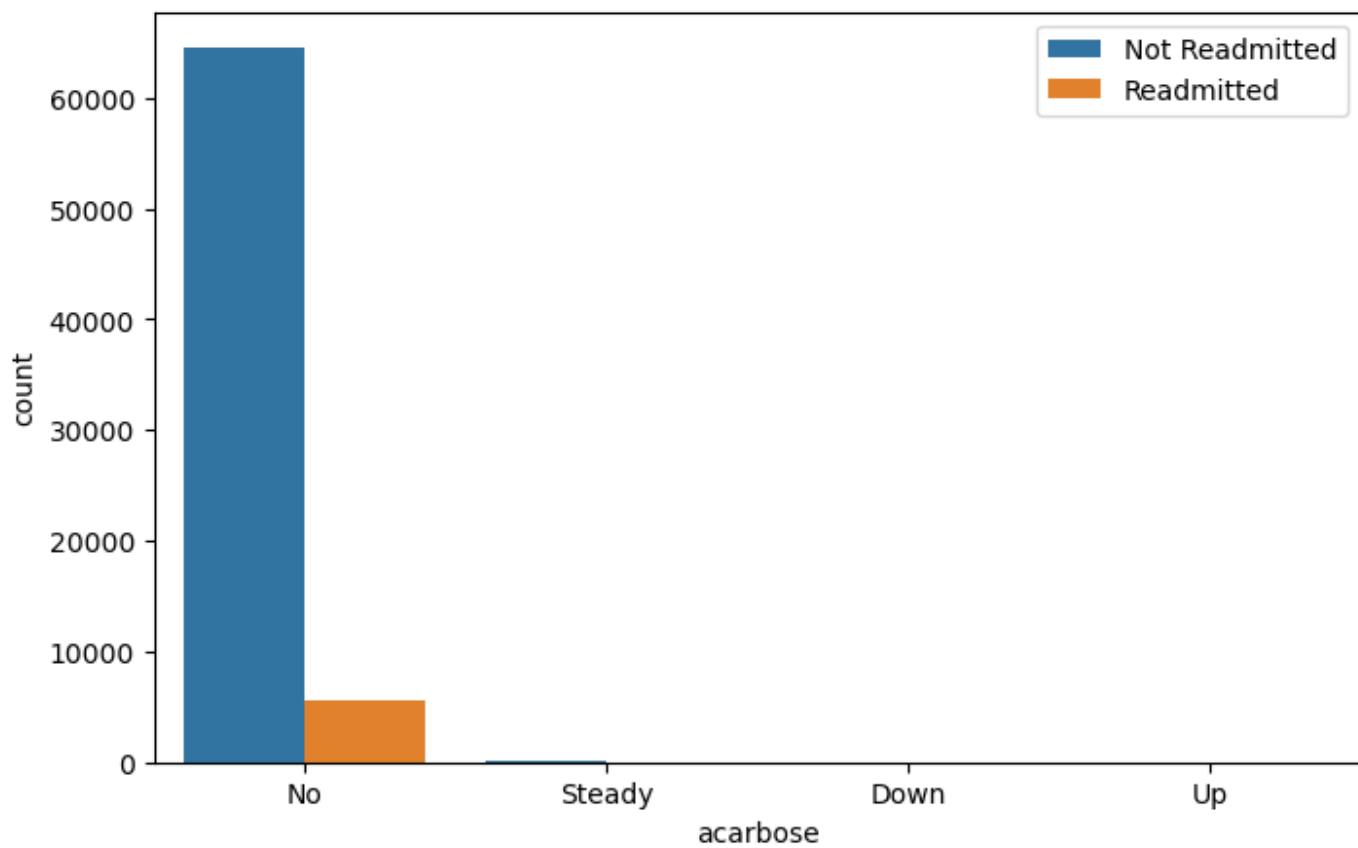
```
In [37]: legend_label = ['Not Readmitted','Readmitted']

for var in honed_cat_var:
    if var != 'binary_readmitted':
        plt.figure(figsize=(8,5))
```

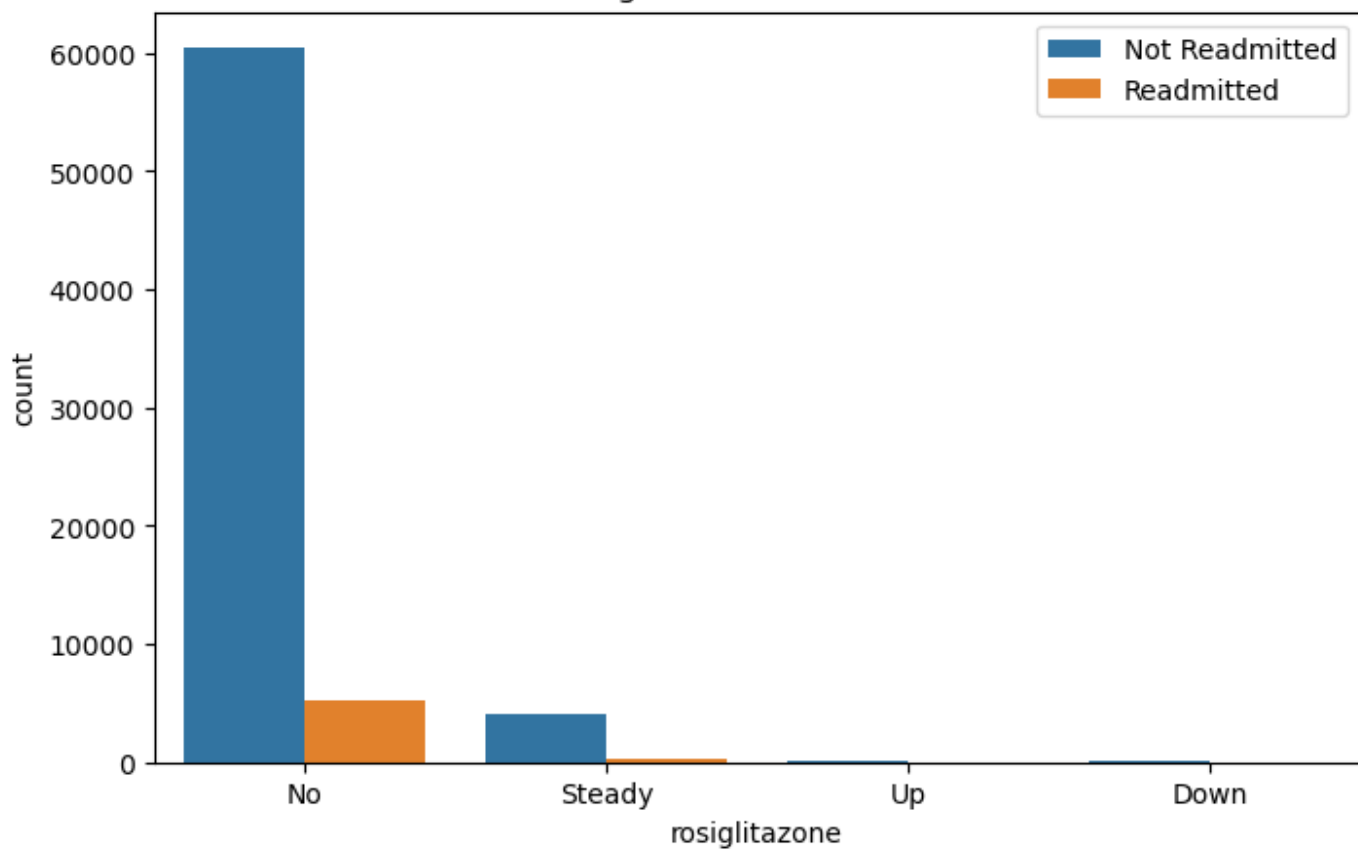
```
ax = sns.countplot(data = df, x = var, hue = 'binary_readmitted')
ax.legend(legend_label)
plt.title(f'{var.capitalize()} - Count Plot')
plt.show()
```



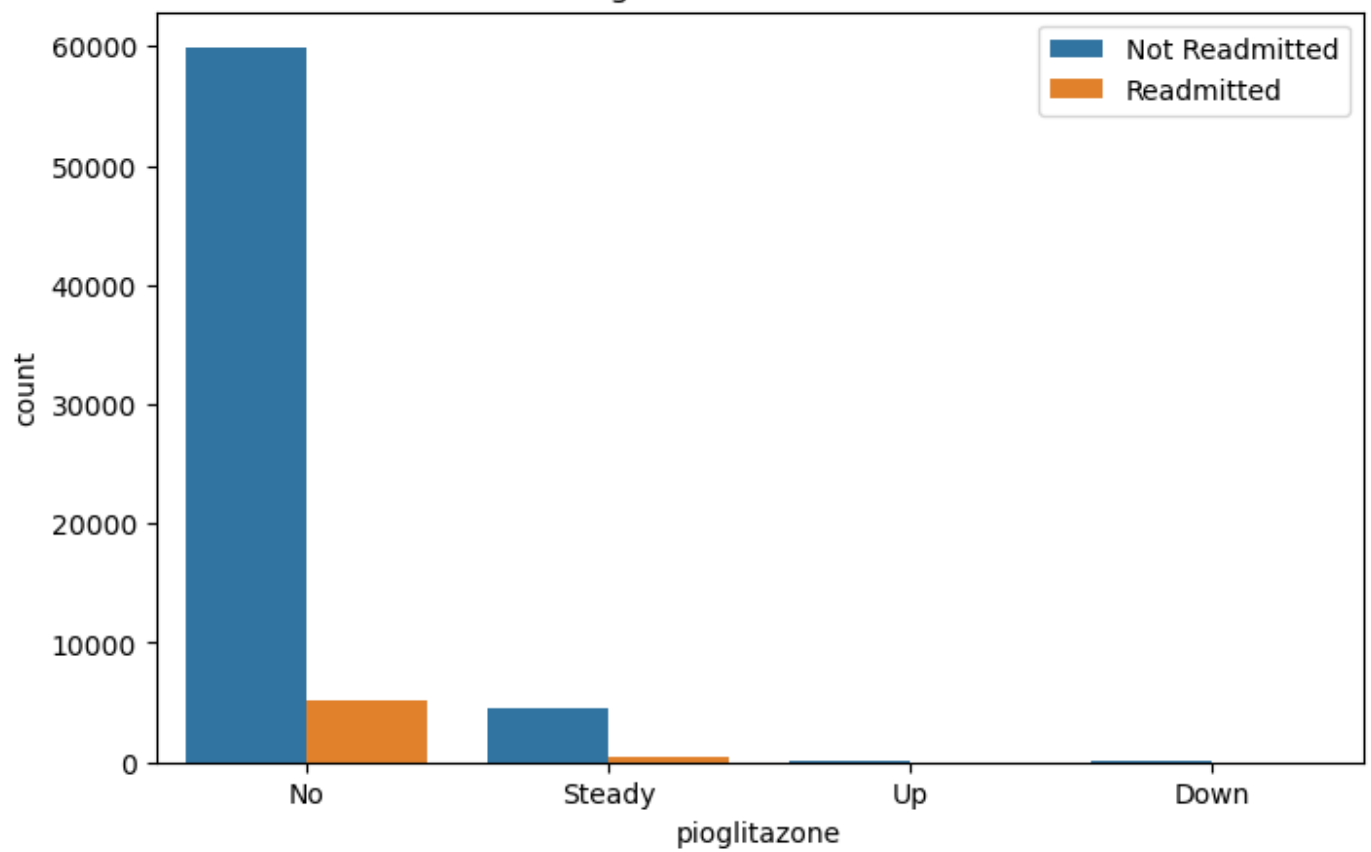
Acarbose - Count Plot



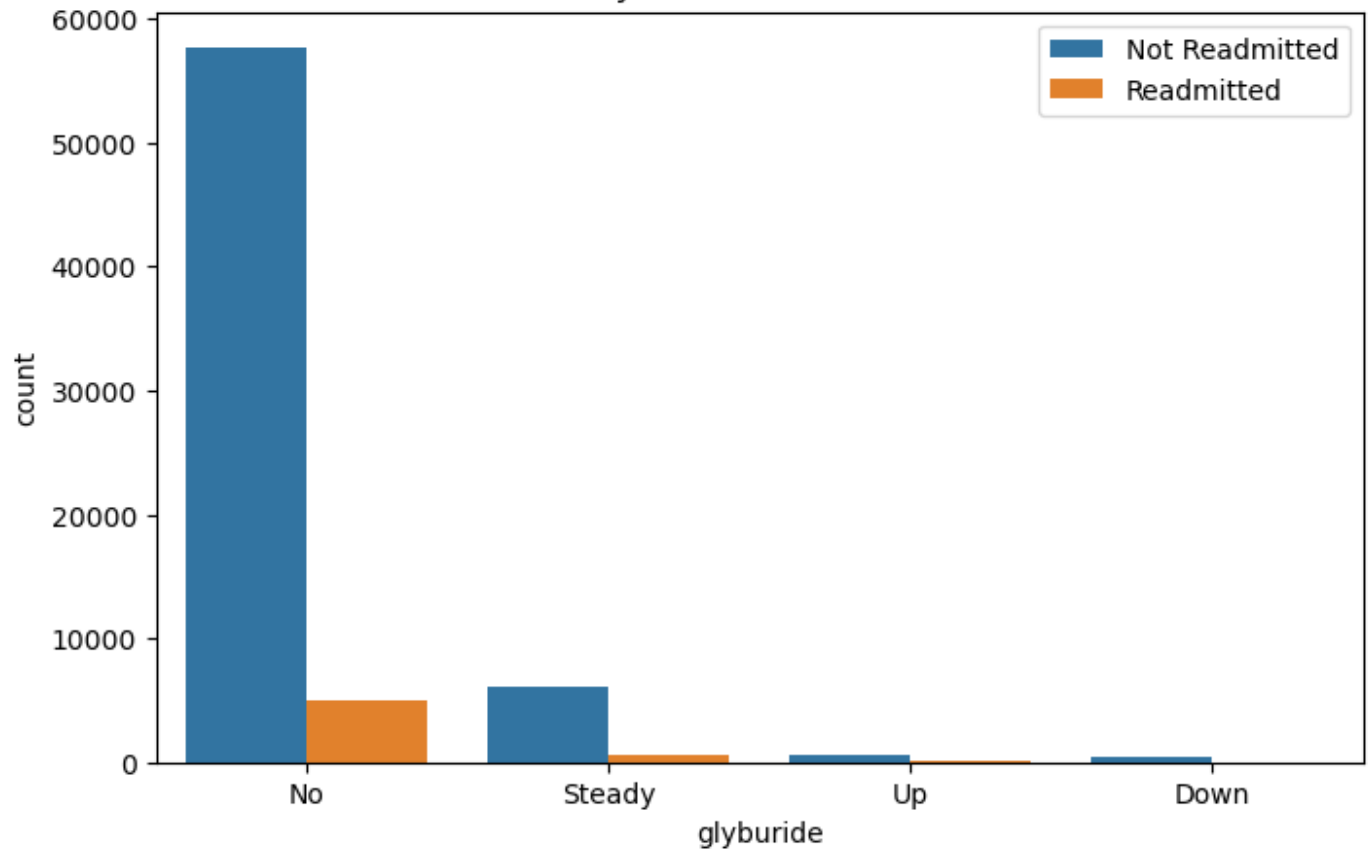
Rosiglitazone - Count Plot



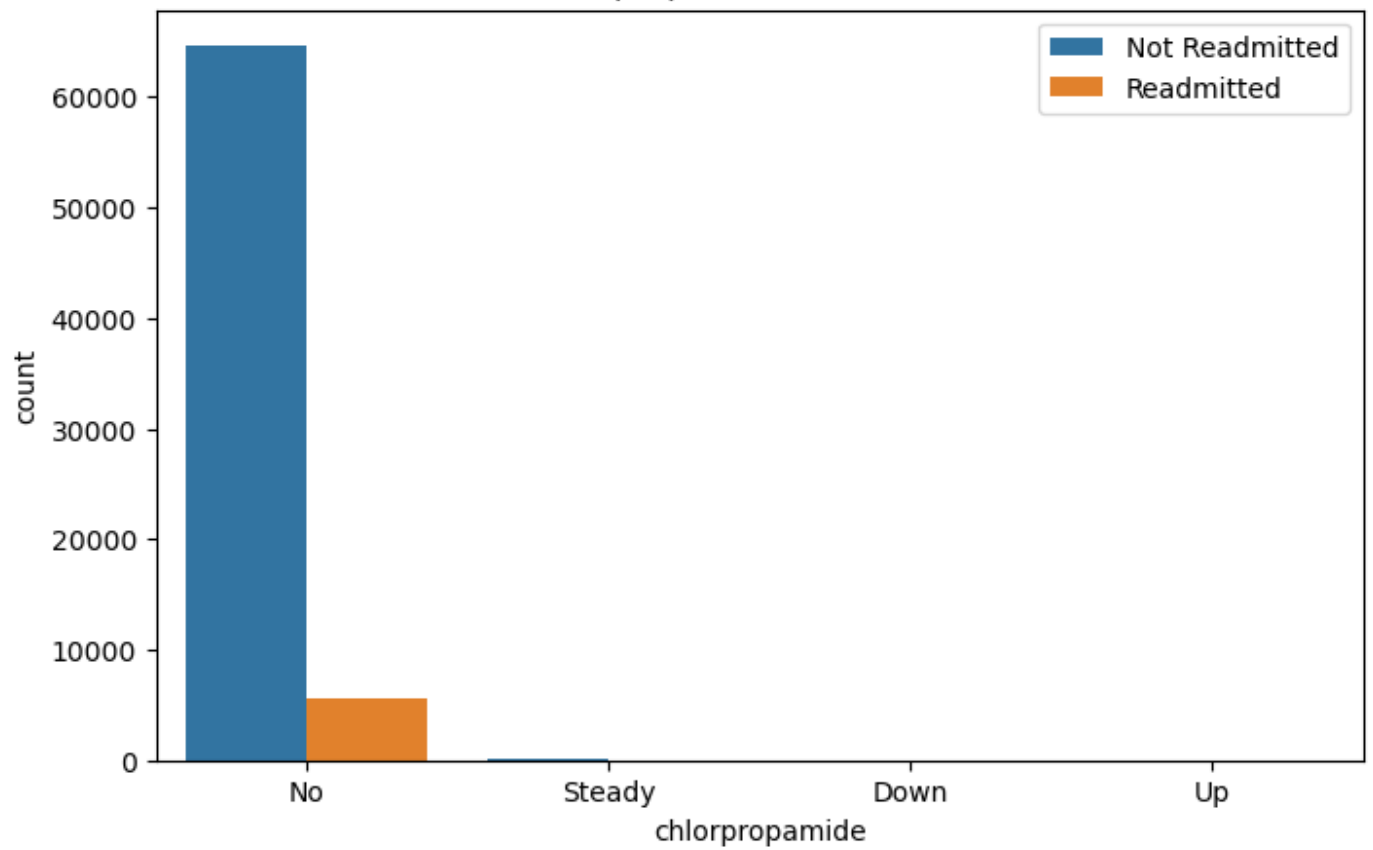
Pioglitazone - Count Plot



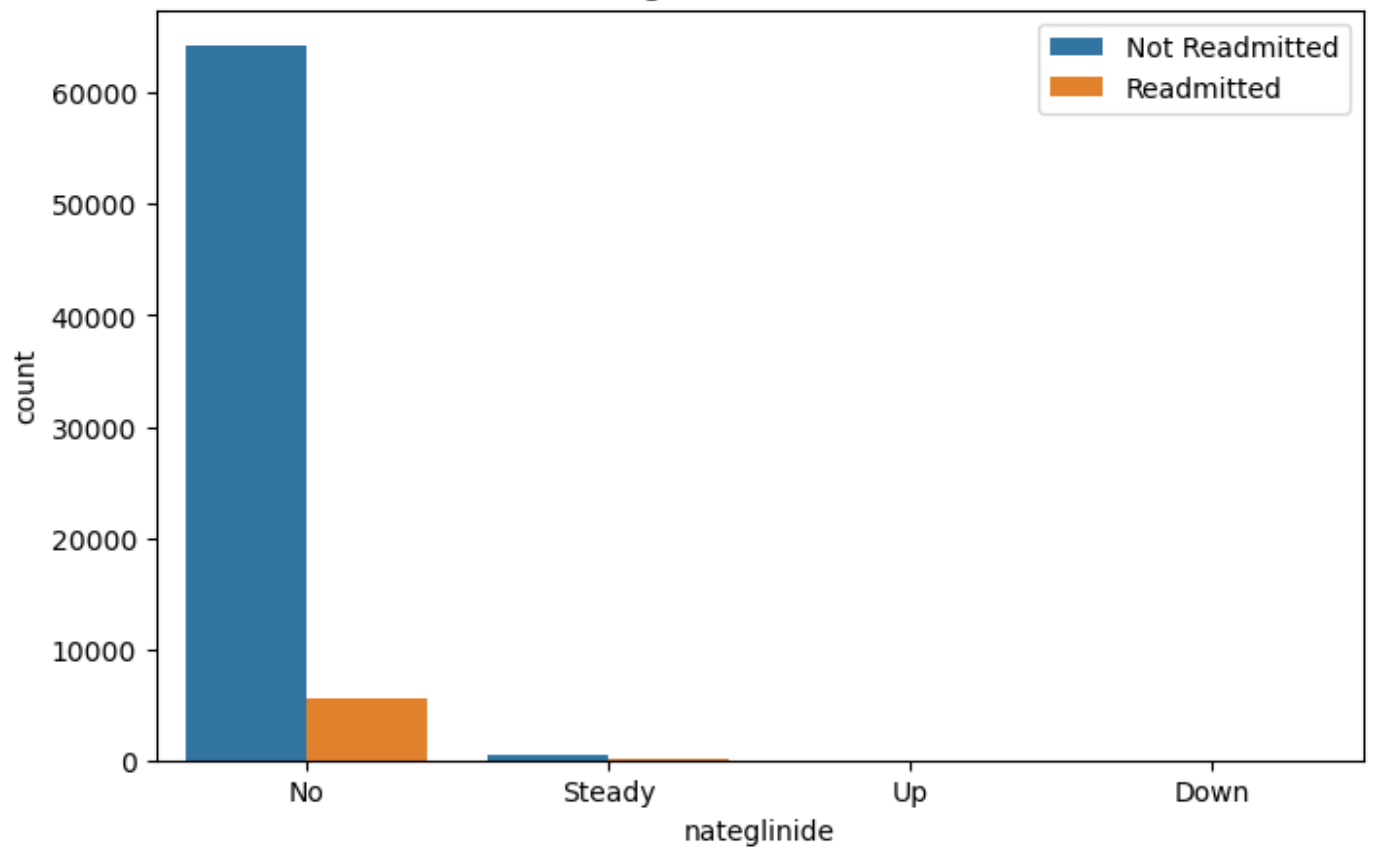
Glyburide - Count Plot

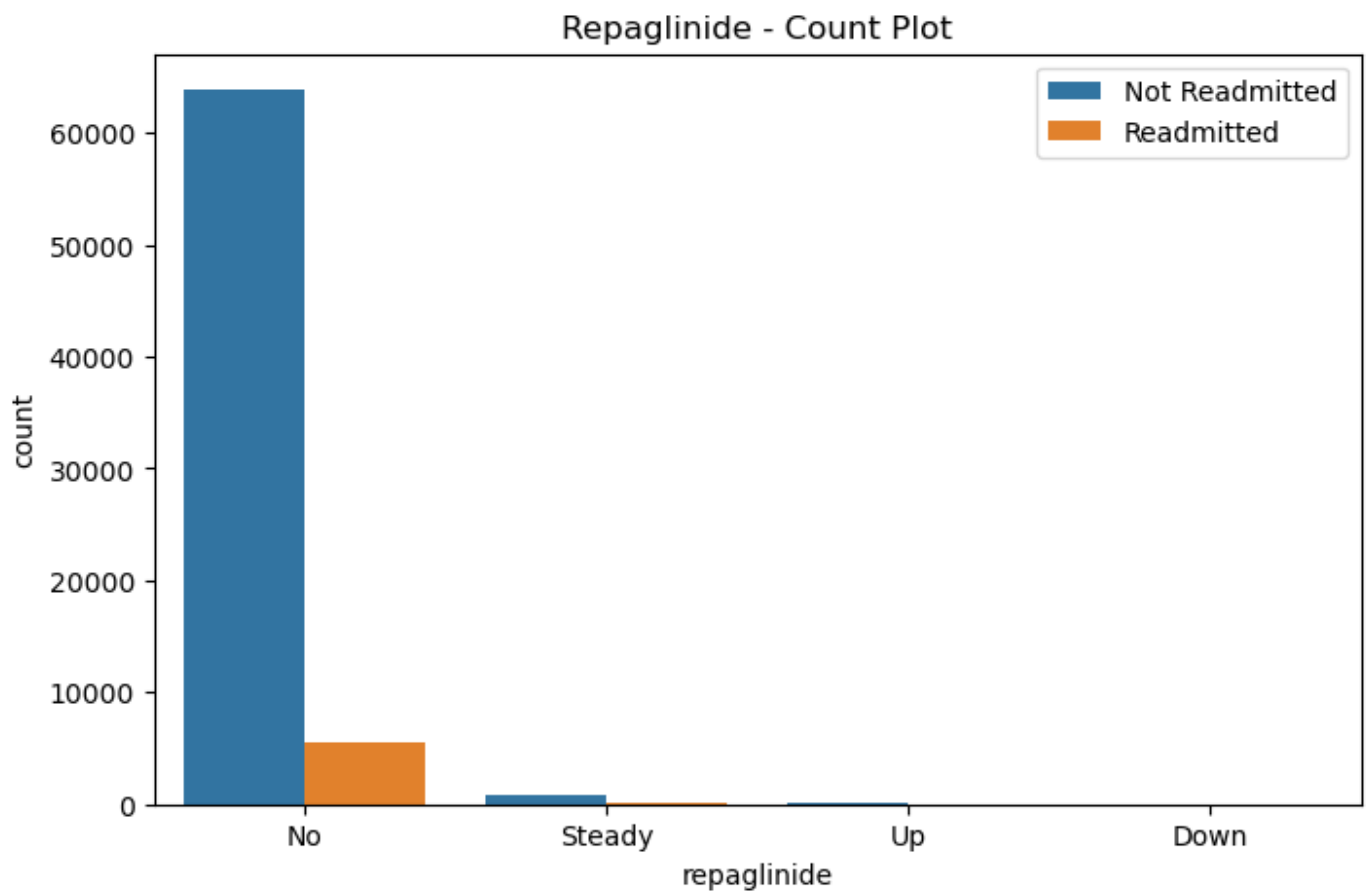


Chlorpropamide - Count Plot



Nateglinide - Count Plot



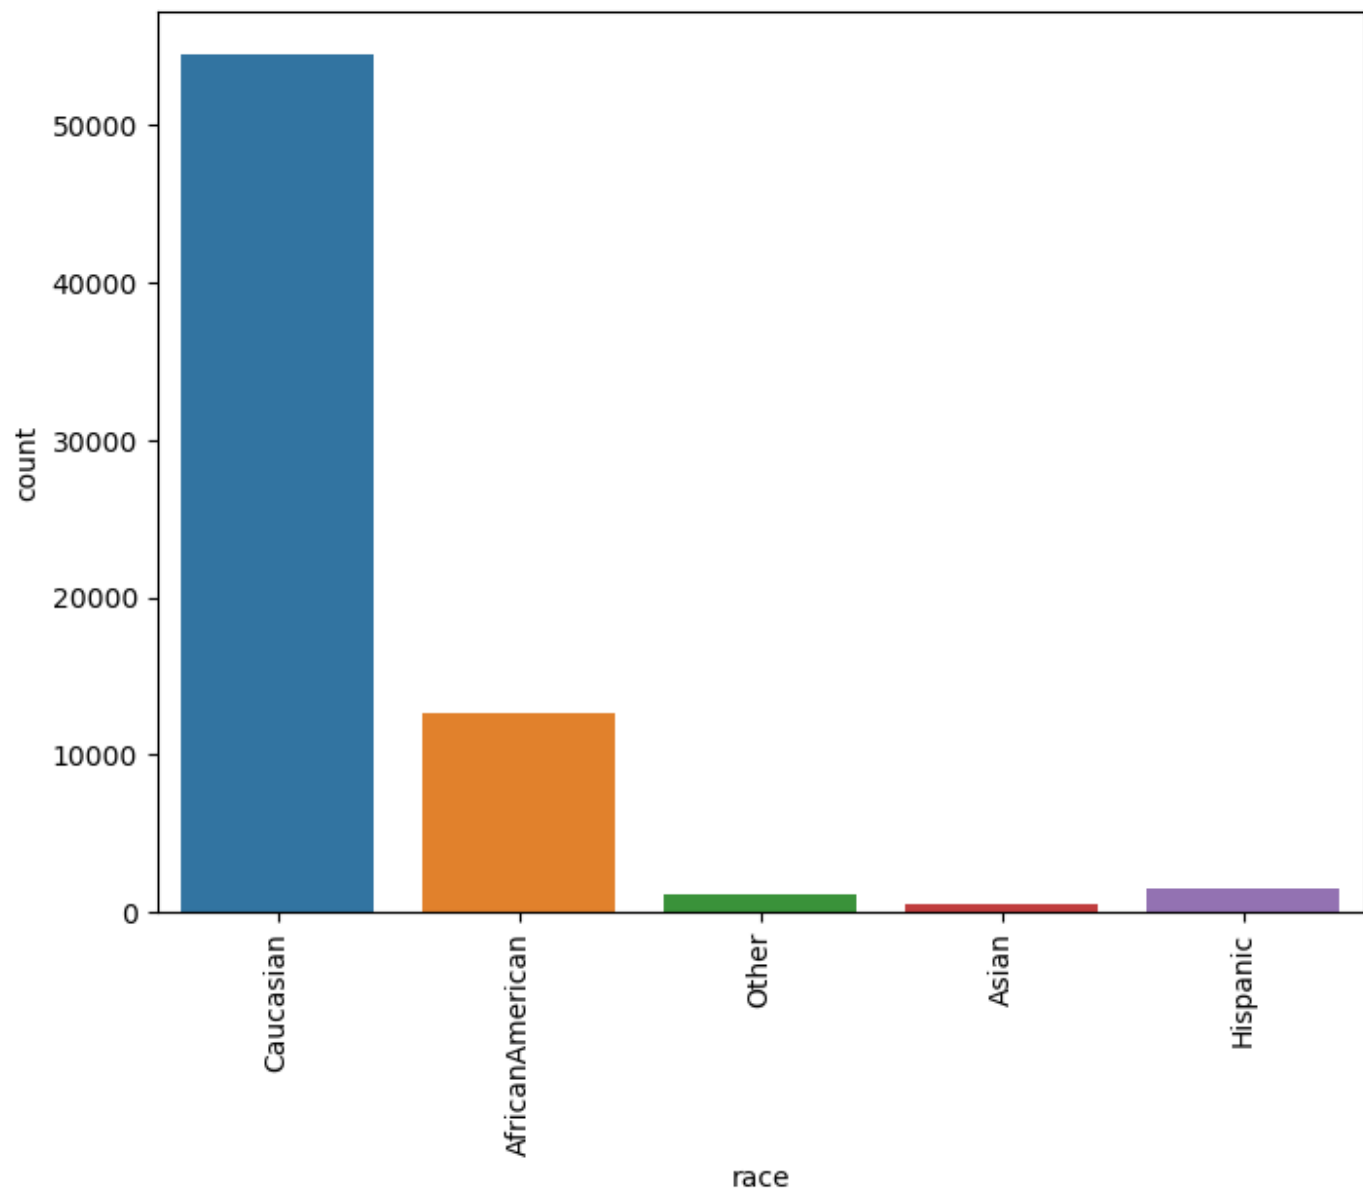


More Countplots for Categorical Variables

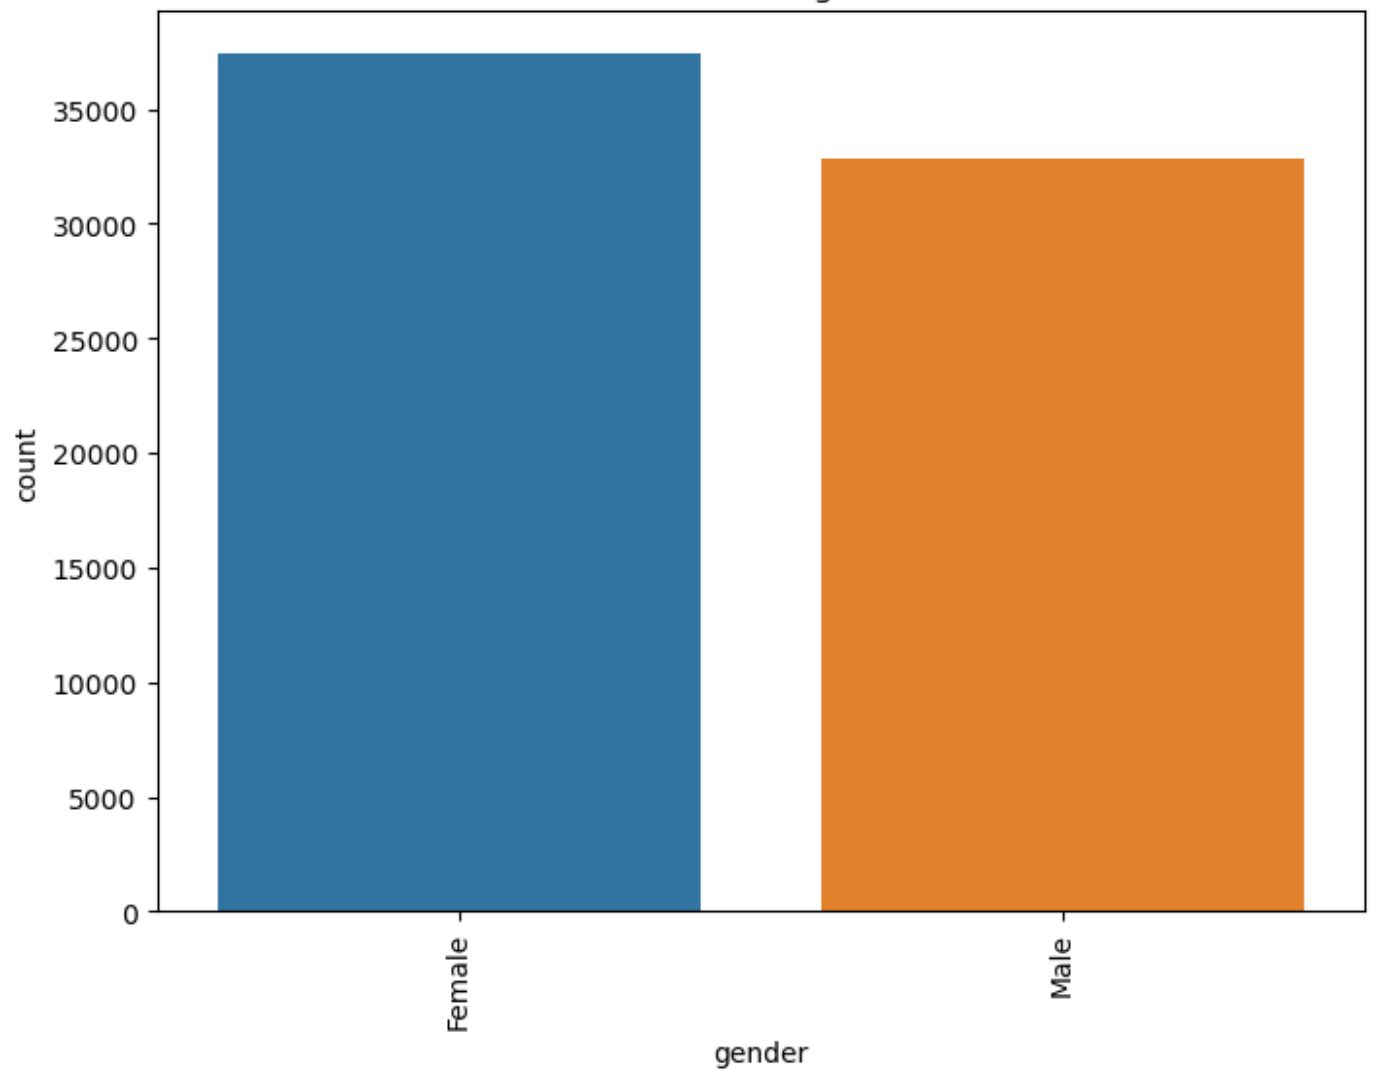
```
In [38]: categorical_vars = df.select_dtypes(include=['object', 'category']).columns

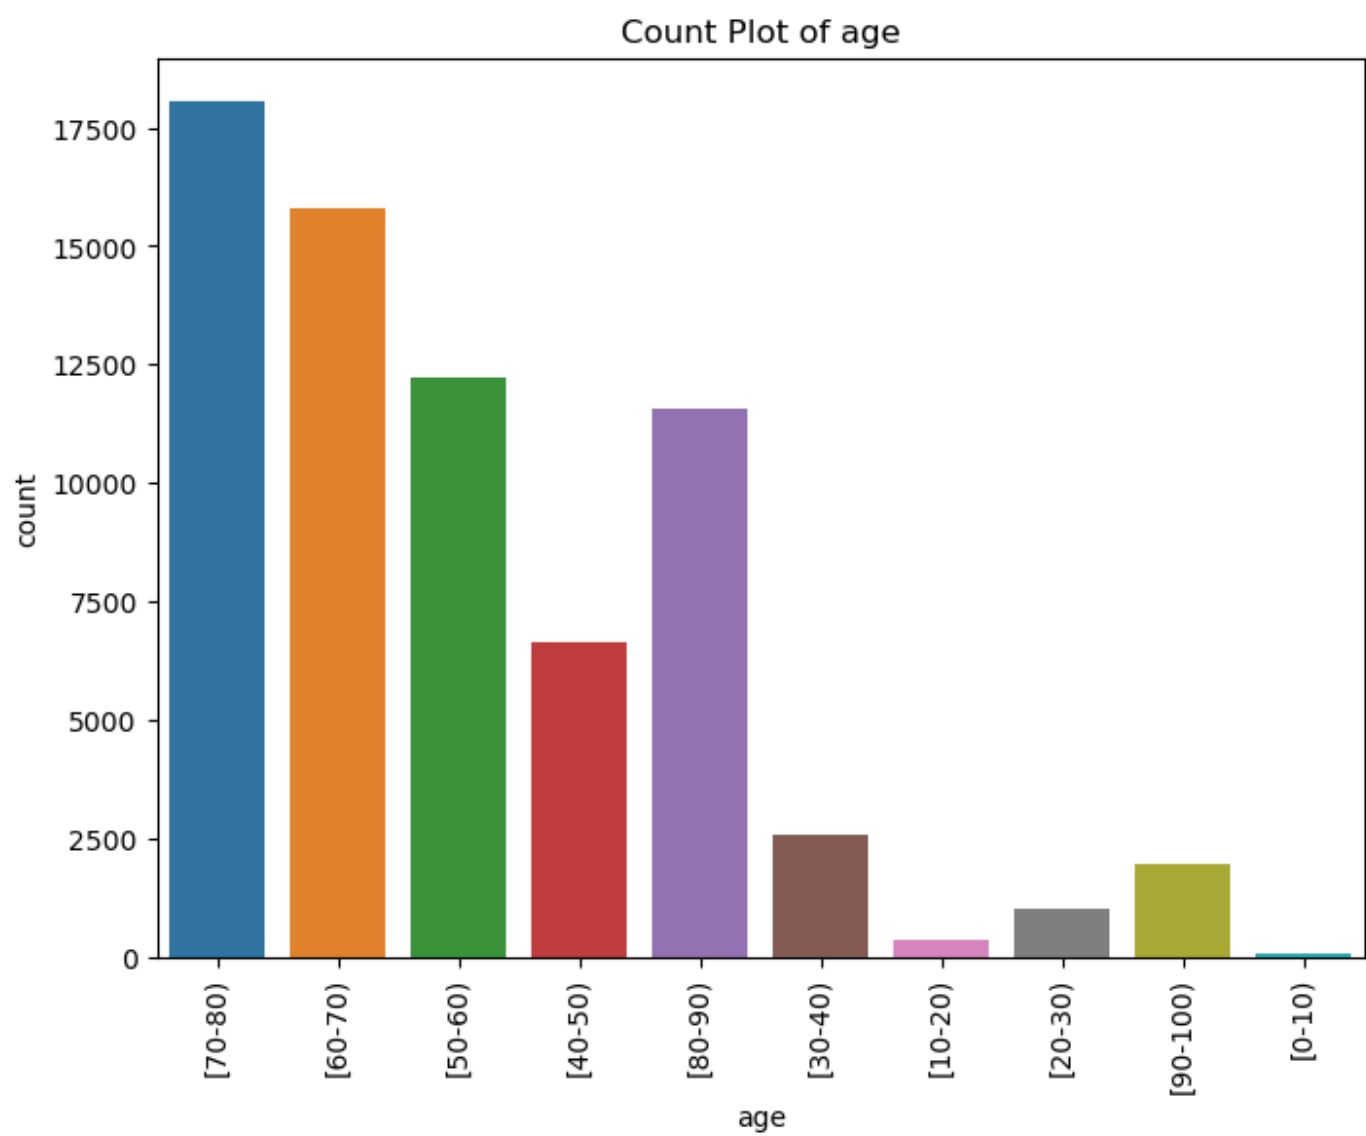
# Create count plots for each categorical variable
for var in categorical_vars:
    plt.figure(figsize=(8, 6)) # Adjust the figure size as needed
    sns.countplot(x=var, data=df)
    plt.title(f'Count Plot of {var}')
    plt.xticks(rotation=90) # Rotate x-axis labels for better readability
    plt.show()
```


Count Plot of race

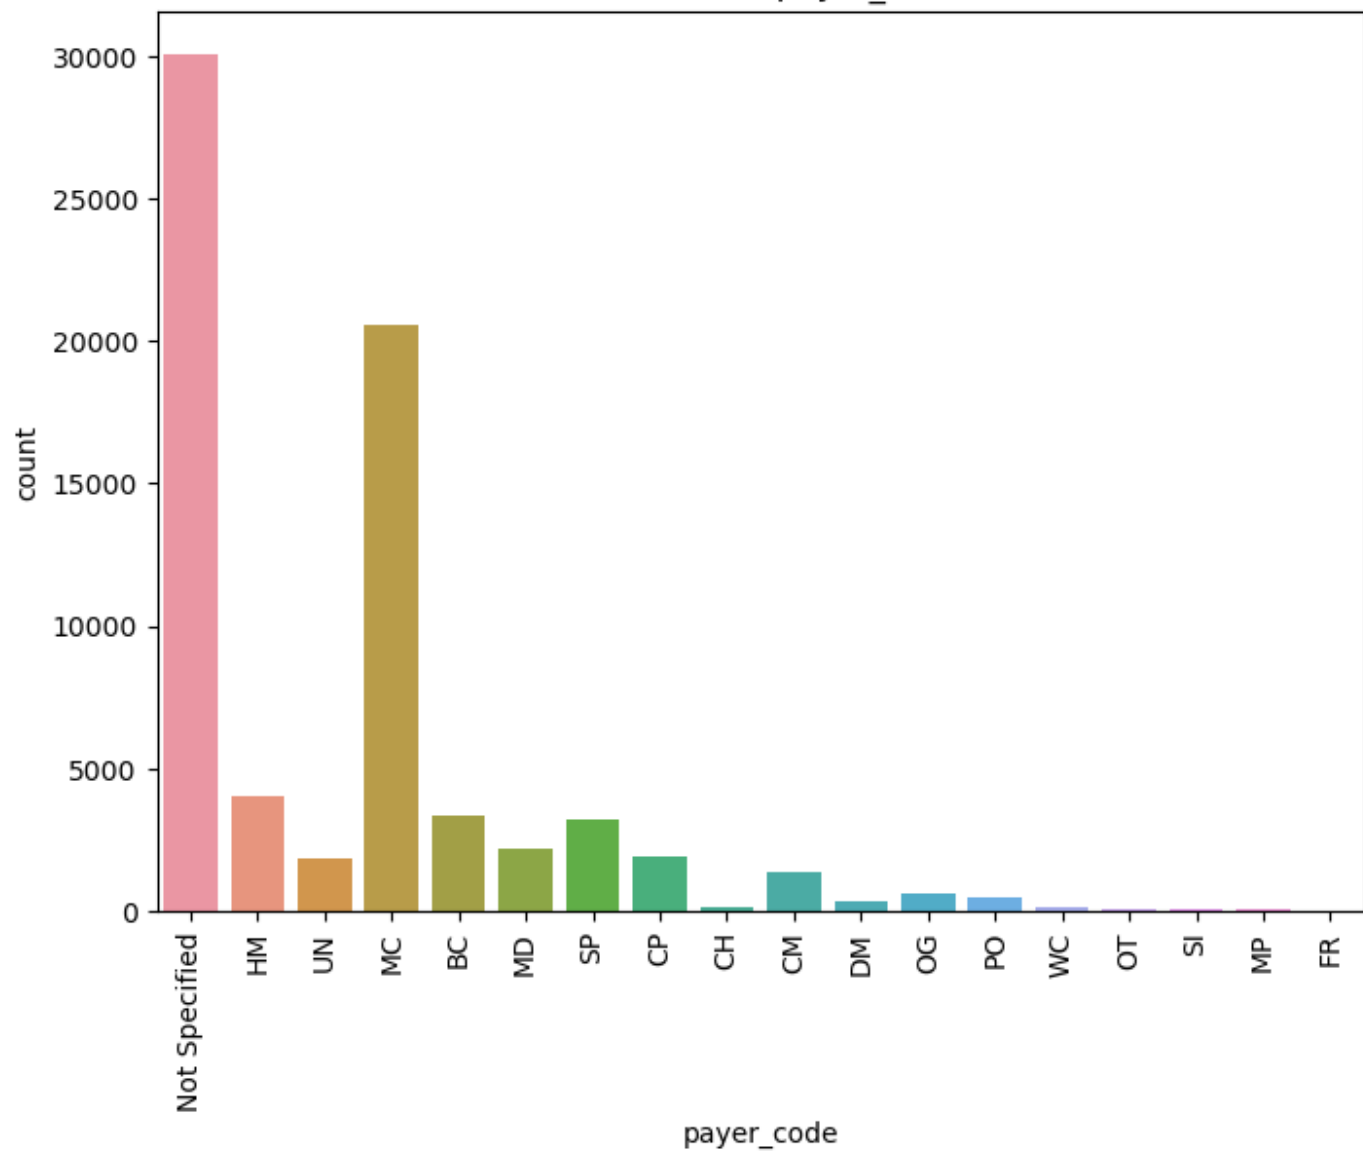


Count Plot of gender

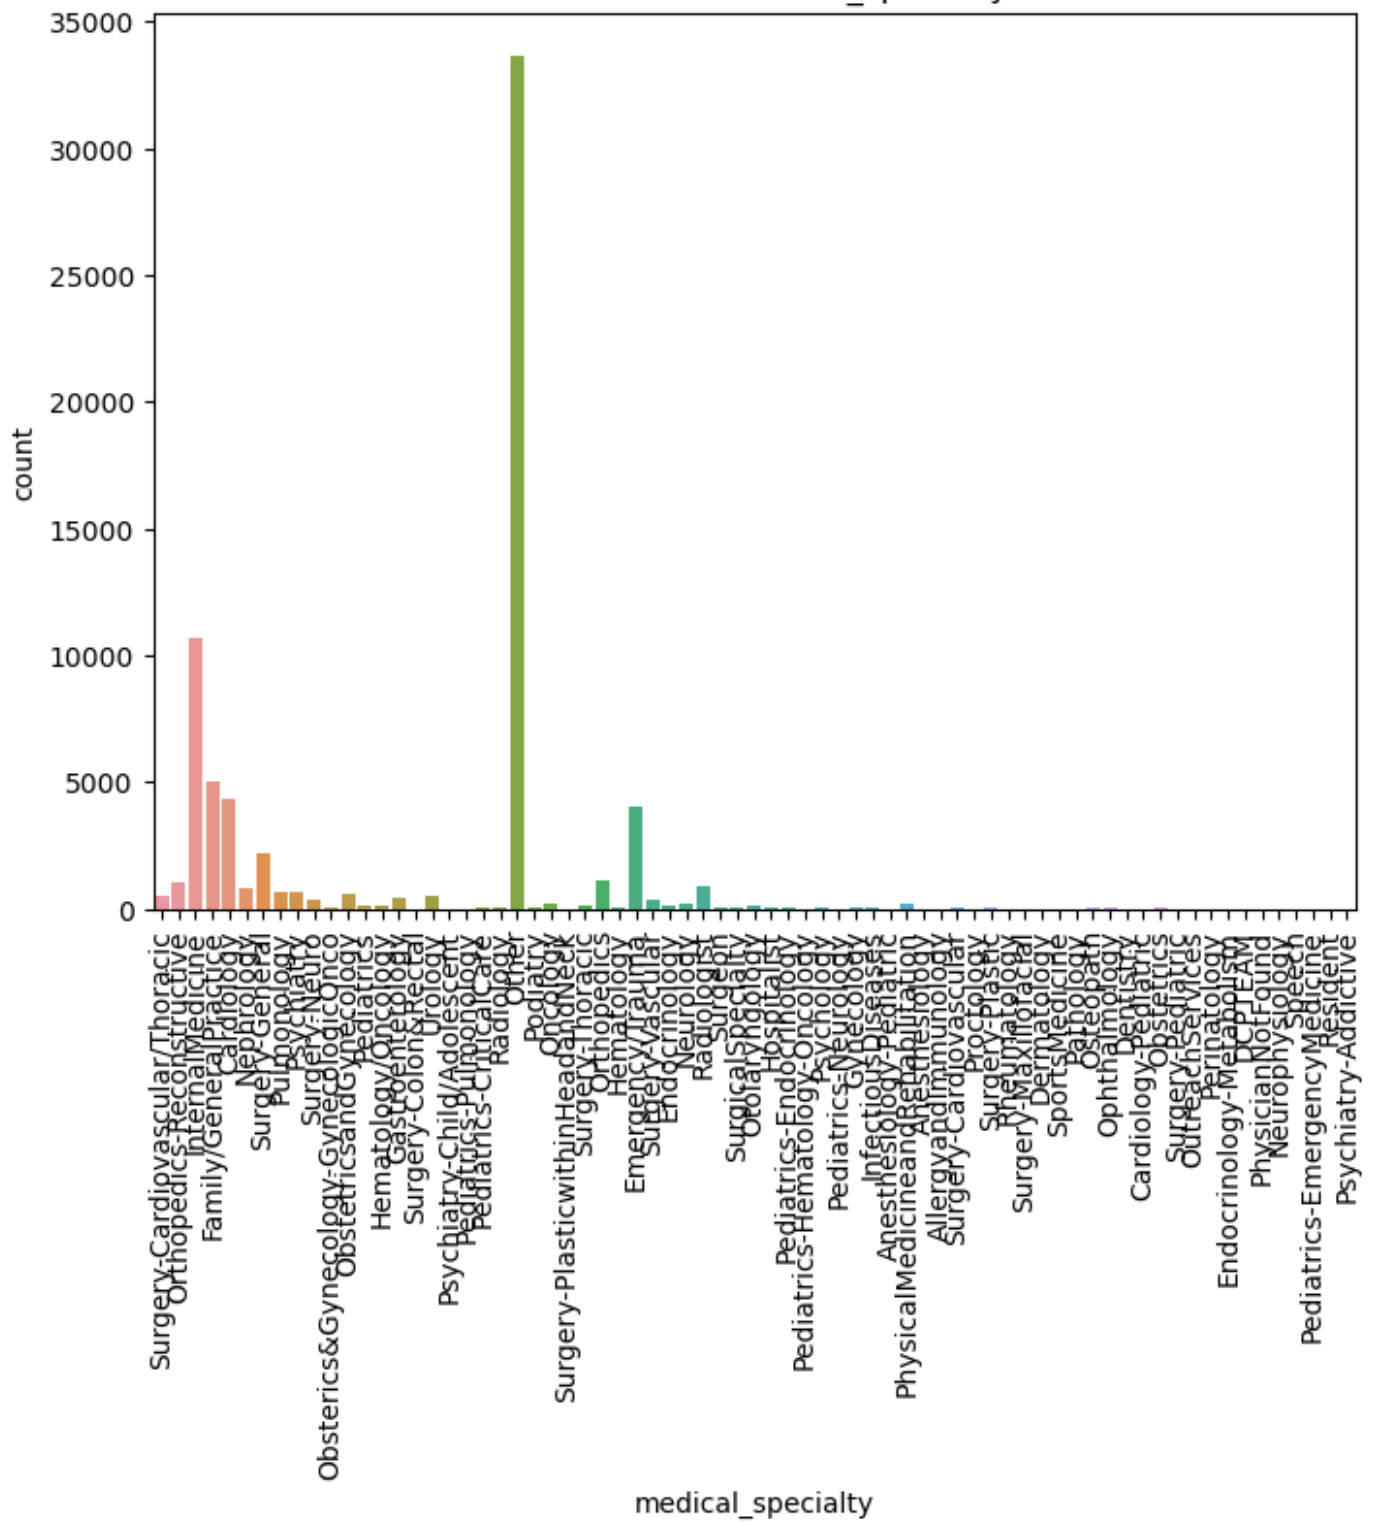




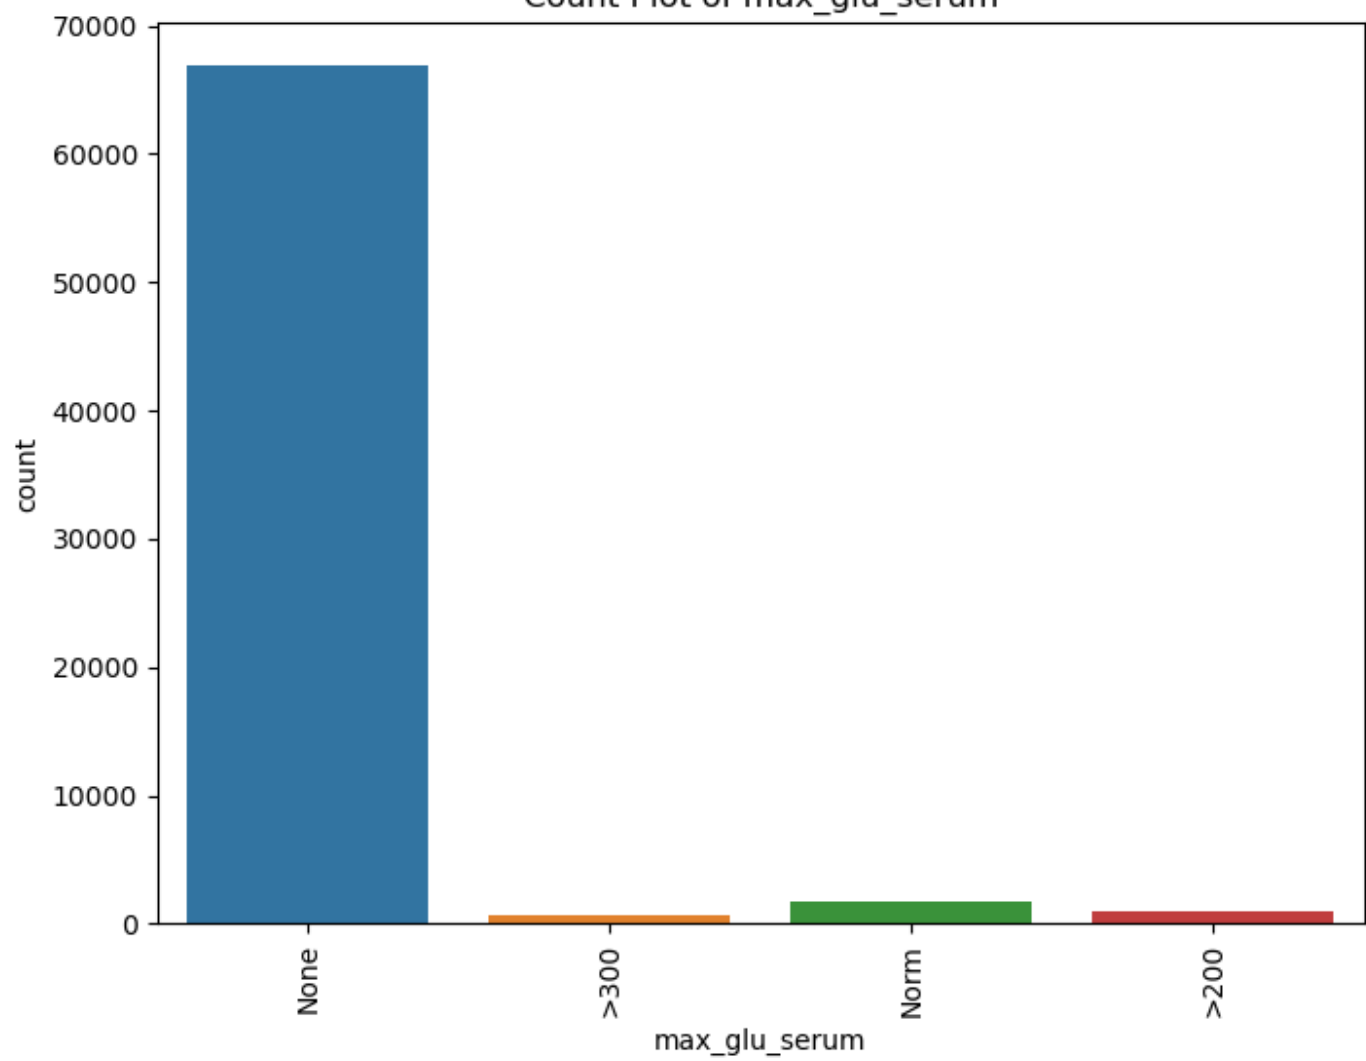
Count Plot of payer_code



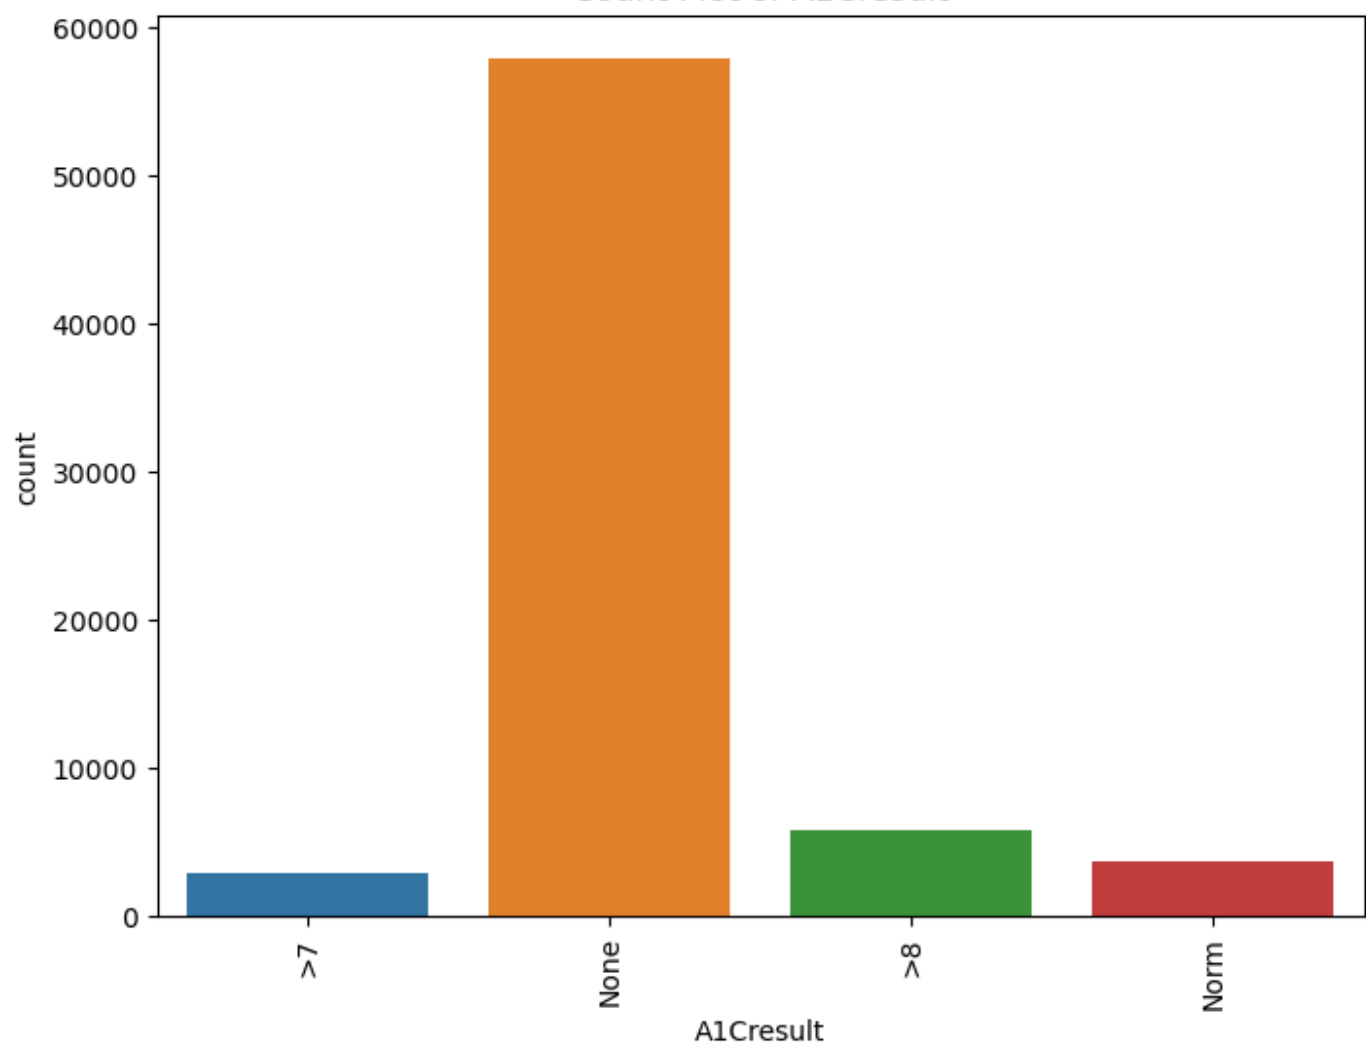
Count Plot of medical_specialty



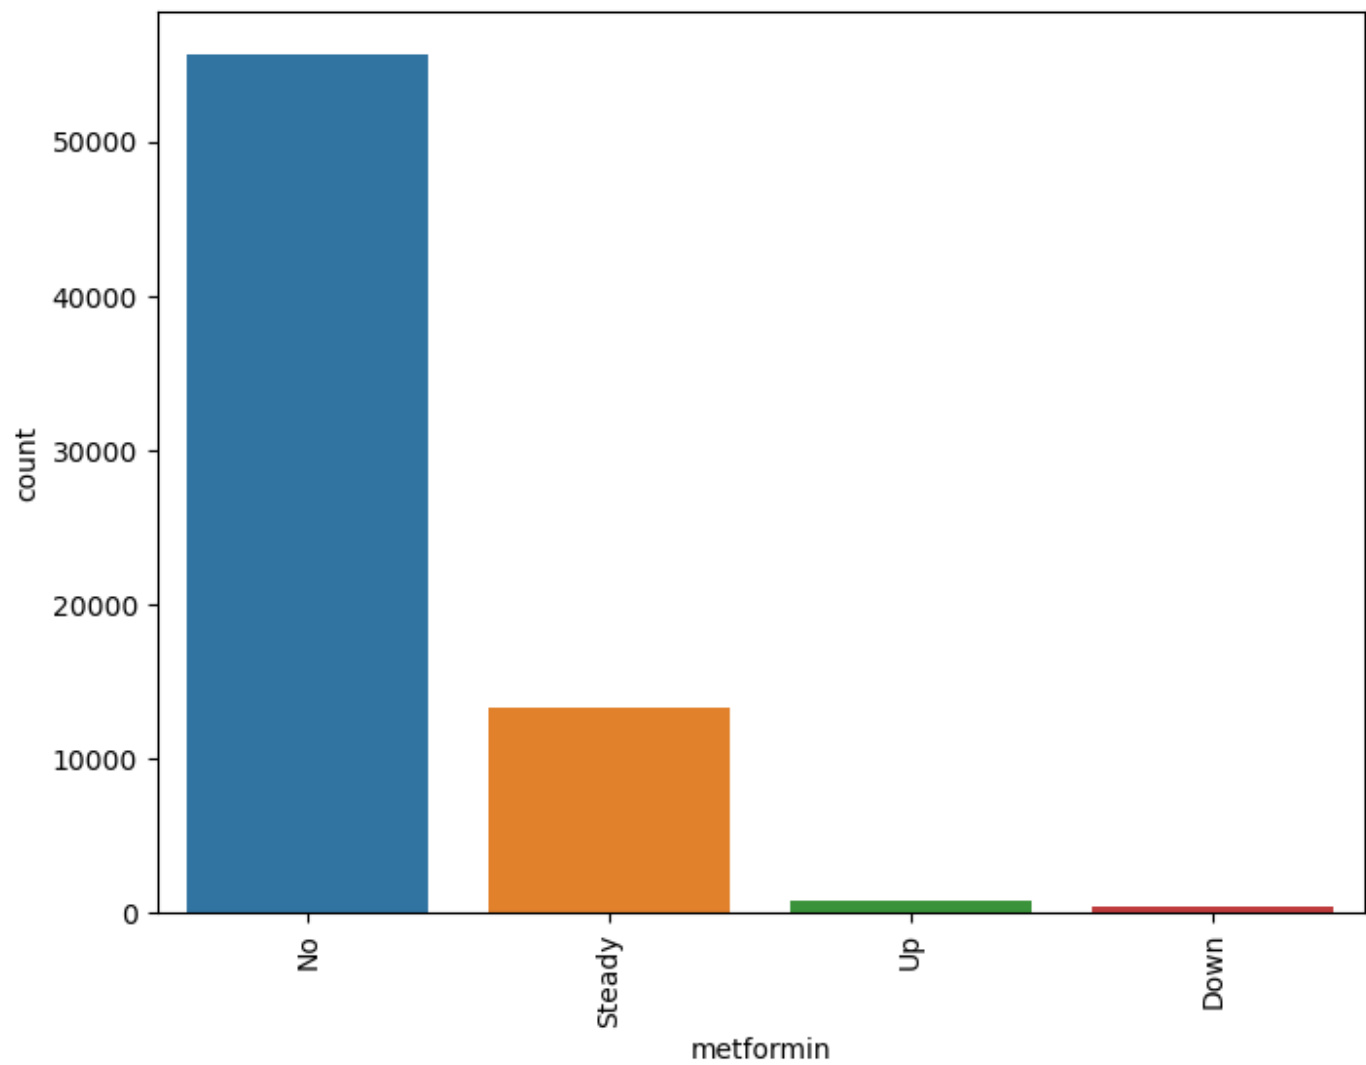
Count Plot of max_glu_serum



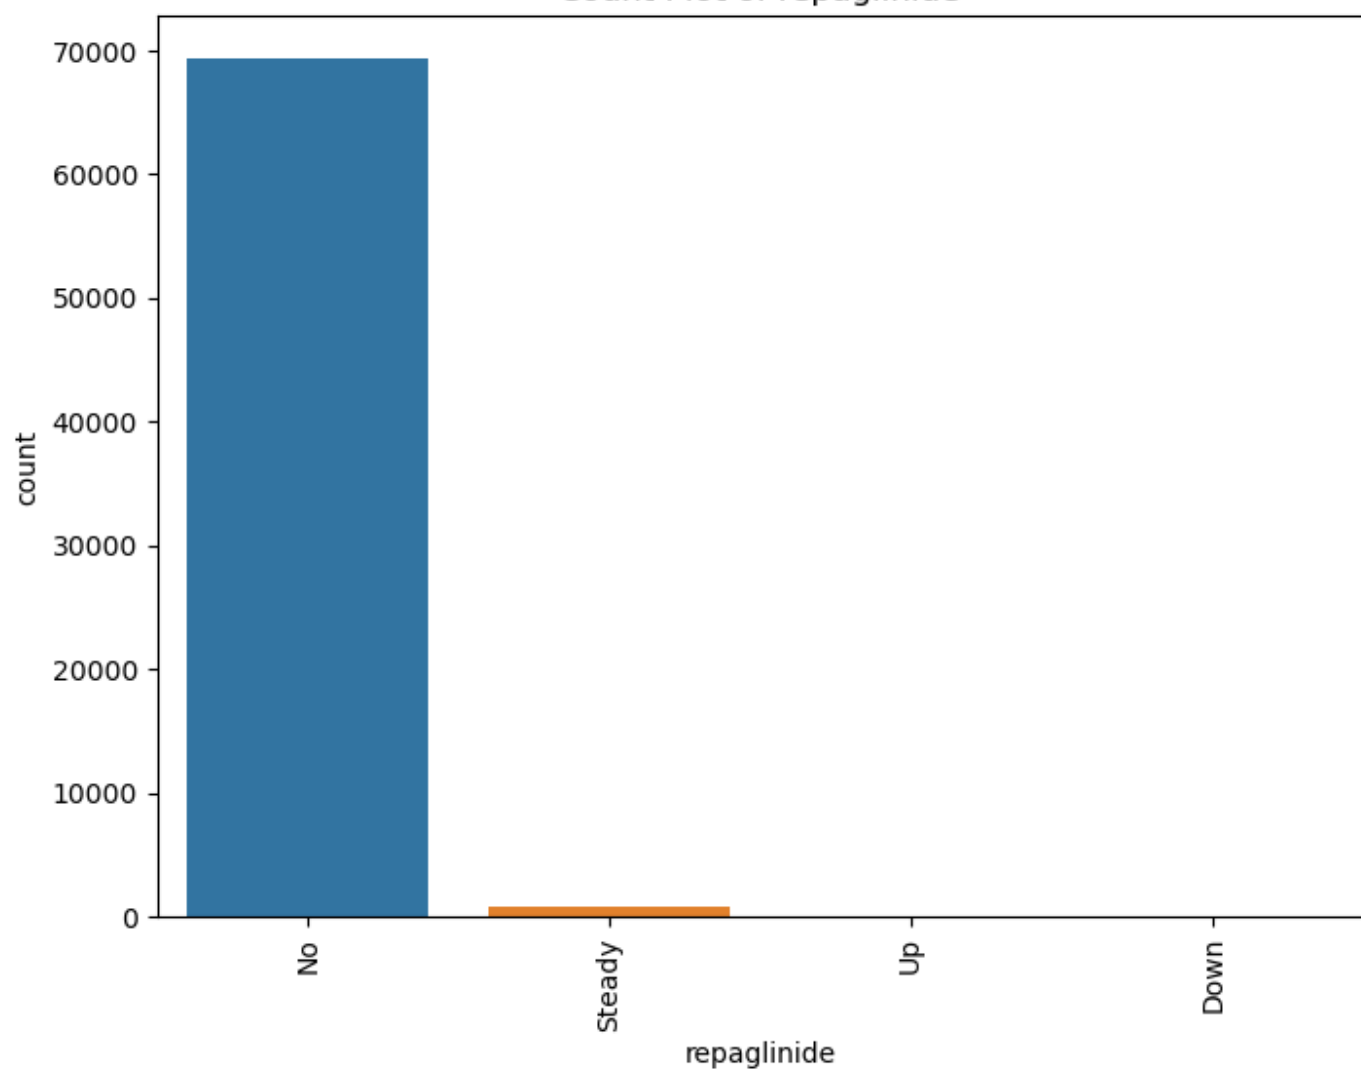
Count Plot of A1Cresult



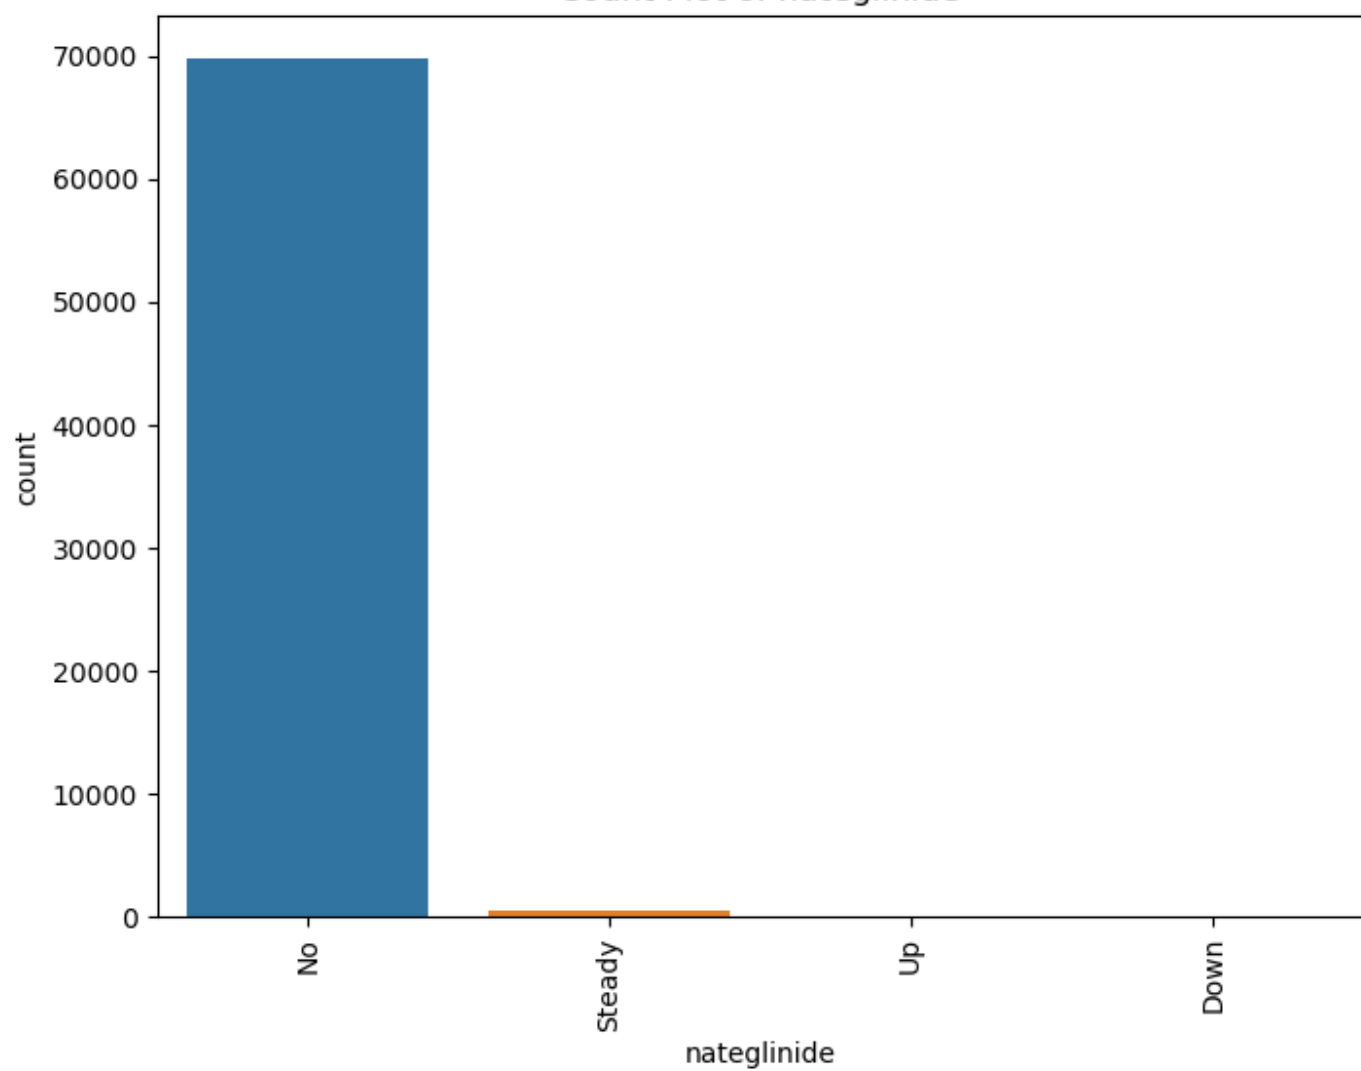
Count Plot of metformin



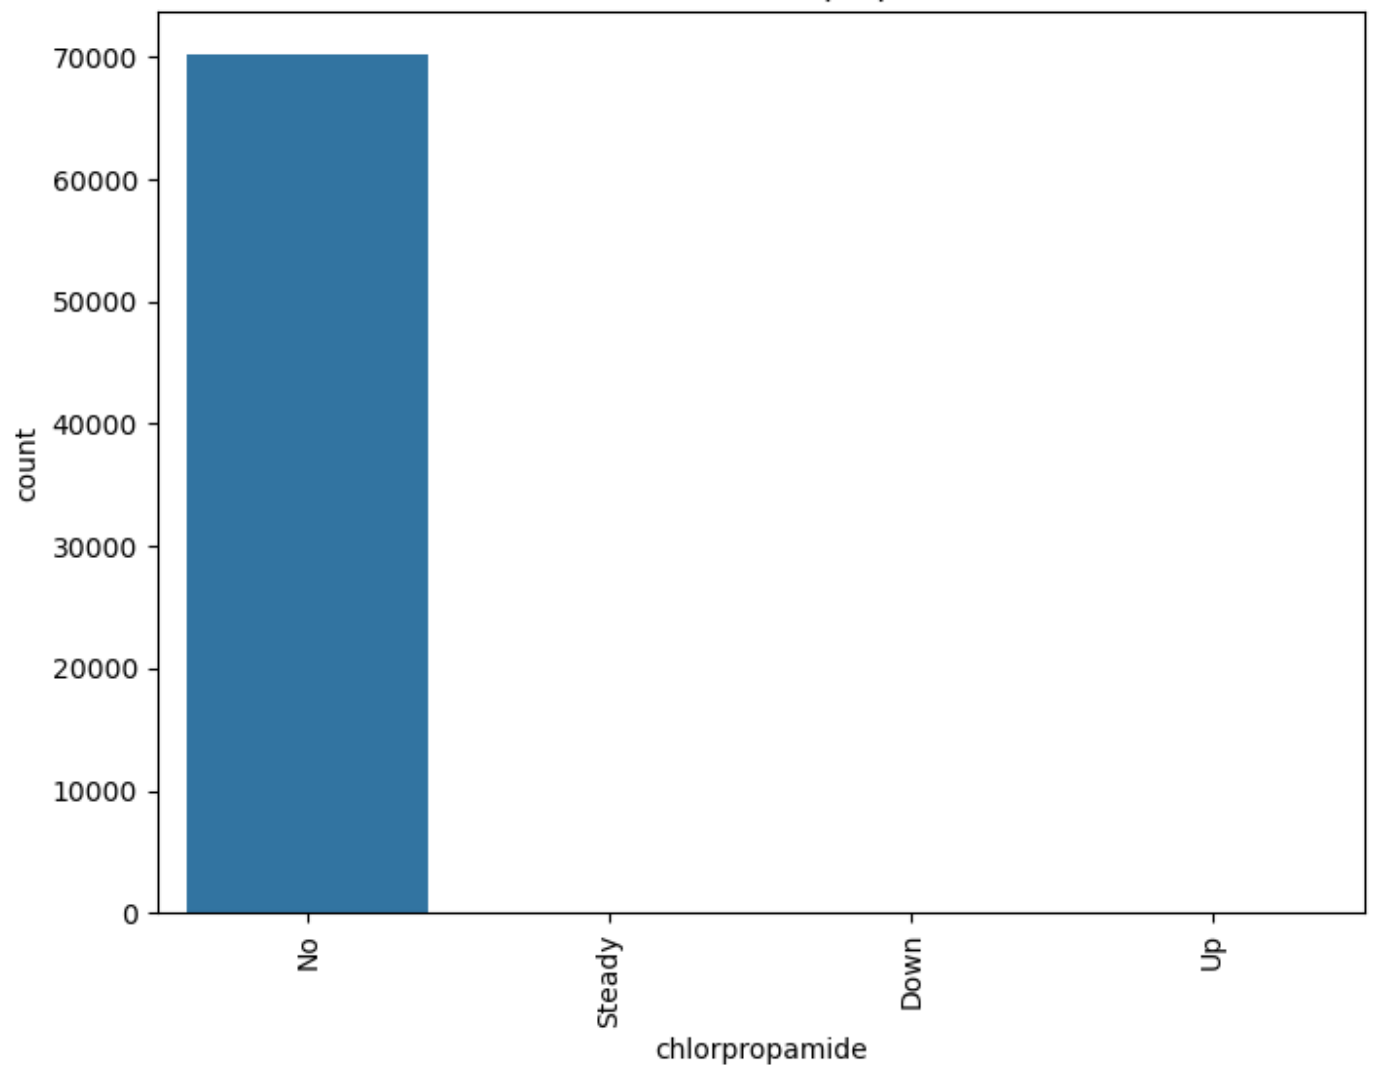
Count Plot of repaglinide



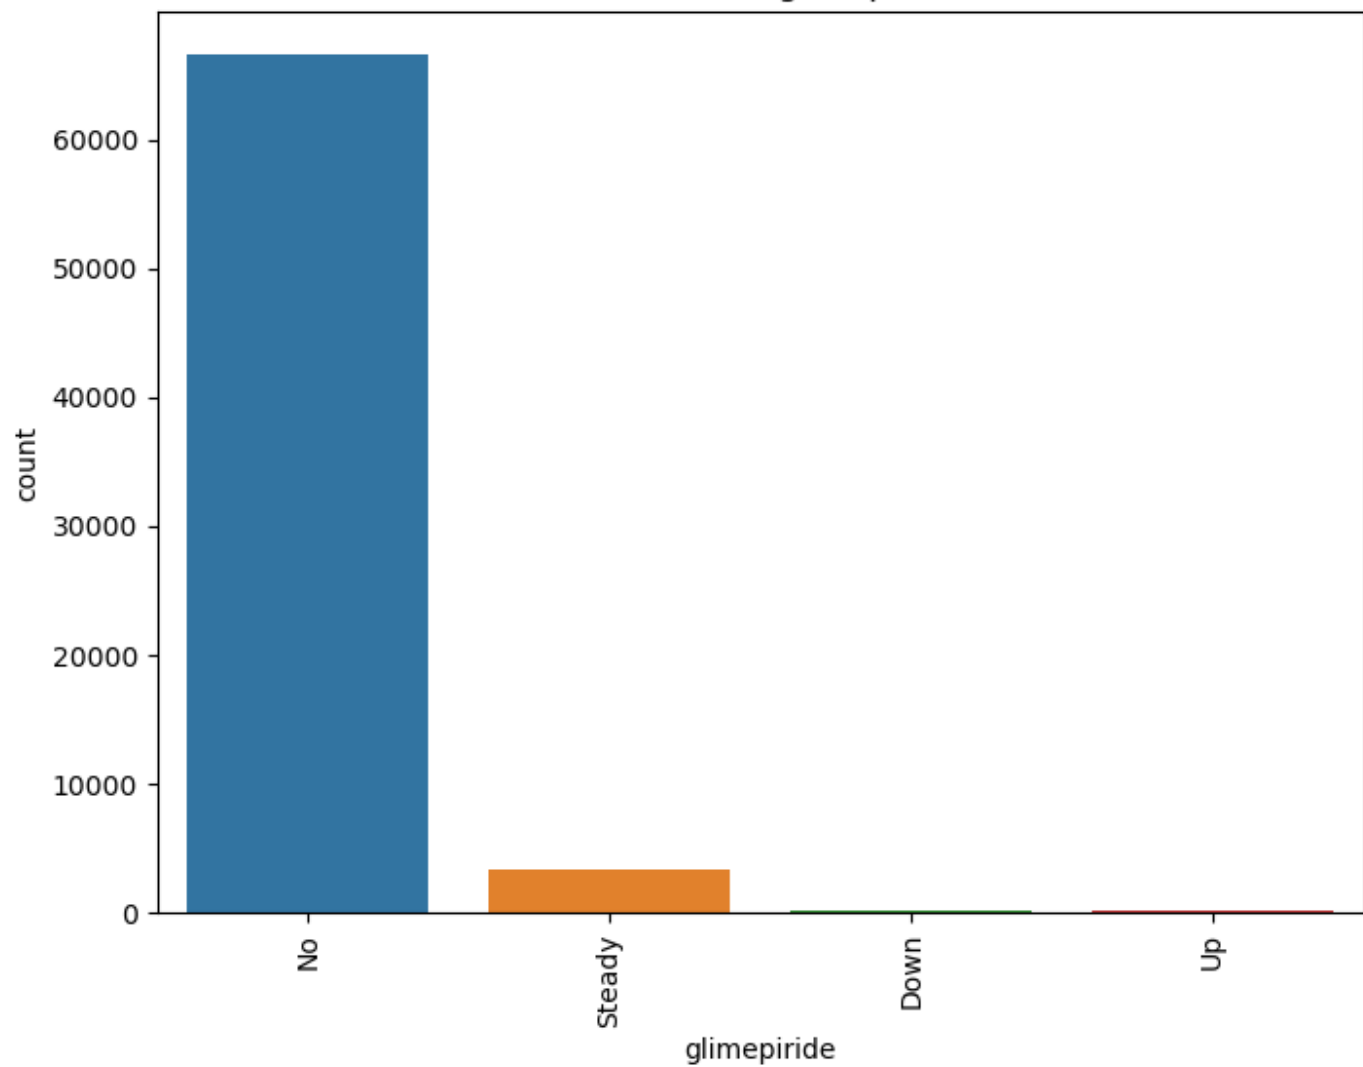
Count Plot of nateglinide



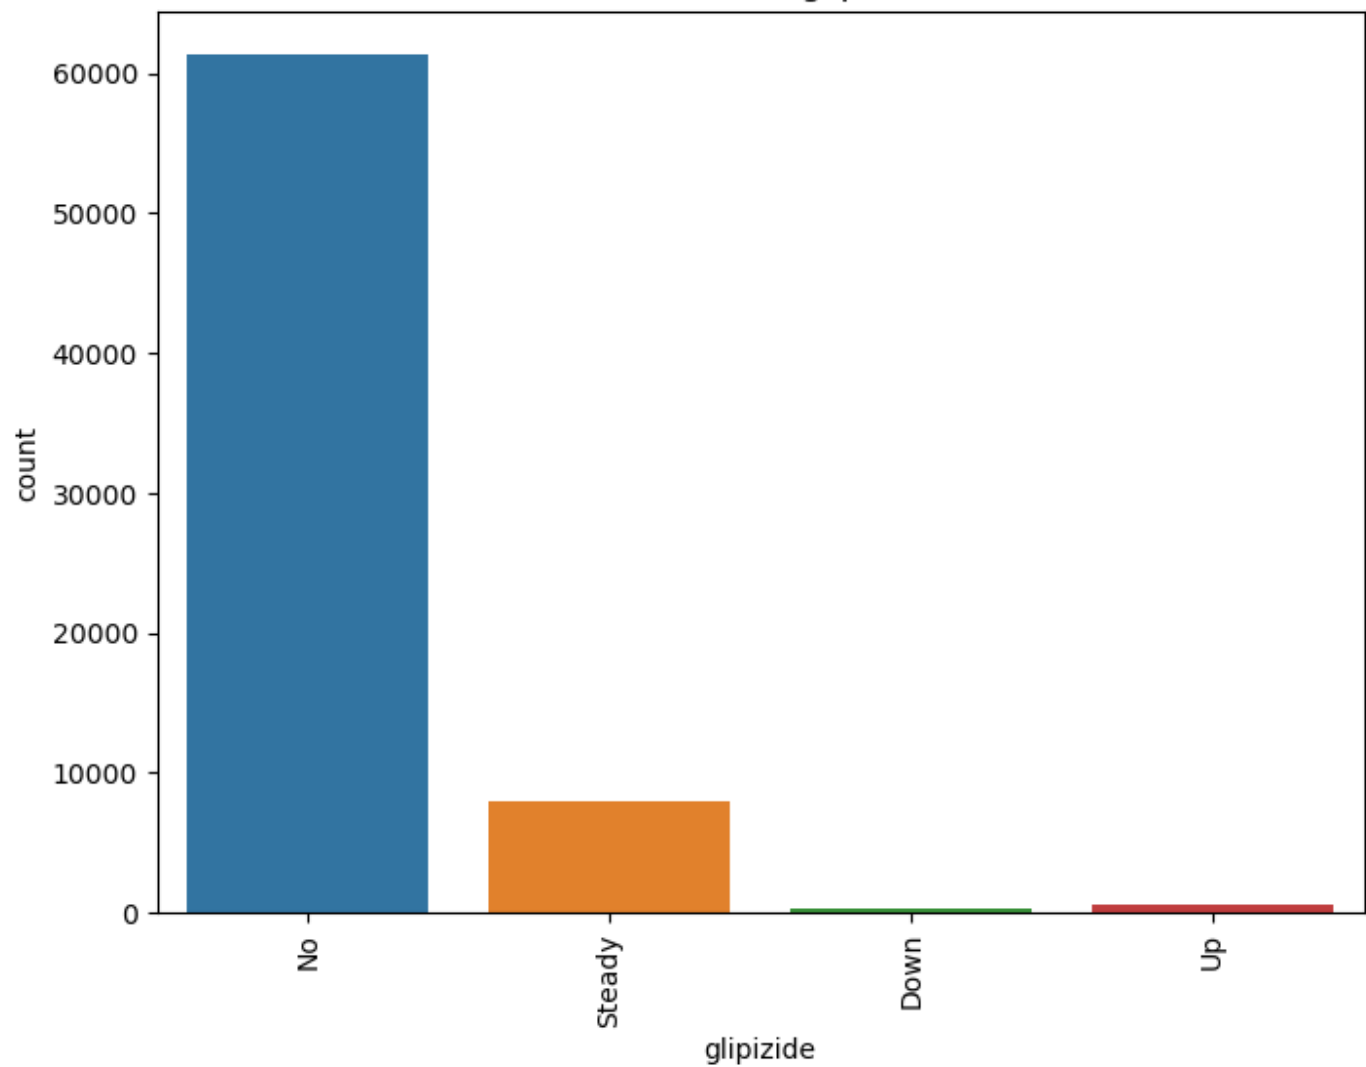
Count Plot of chlorpropamide



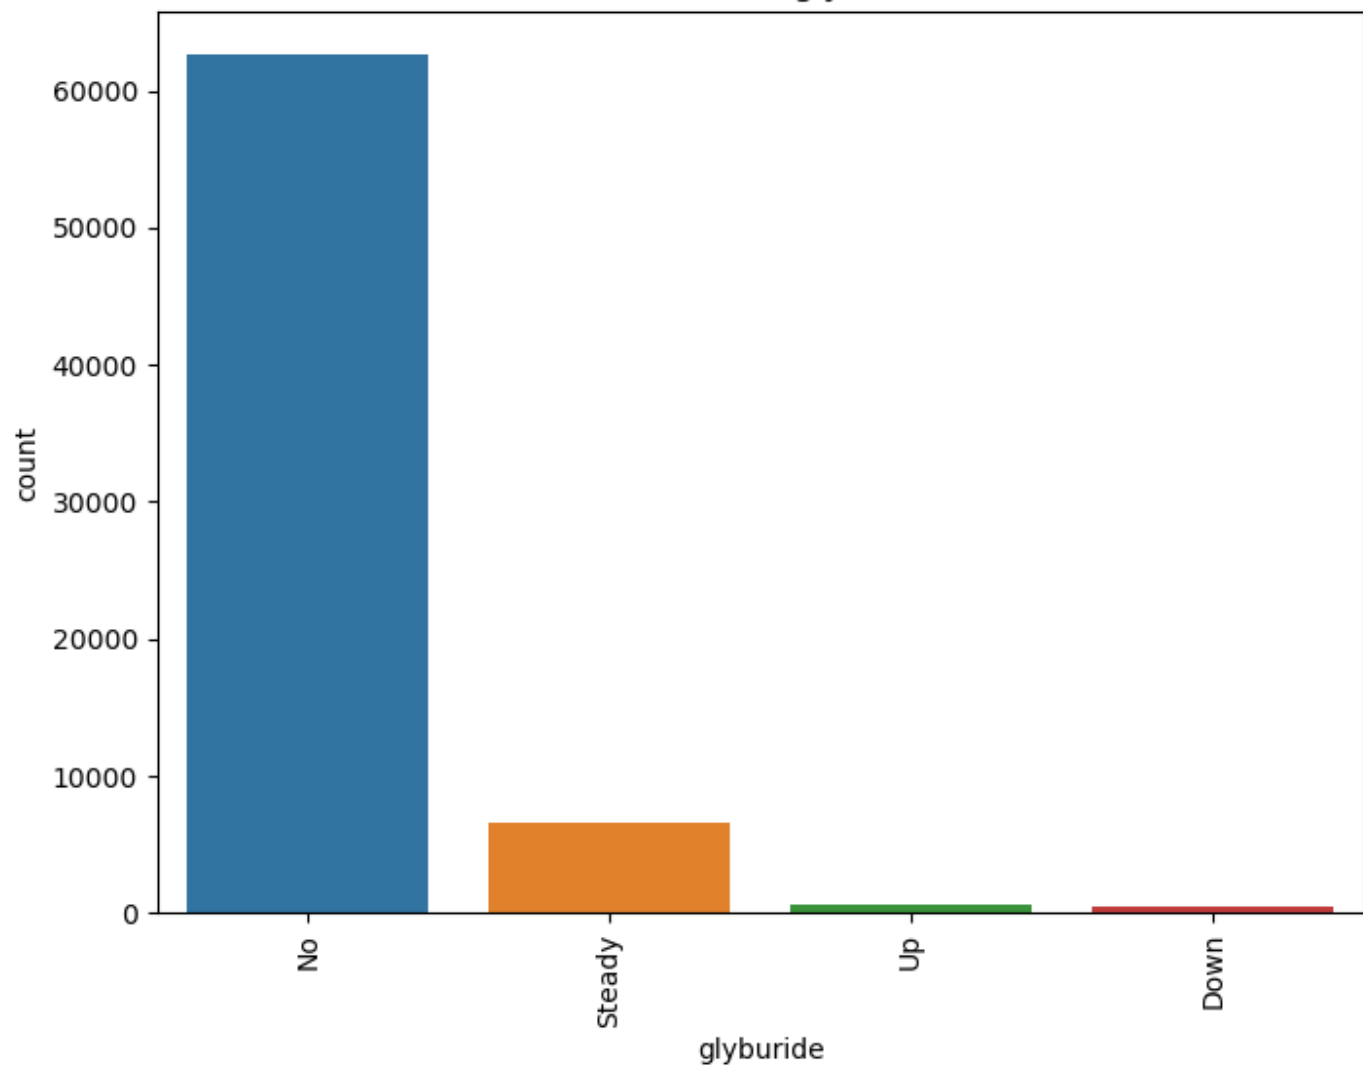
Count Plot of glimepiride



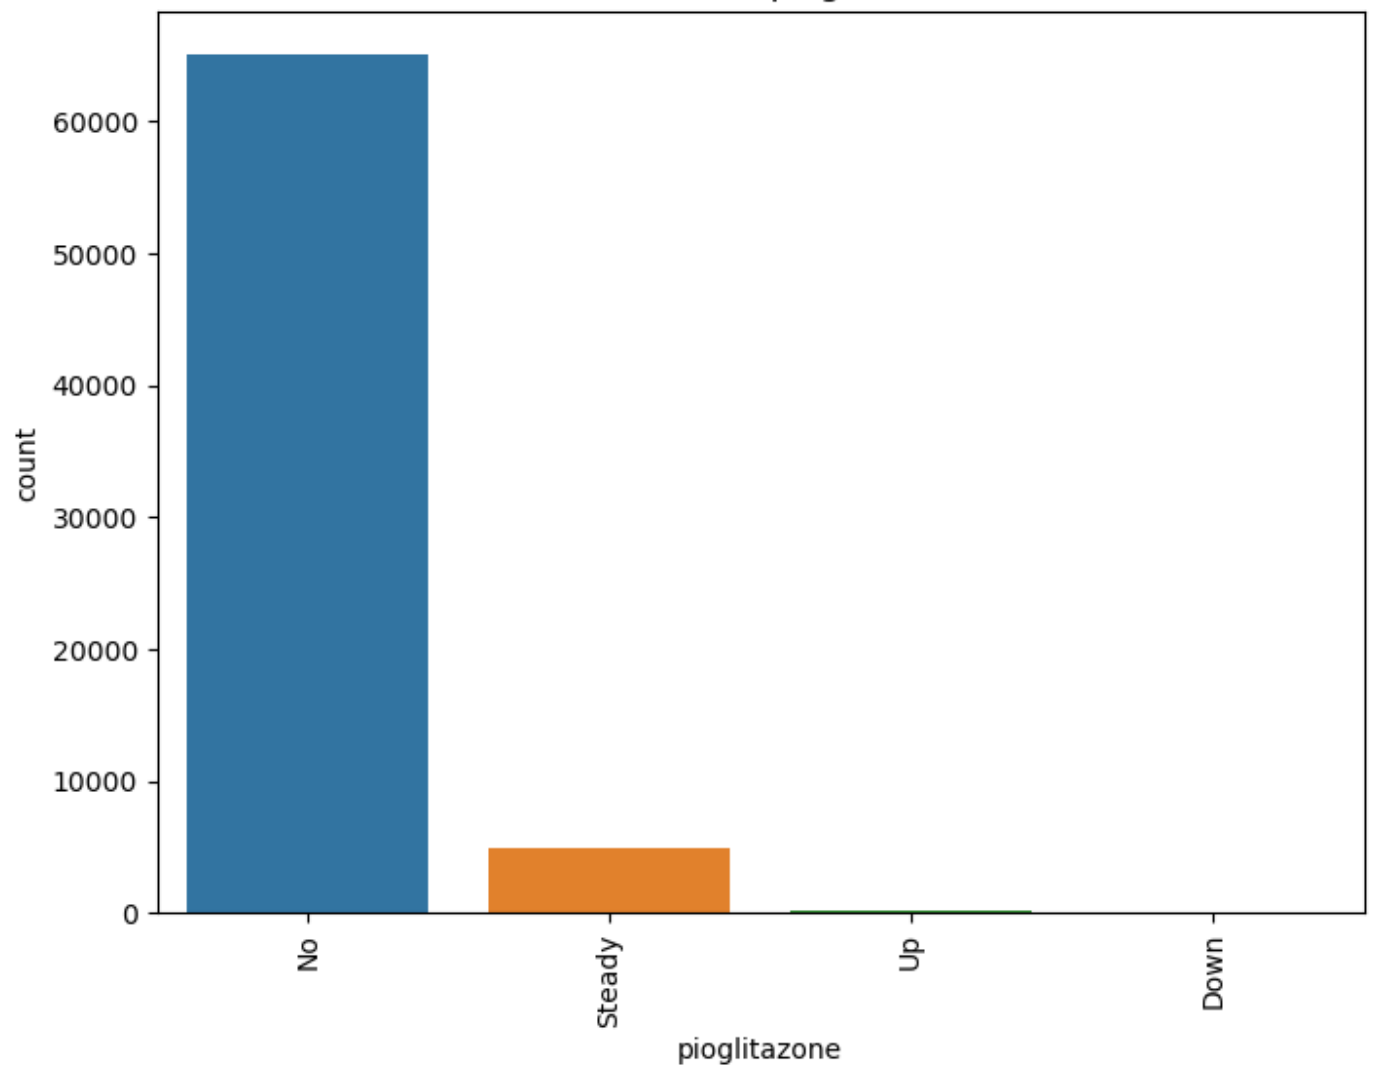
Count Plot of glipizide



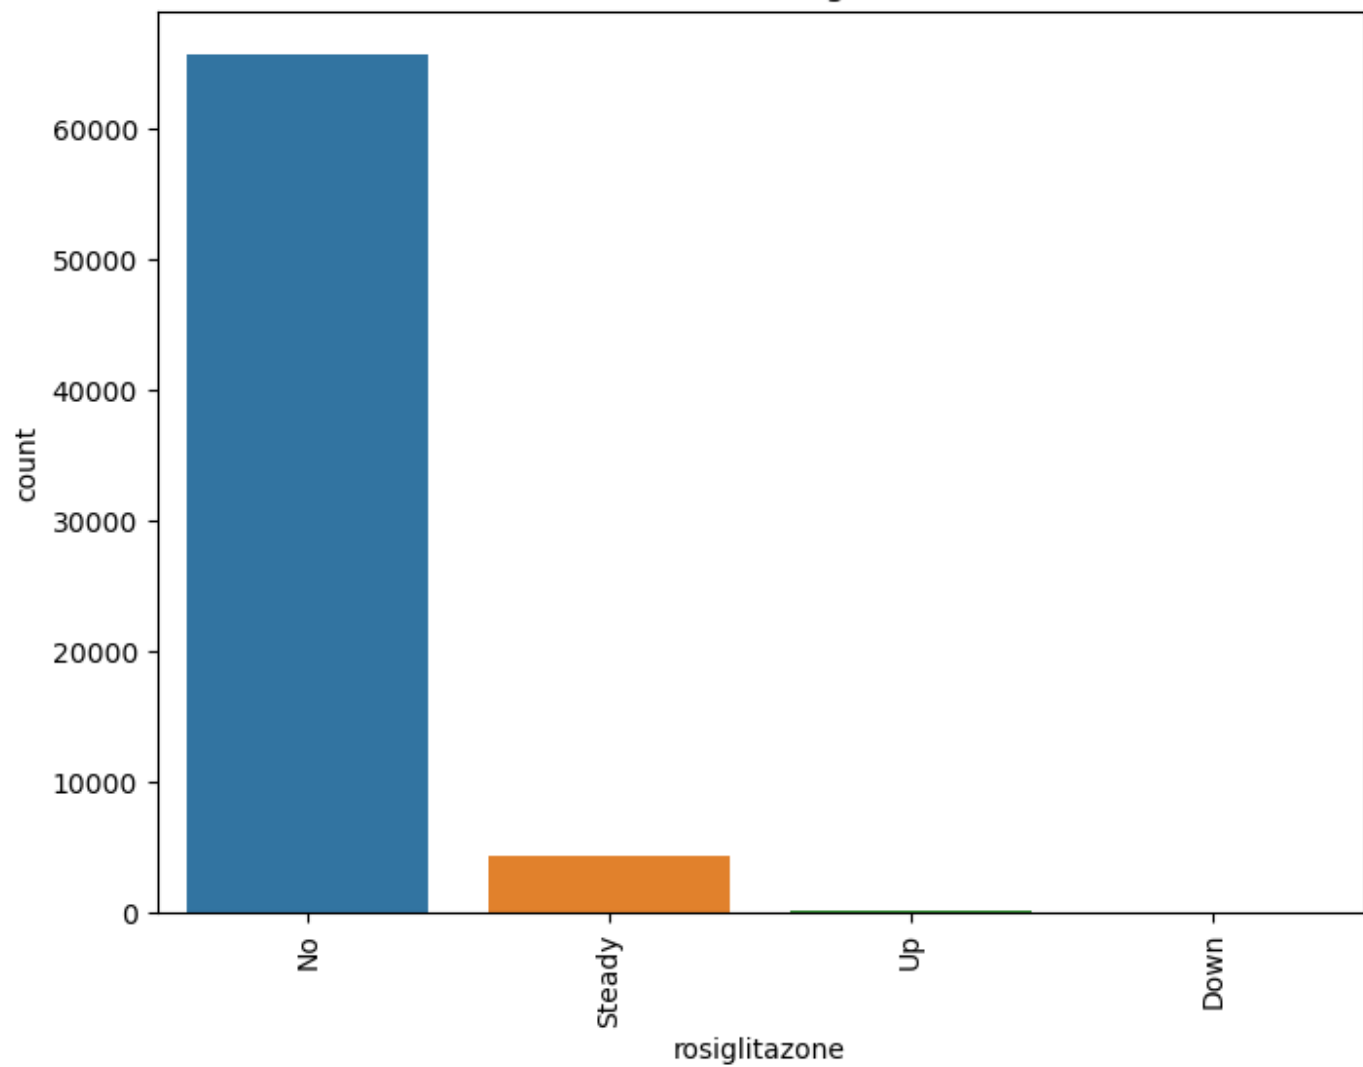
Count Plot of glyburide



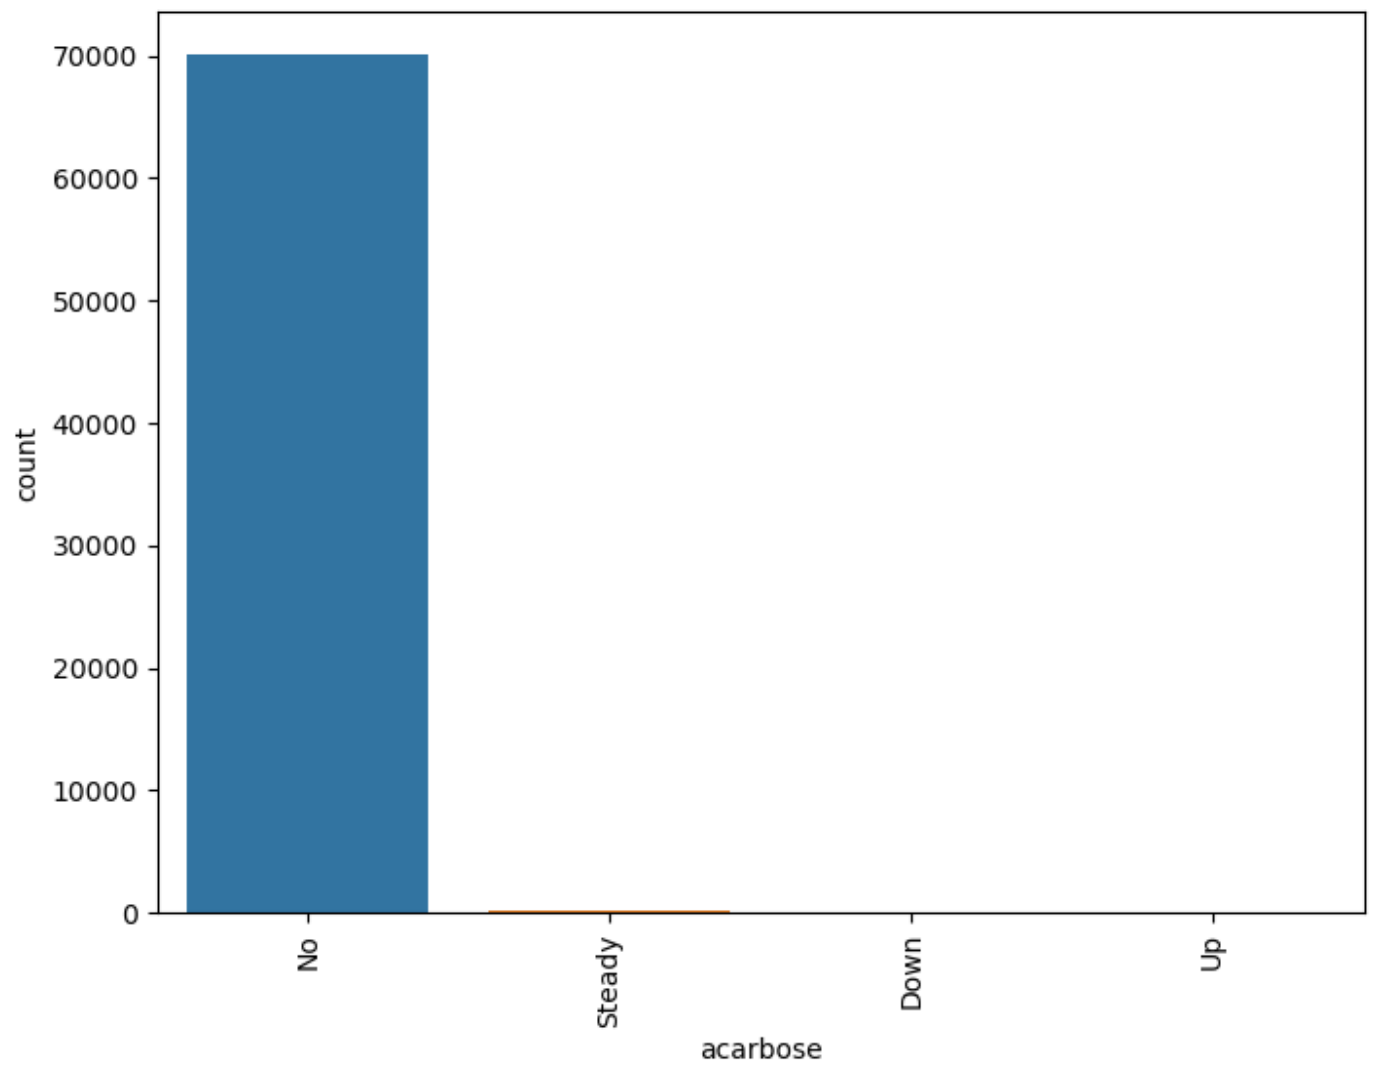
Count Plot of pioglitazone



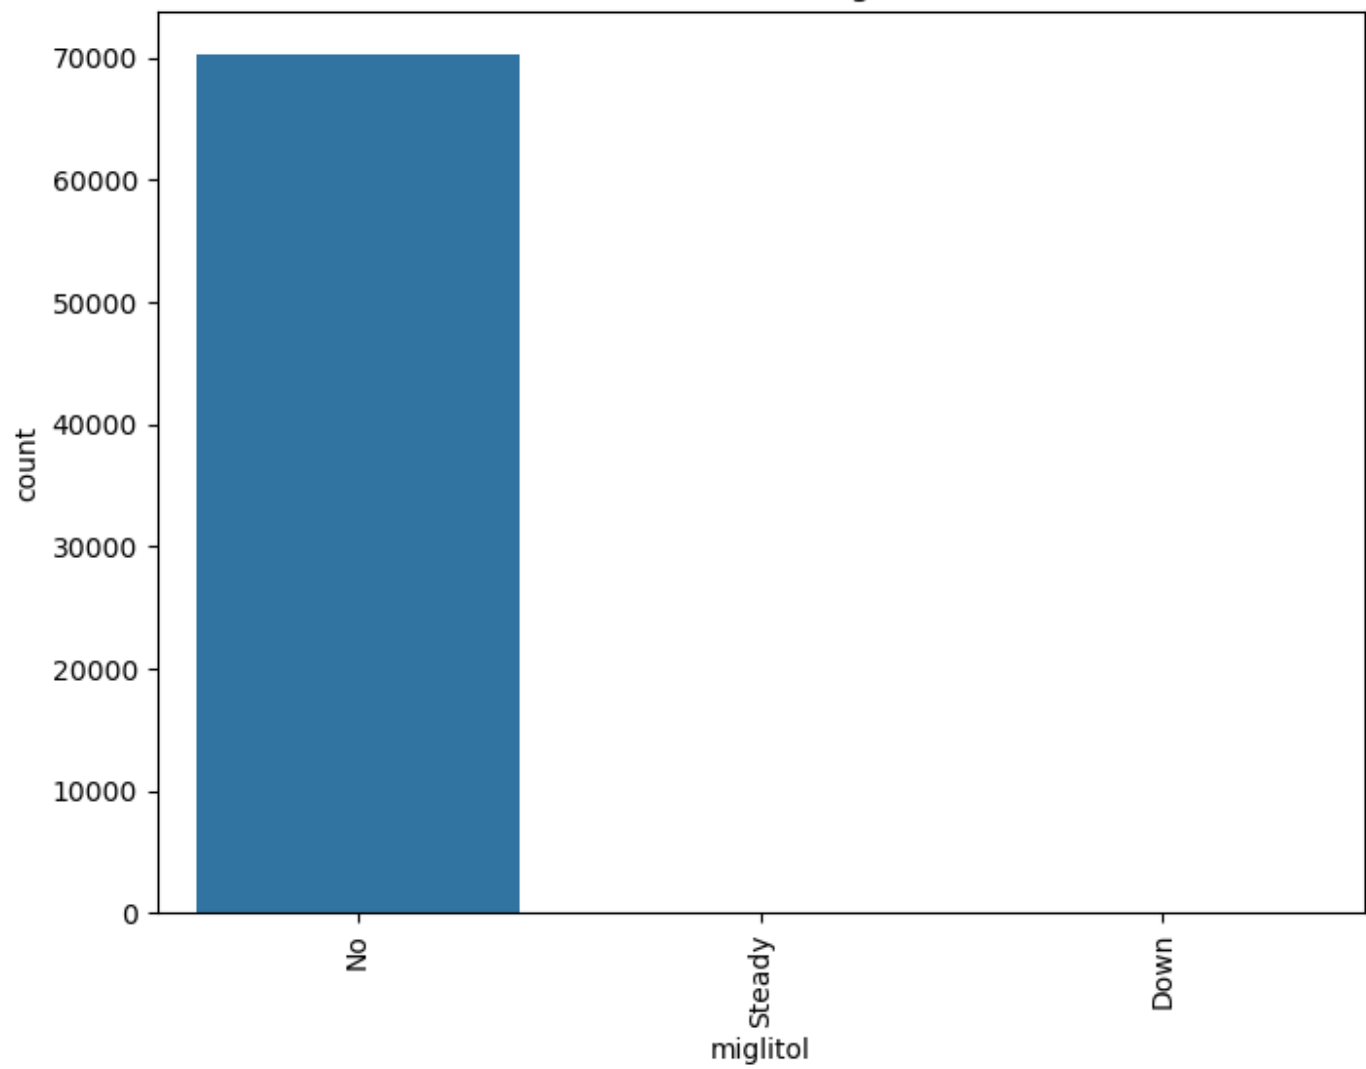
Count Plot of rosiglitazone



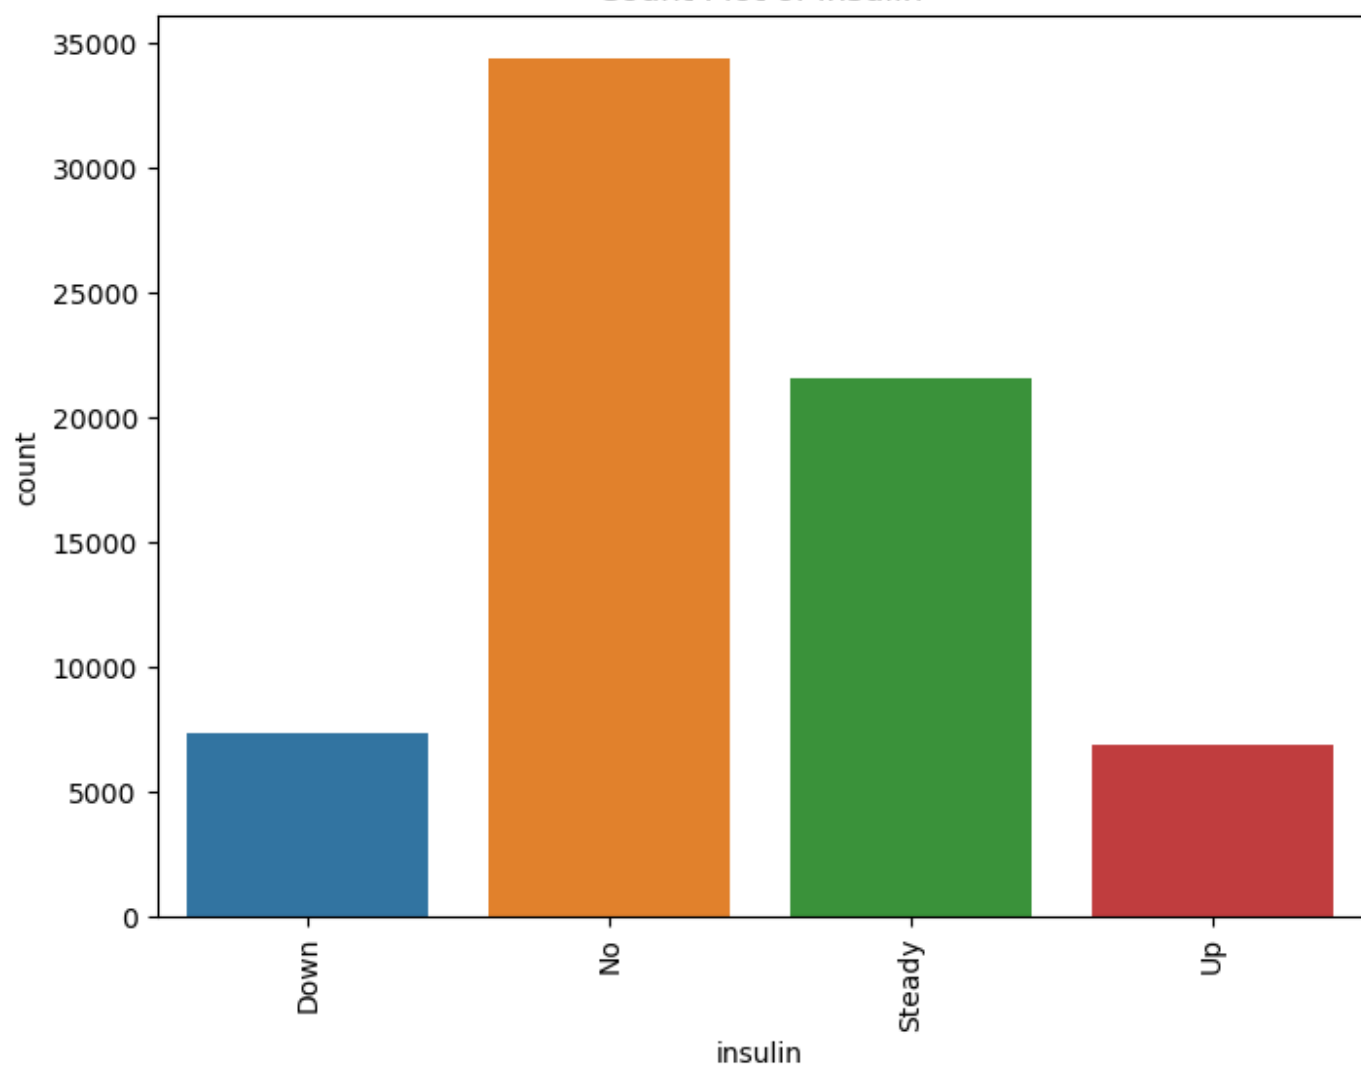
Count Plot of acarbose



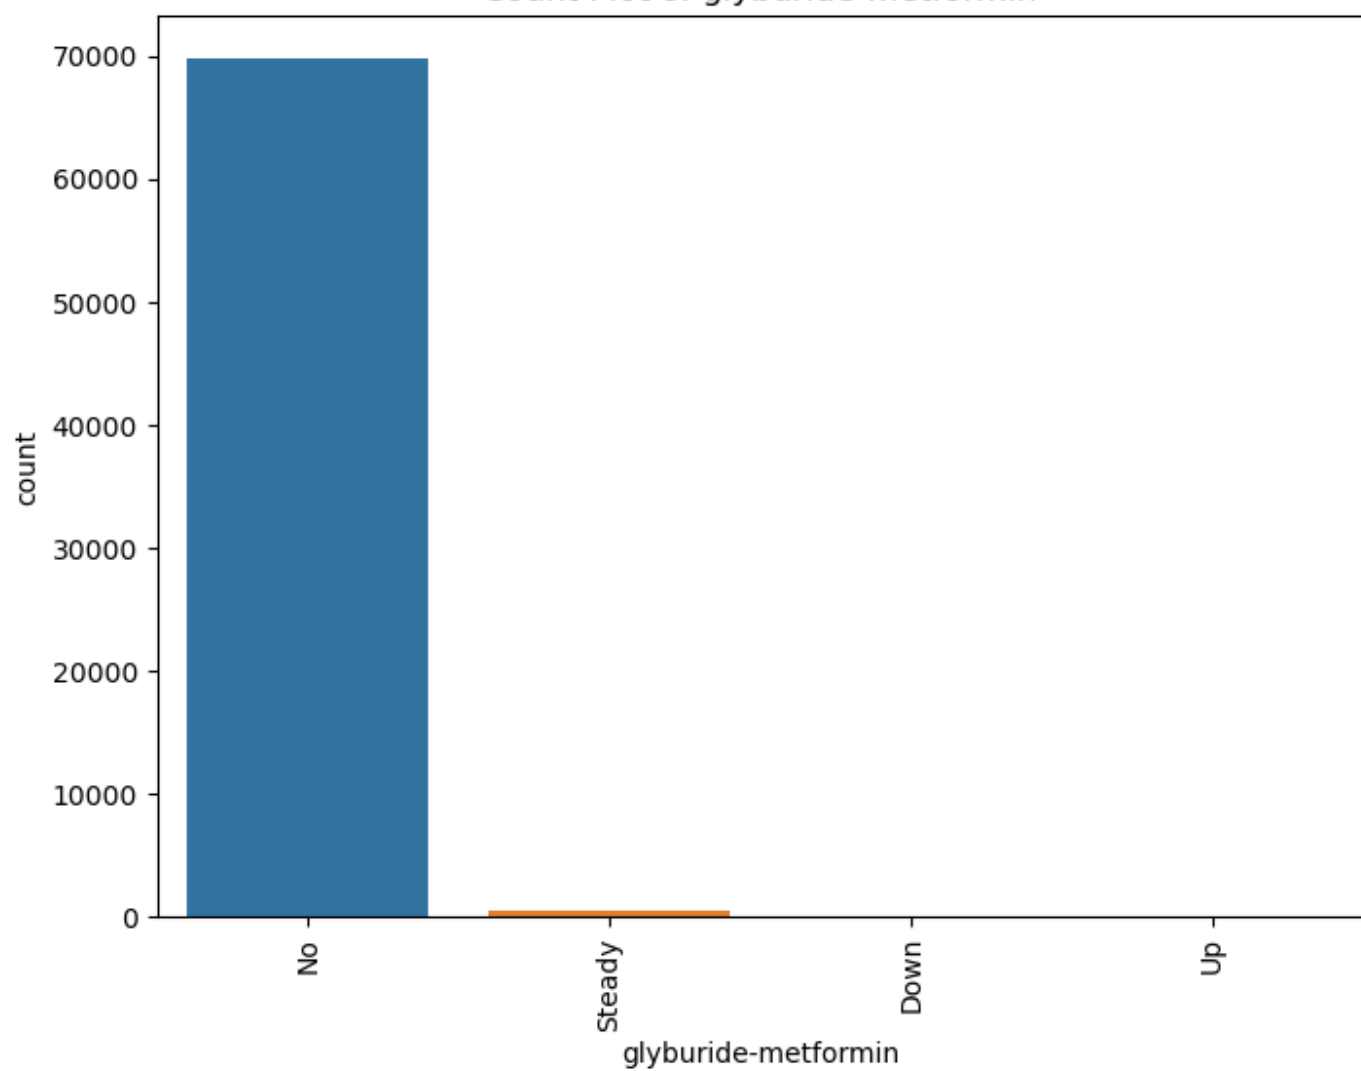
Count Plot of miglitol



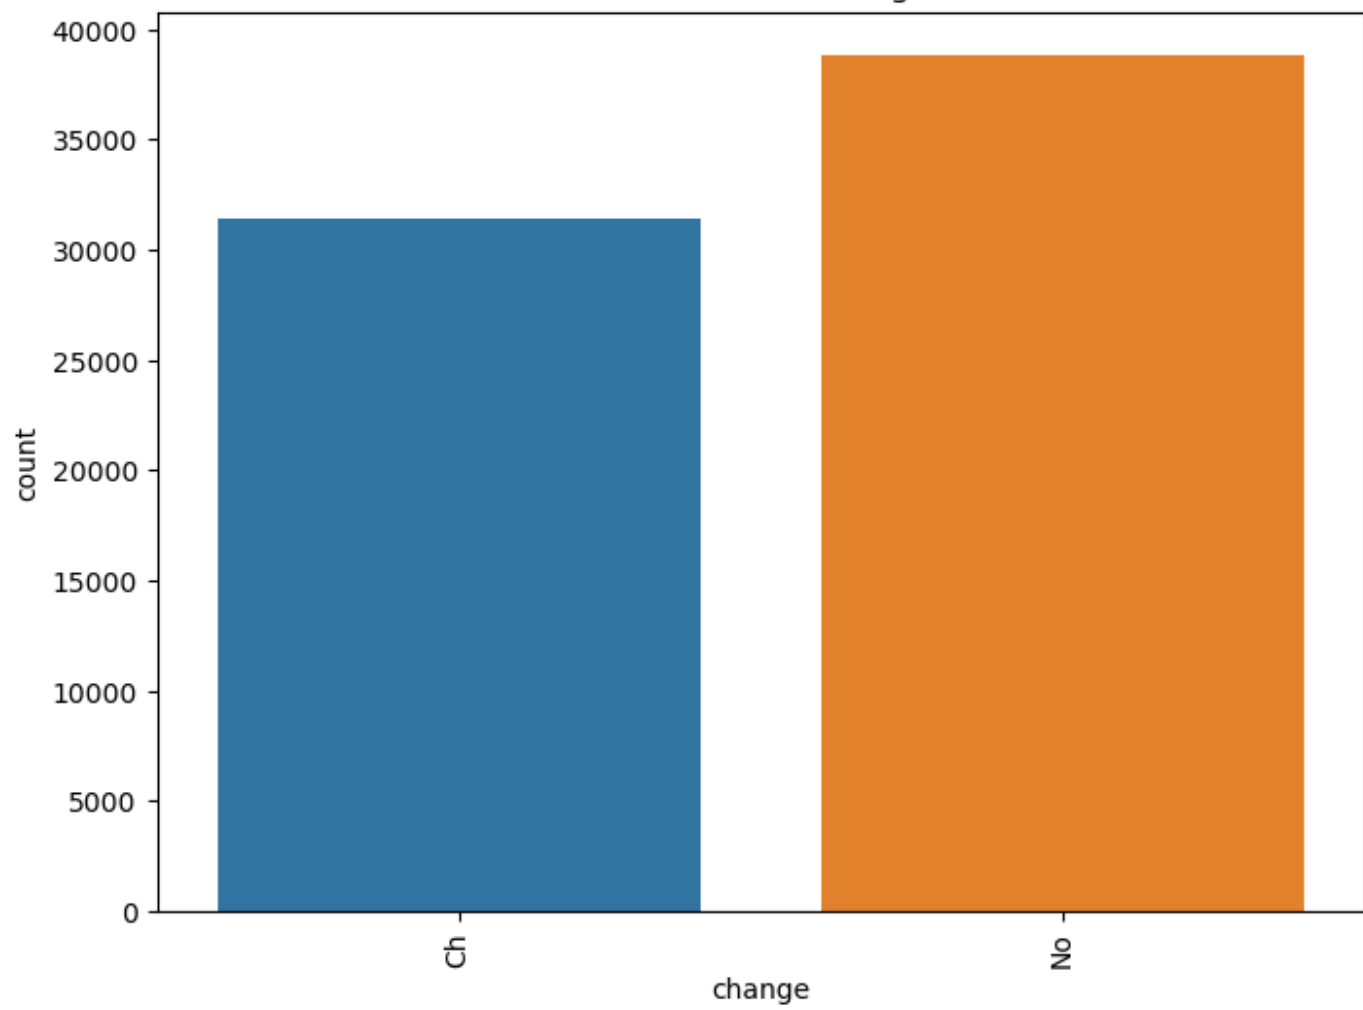
Count Plot of insulin



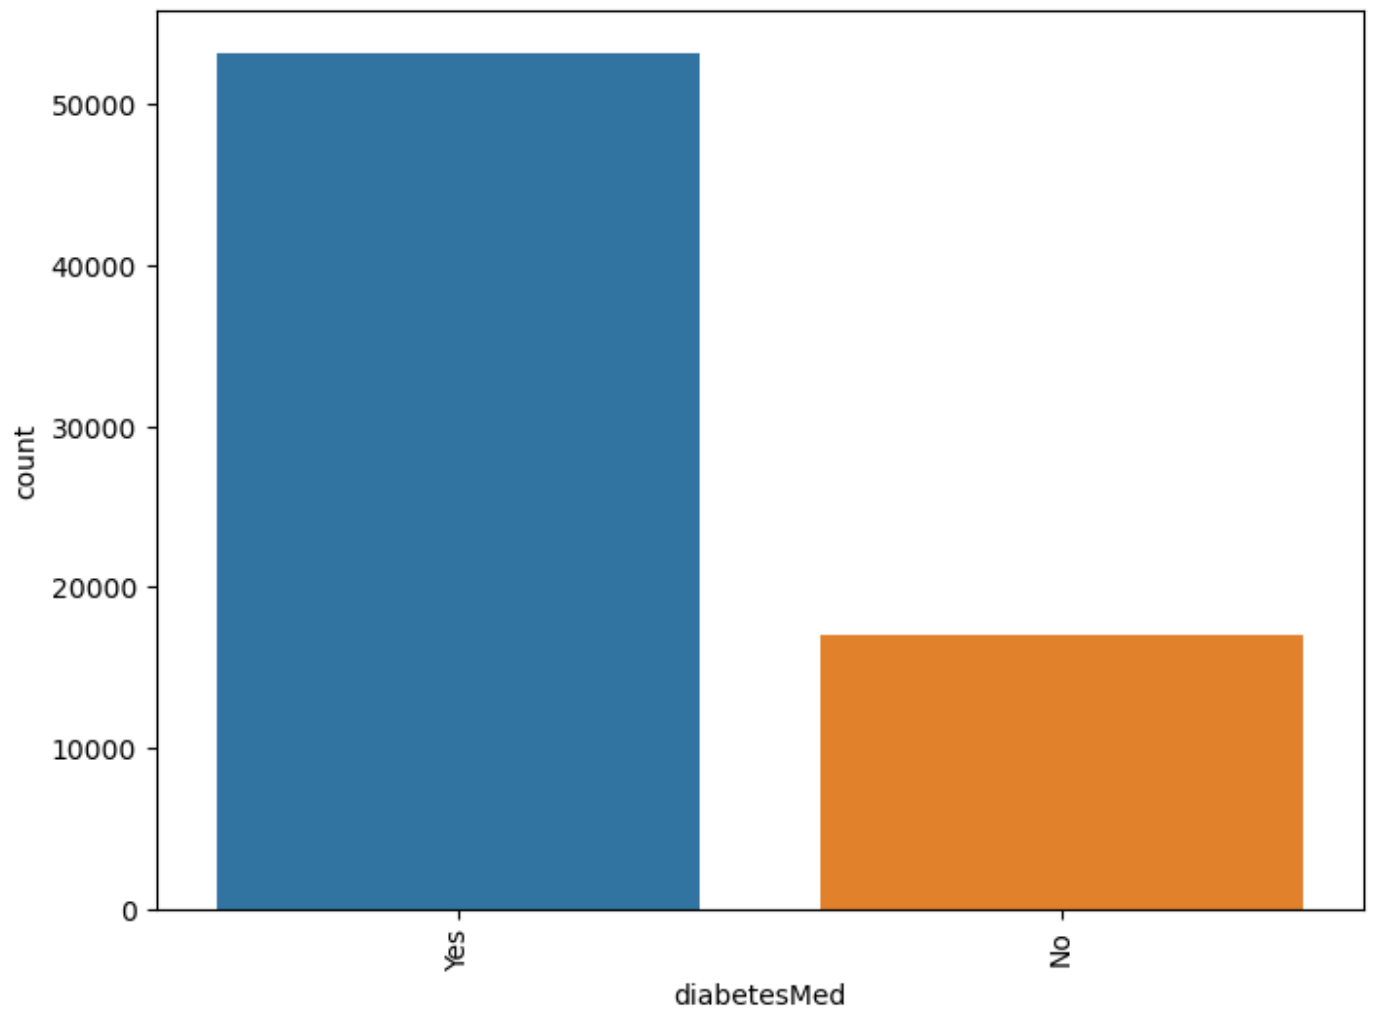
Count Plot of glyburide-metformin



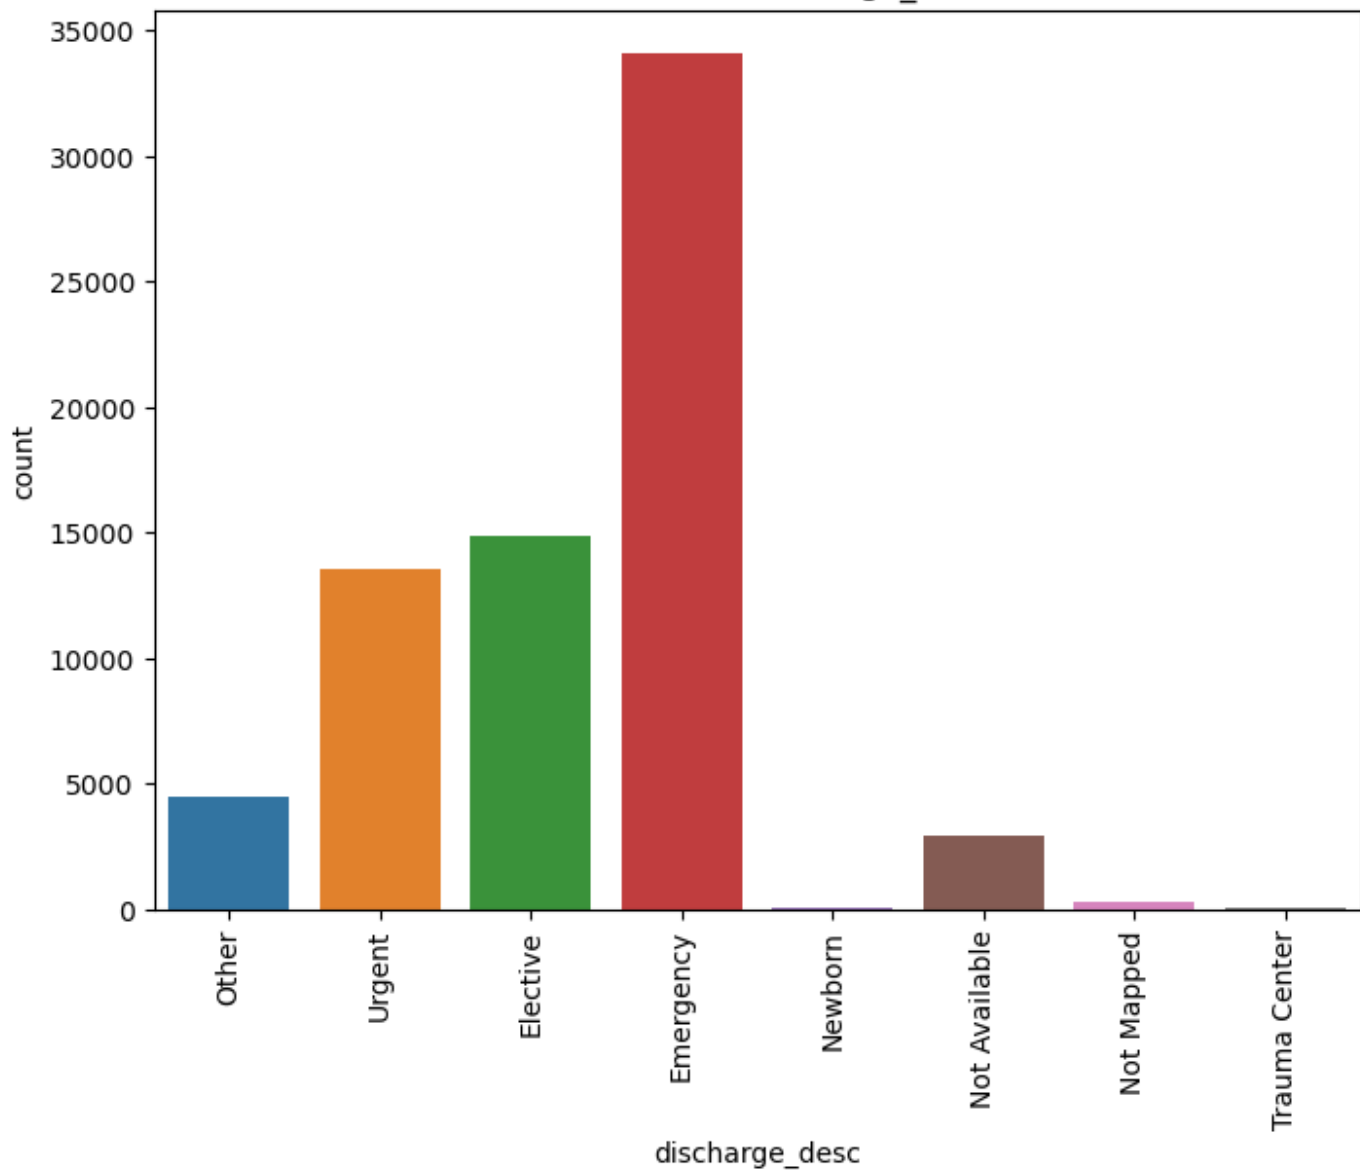
Count Plot of change



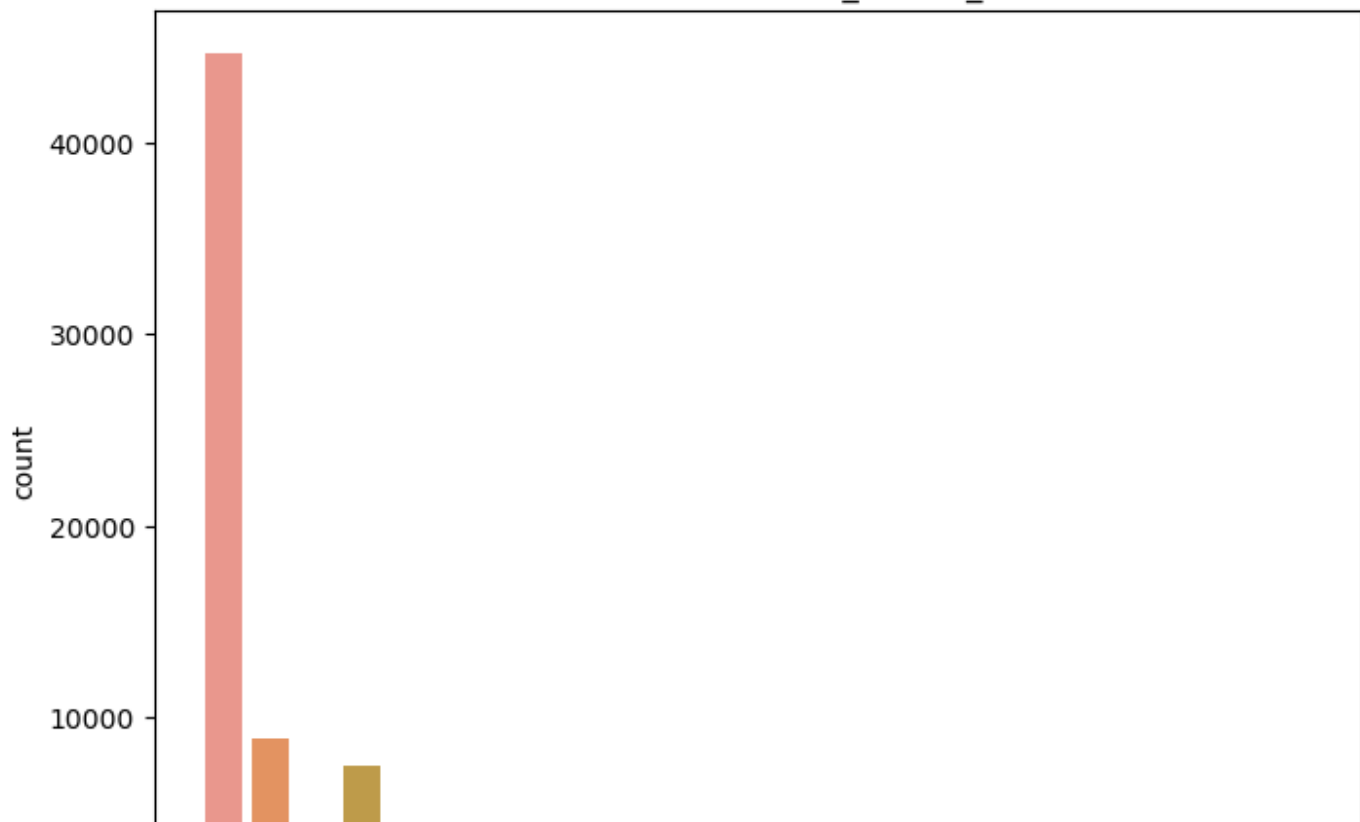
Count Plot of diabetesMed

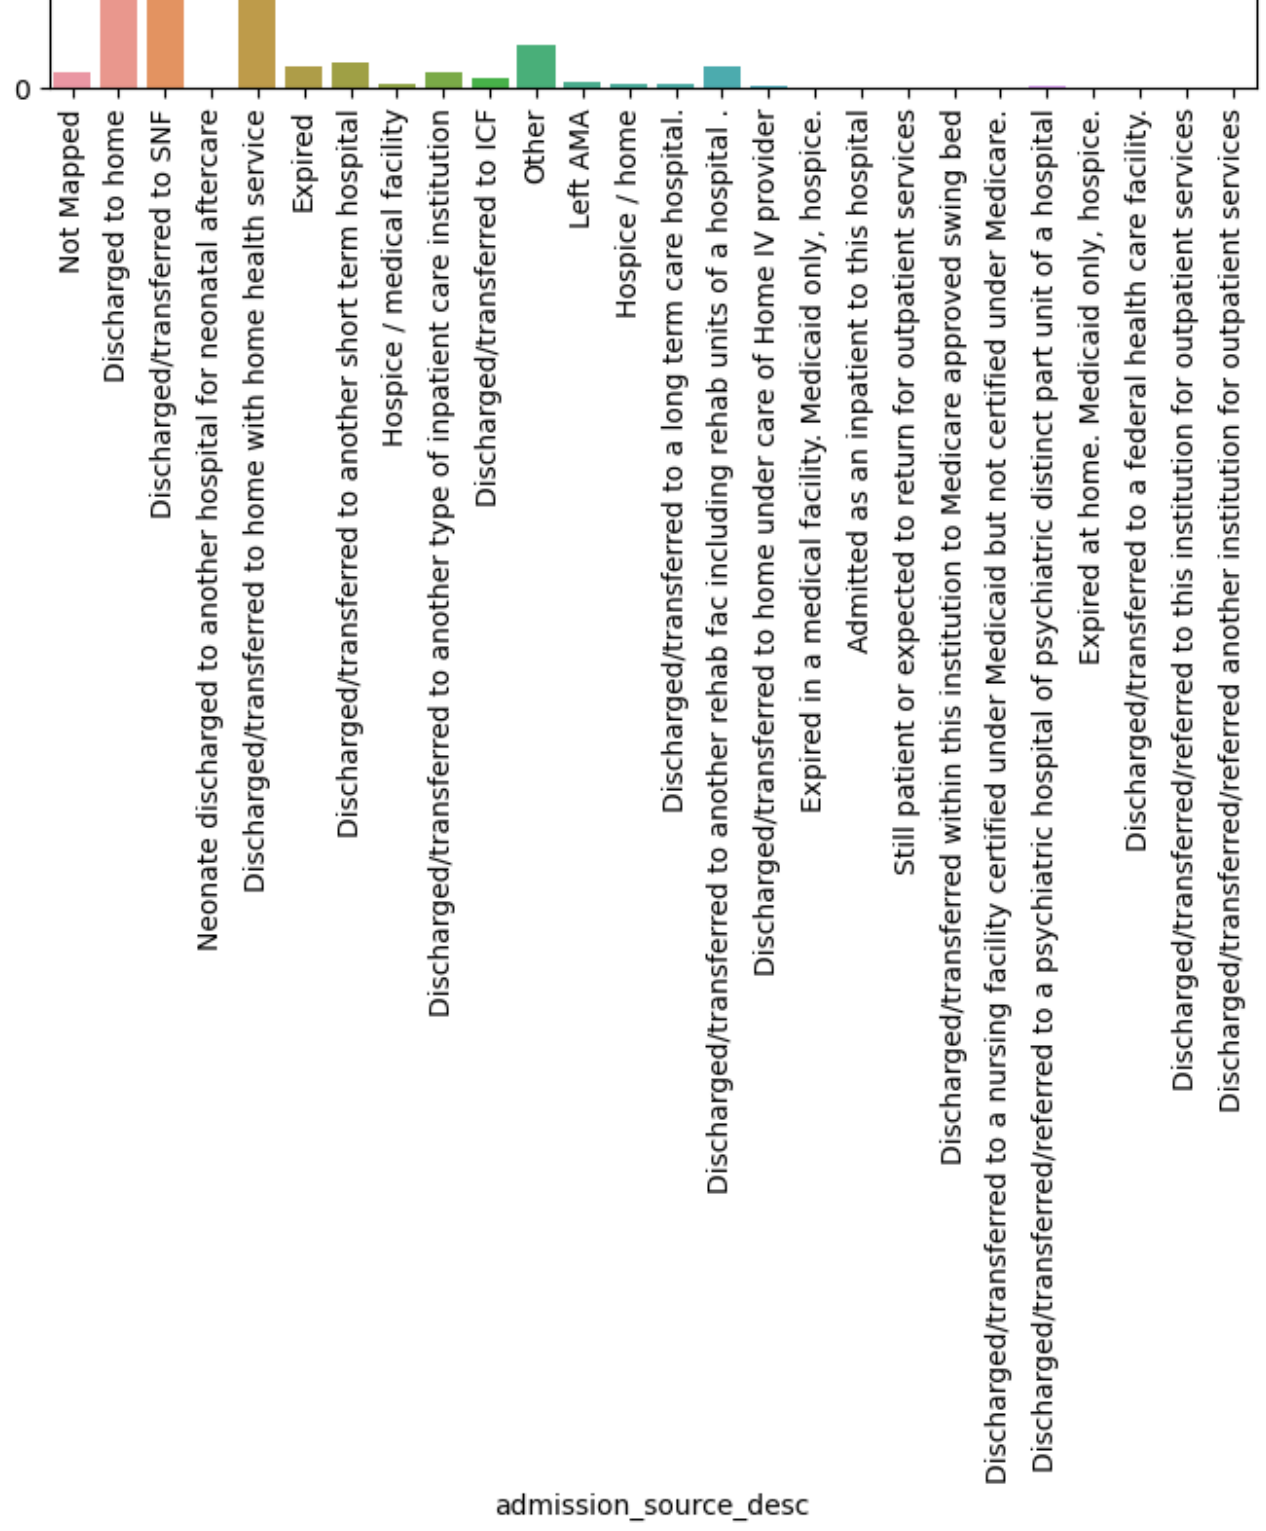


Count Plot of discharge_desc

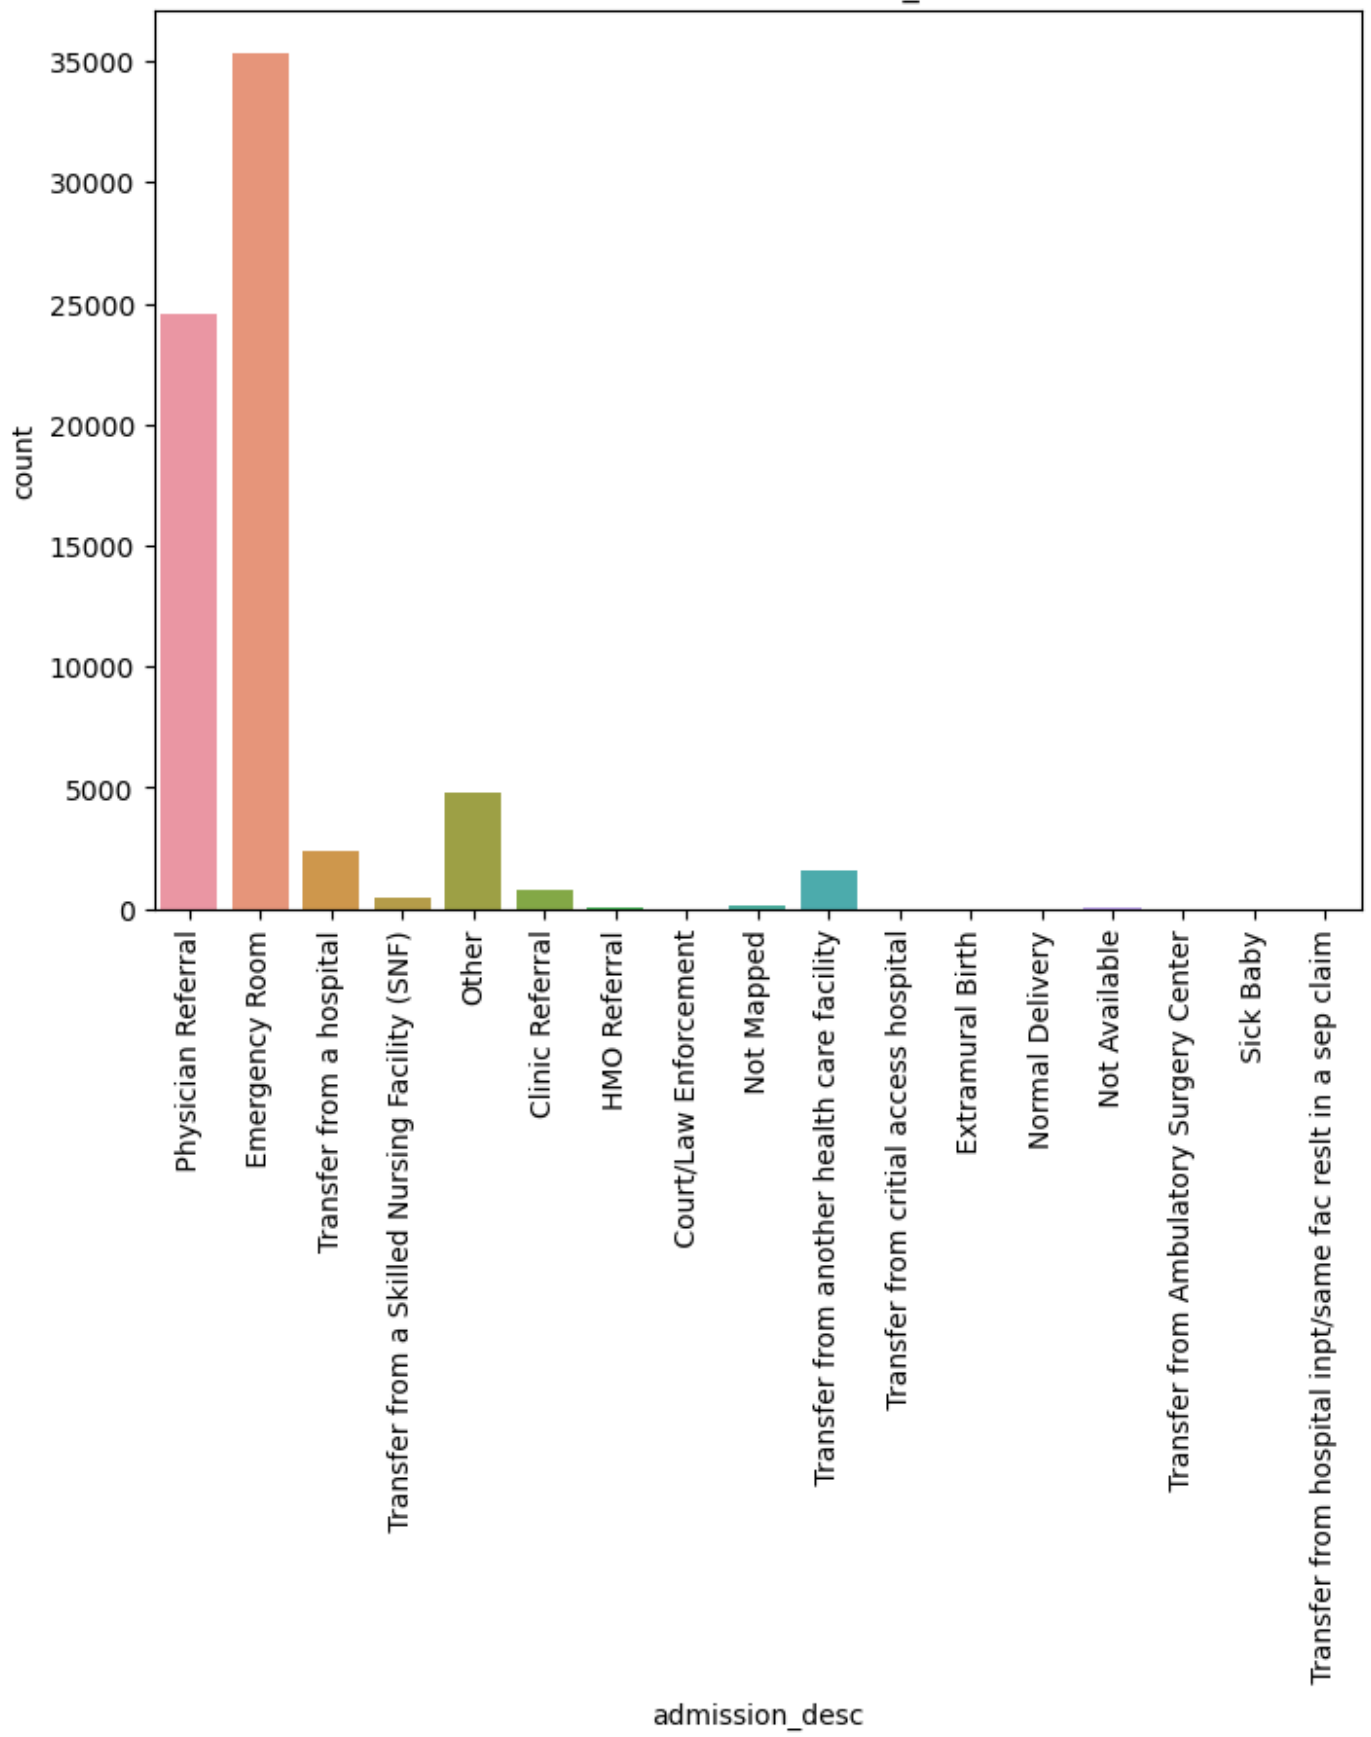


Count Plot of admission_source_desc

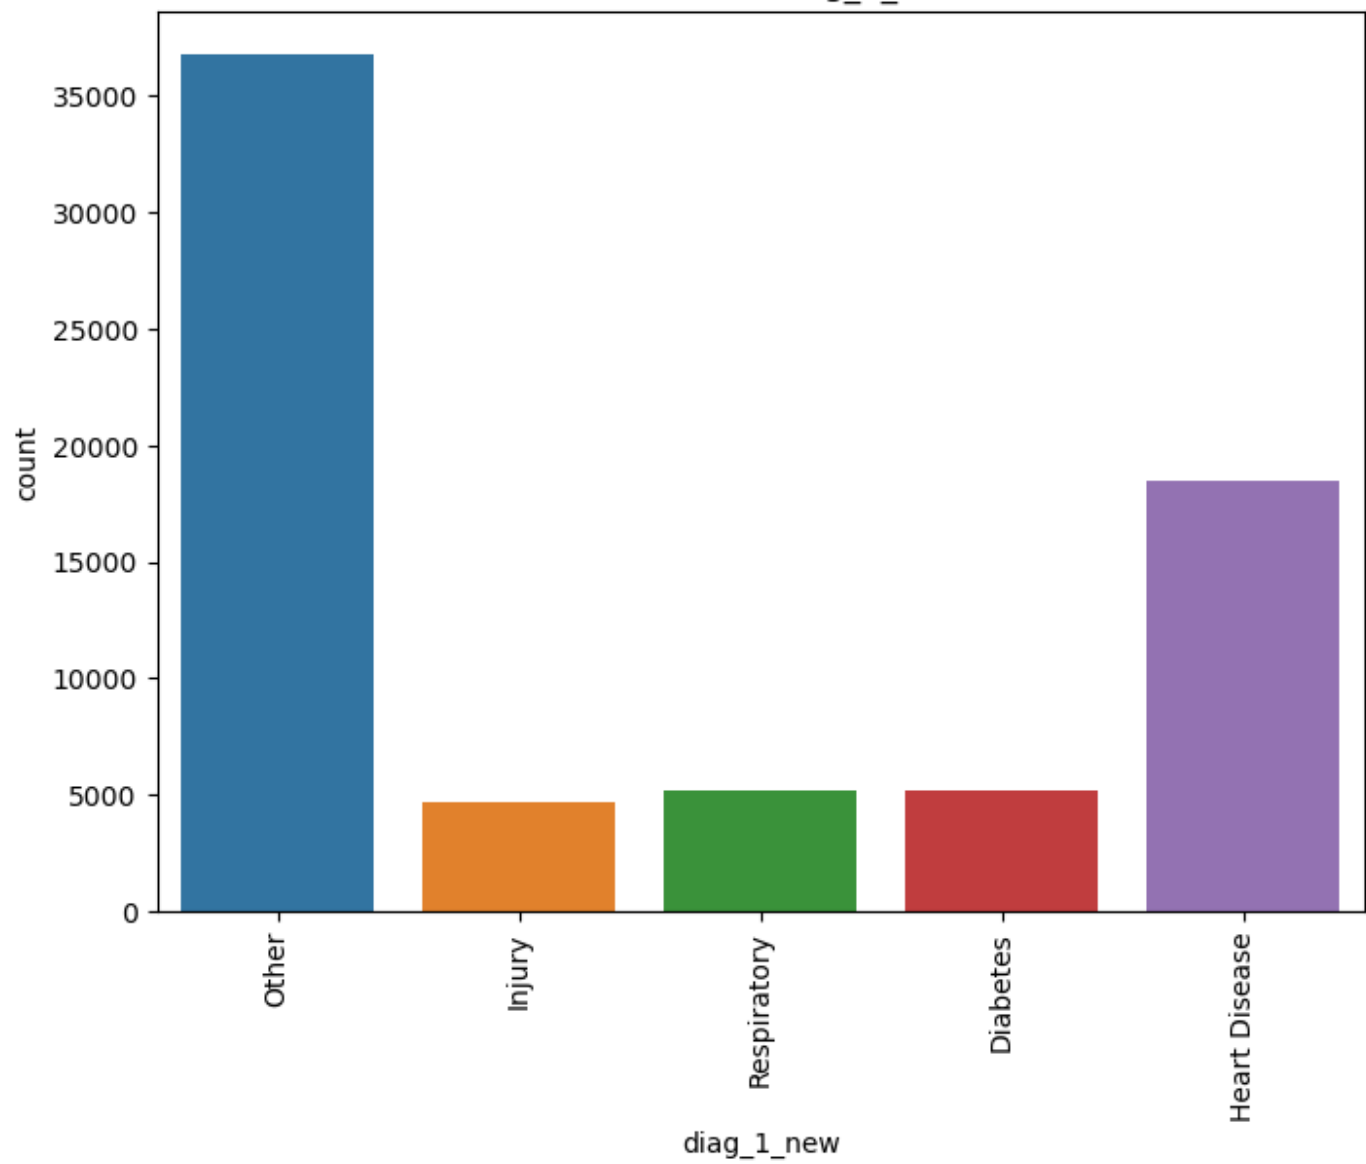




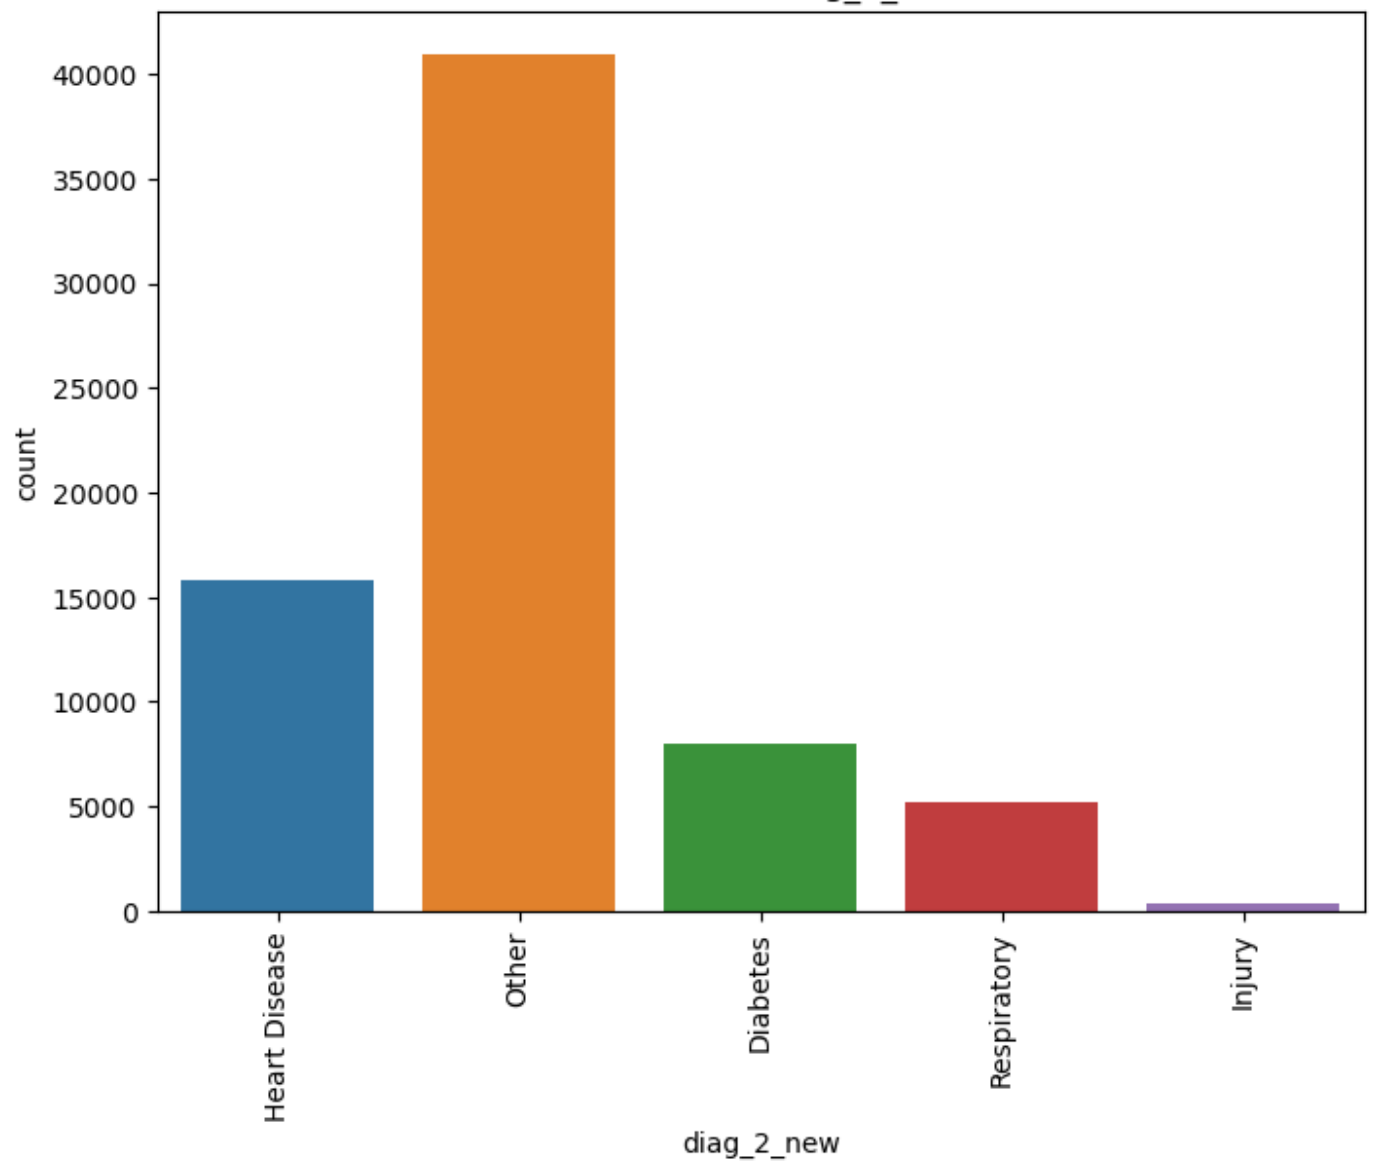
Count Plot of admission_desc

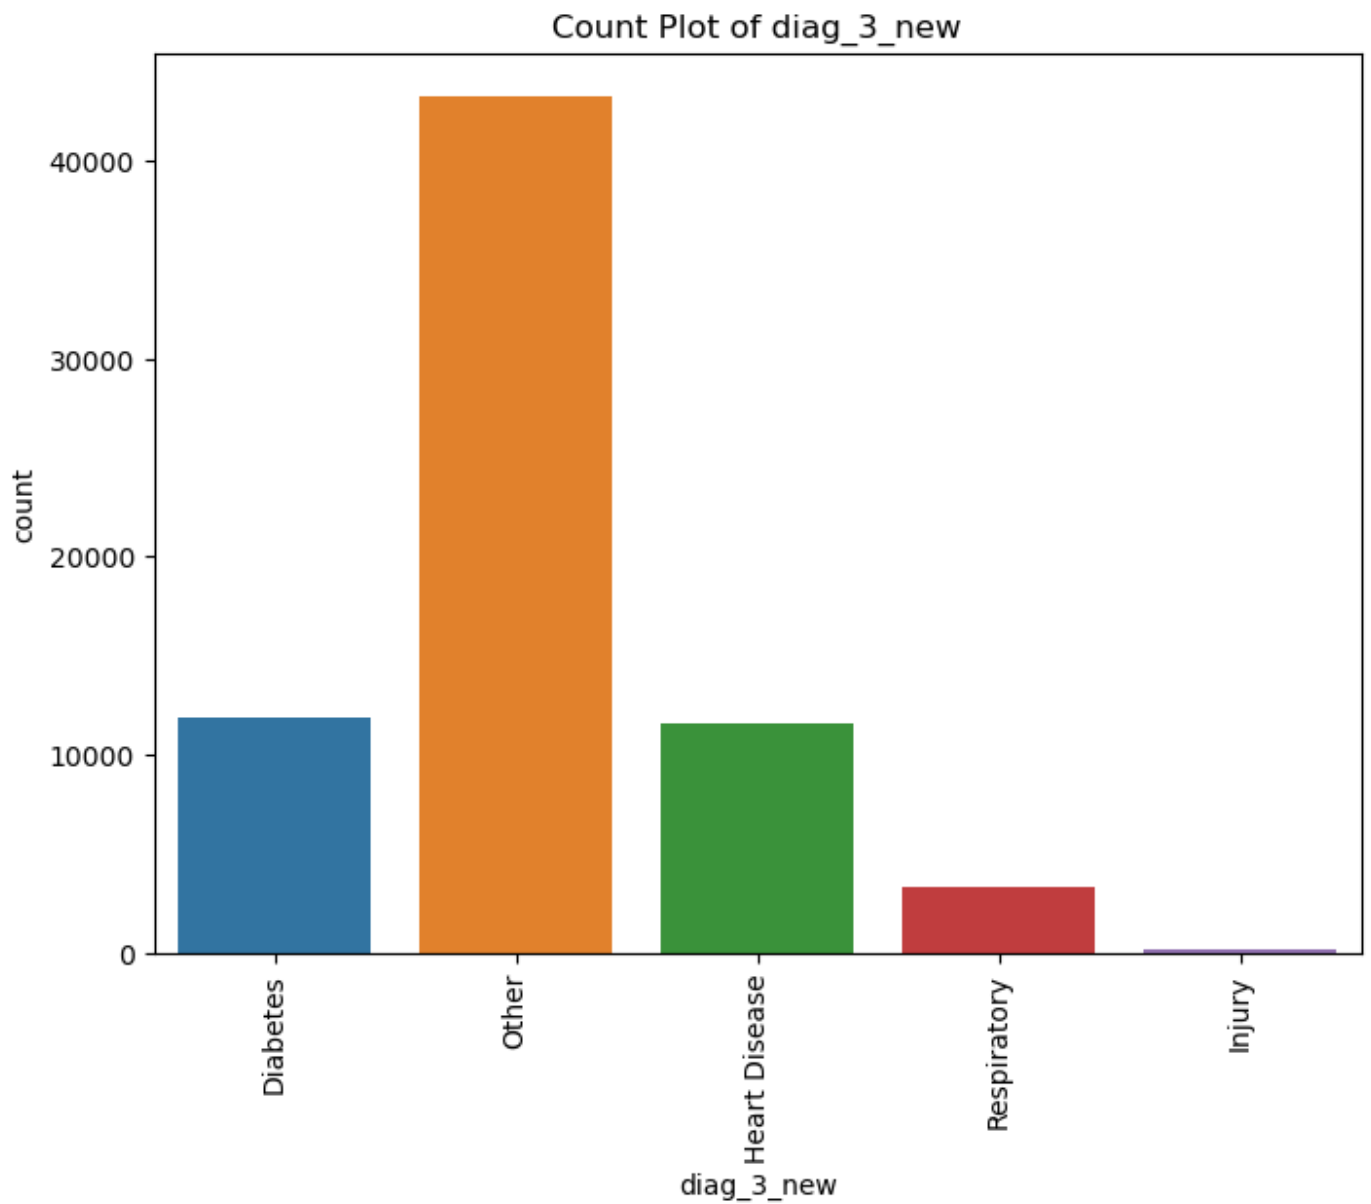


Count Plot of diag_1_new



Count Plot of diag_2_new





Data Statistics

```
In [39]: # des. stat. - numeric only  
df.describe().T
```

```
Out[39]:
```

	count	mean	std	min	25%	50%	75%	max
time_in_hospital	70256.0	4.307020	2.956366	1.0	2.0	4.0	6.0	14.0
num_lab_procedures	70256.0	42.961085	19.914422	1.0	31.0	44.0	57.0	132.0
num_procedures	70256.0	1.444873	1.761094	0.0	0.0	1.0	2.0	6.0
num_medications	70256.0	15.850675	8.264550	1.0	10.0	14.0	20.0	81.0
number_outpatient	70256.0	0.301668	1.117115	0.0	0.0	0.0	0.0	38.0
number_emergency	70256.0	0.116816	0.619245	0.0	0.0	0.0	0.0	76.0
number_inpatient	70256.0	0.274909	0.745316	0.0	0.0	0.0	0.0	15.0
number_diagnoses	70256.0	7.346462	1.887392	3.0	6.0	8.0	9.0	16.0
binary_readmitted	70256.0	0.079096	0.269891	0.0	0.0	0.0	0.0	1.0

Boxplot of Int Features

```
In [40]: fig, axes = plt.subplots(nrows=3,ncols=3,figsize = (10,10))

axes[0,0].boxplot(df['time_in_hospital'])
axes[0,0].set_title('Time in Hospital')

axes[0,1].boxplot(df['num_lab_procedures'])
axes[0,1].set_title('Number of Lab Procedures')

axes[0,2].boxplot(df['num_procedures'])
axes[0,2].set_title('Number of Procedures')

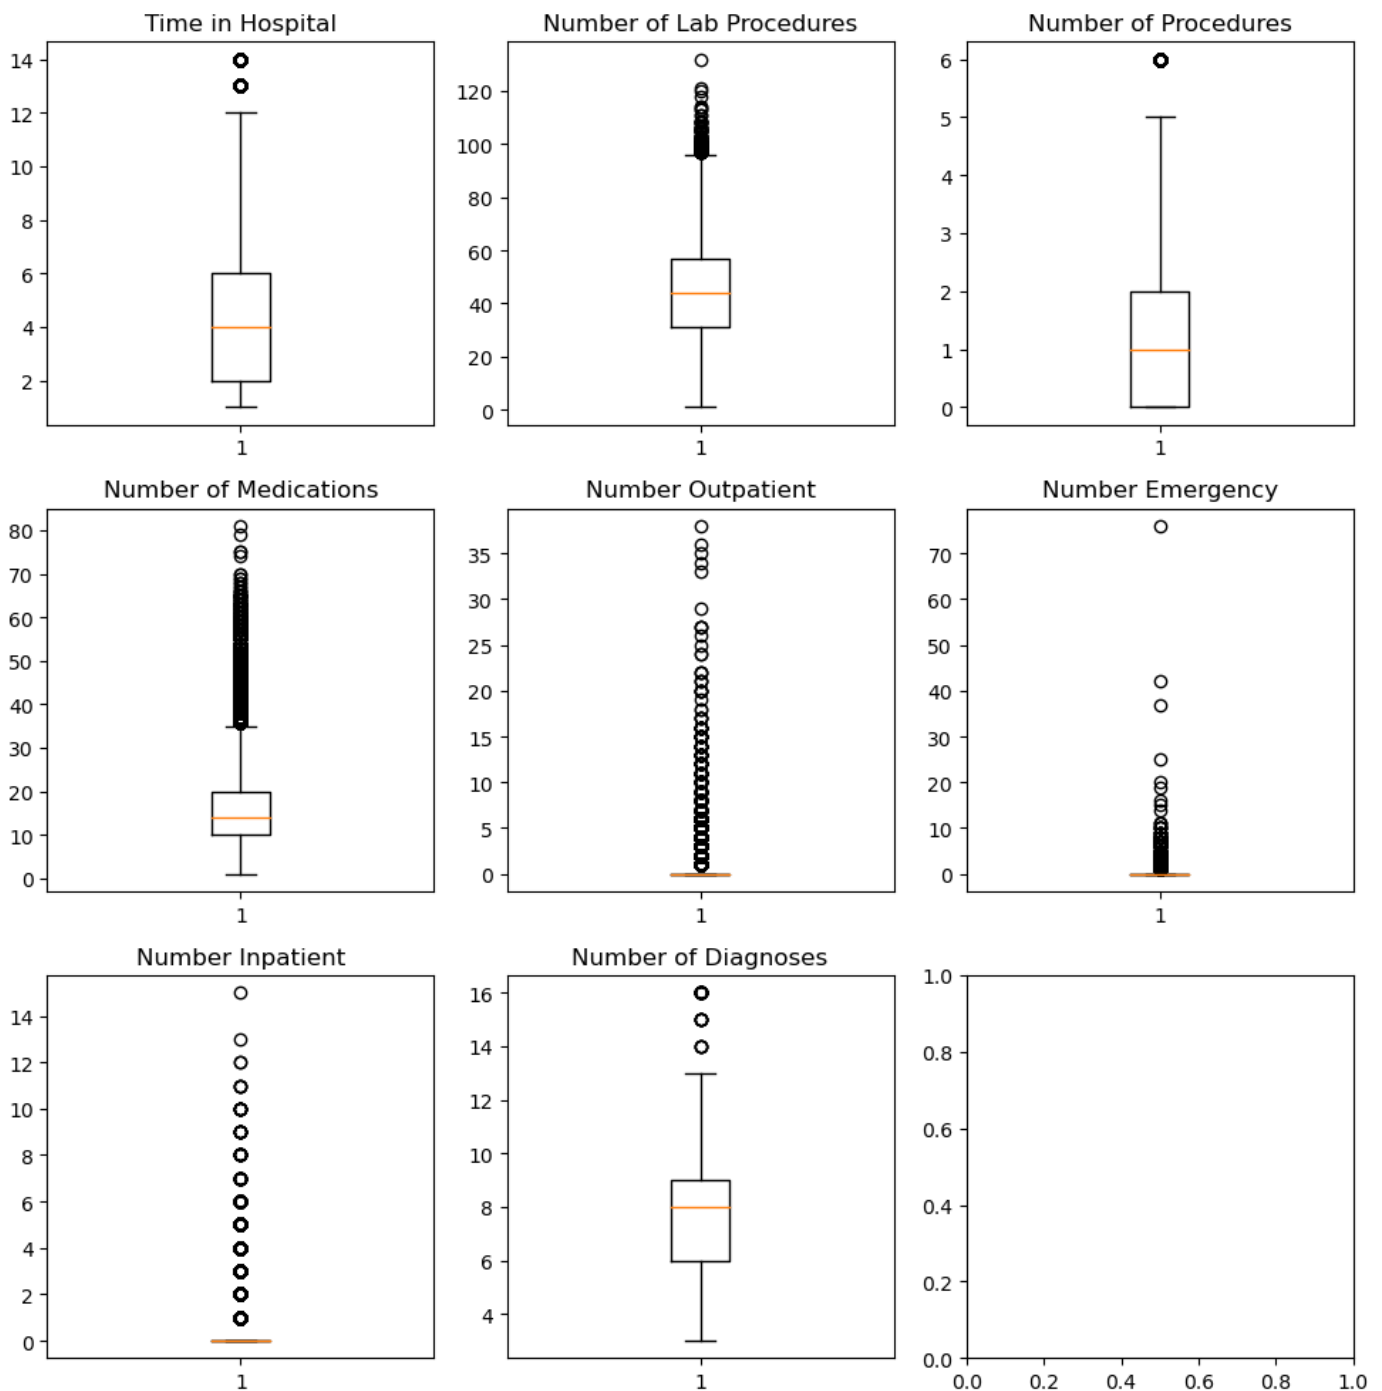
axes[1,0].boxplot(df['num_medications'])
axes[1,0].set_title('Number of Medications')

axes[1,1].boxplot(df['number_outpatient'])
axes[1,1].set_title('Number Outpatient')

axes[1,2].boxplot(df['number_emergency'])
axes[1,2].set_title('Number Emergency')

axes[2,0].boxplot(df['number_inpatient'])
axes[2,0].set_title('Number Inpatient')

axes[2,1].boxplot(df['number_diagnoses'])
axes[2,1].set_title('Number of Diagnoses')
plt.tight_layout()
```



Histogram of Int Features

```
In [41]: fig, axes = plt.subplots(nrows=3,ncols=3,figsize = (10,10))

axes[0,0].hist(df['time_in_hospital'], bins = 12)
axes[0,0].set_title('Time in Hospital')

axes[0,1].hist(df['num_lab_procedures'], bins = 20)
axes[0,1].set_title('Number of Lab Procedures')

axes[0,2].hist(df['num_procedures'], bins = 5)
axes[0,2].set_title('Number of Procedures')

axes[1,0].hist(df['num_medications'], bins = 12)
axes[1,0].set_title('Number of Medications')

axes[1,1].hist(df['number_outpatient'], bins = 25)
axes[1,1].set_title('Number Outpatient')
```

```

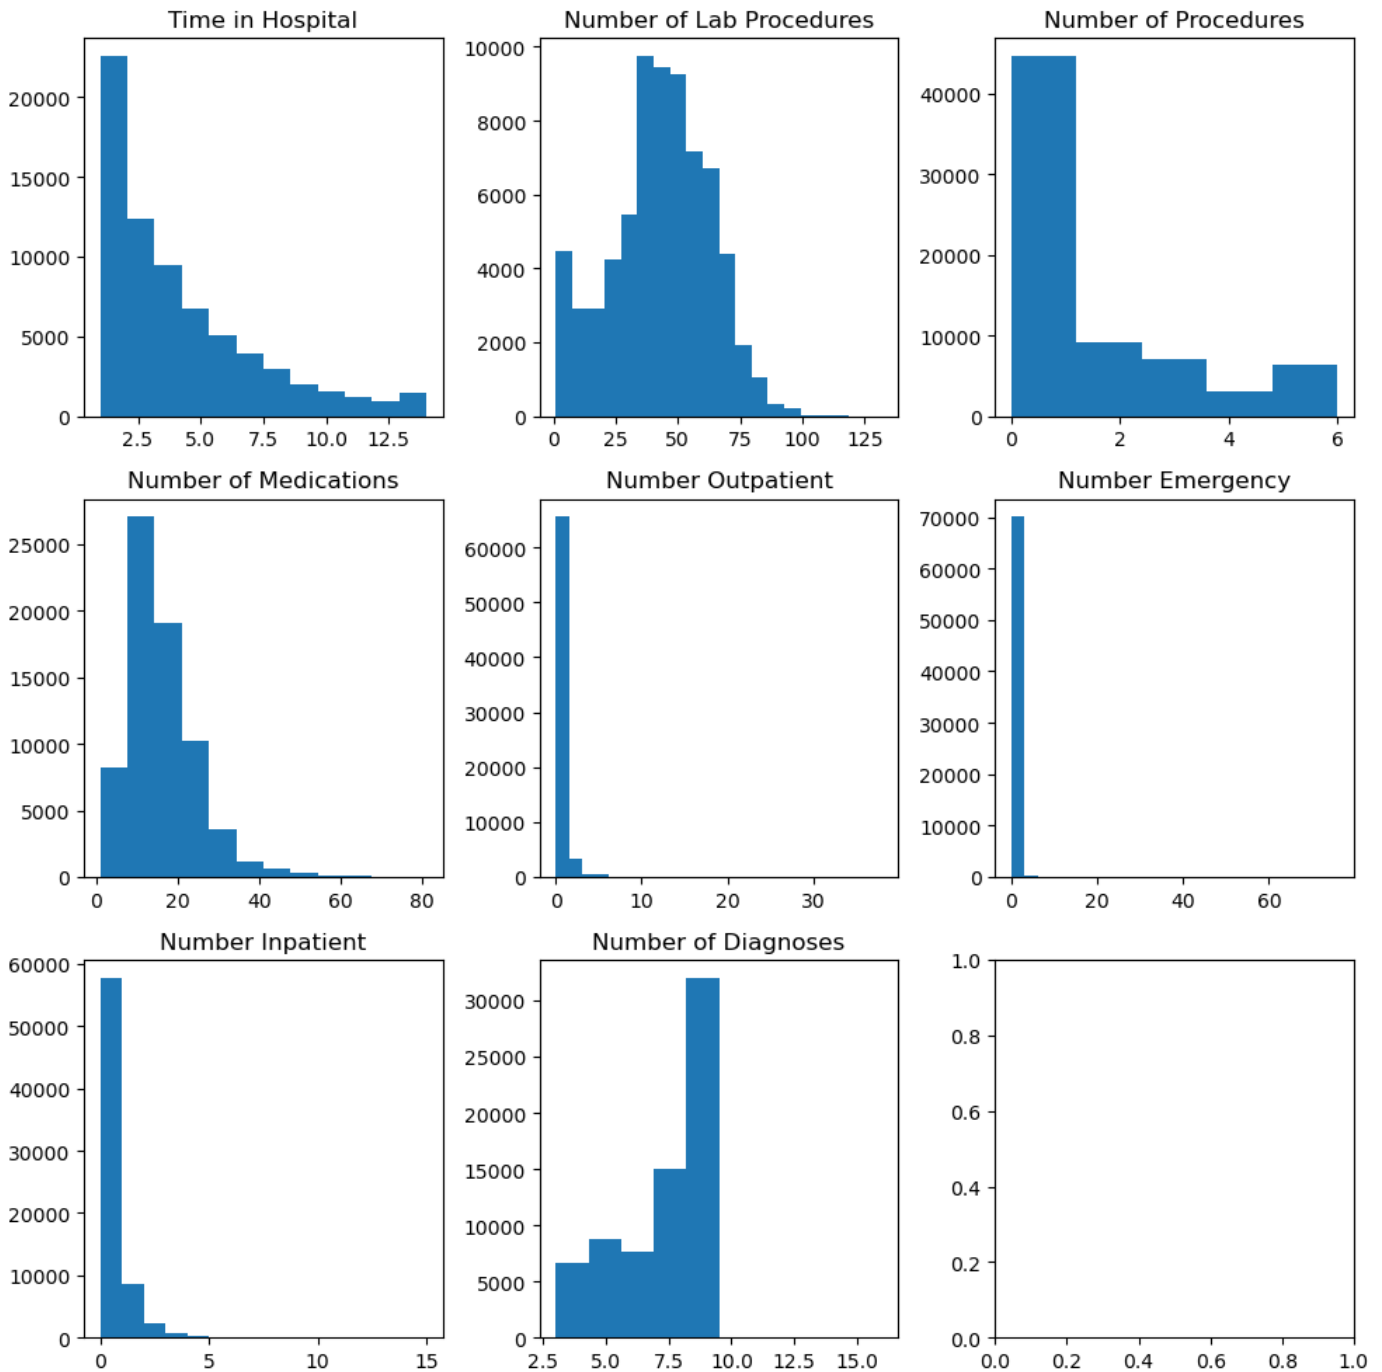
axes[1,2].hist(df['number_emergency'], bins = 25)
axes[1,2].set_title('Number Emergency')

axes[2,0].hist(df['number_inpatient'], bins = 15)
axes[2,0].set_title('Number Inpatient')

axes[2,1].hist(df['number_diagnoses'], bins = 10)
axes[2,1].set_title('Number of Diagnoses')

plt.tight_layout()

```



Dummy and Scaling

```

In [42]: from sklearn.preprocessing import StandardScaler
         from sklearn.model_selection import train_test_split

```

```

X = df.drop(['binary_readmitted'], axis = 1)
y = df['binary_readmitted']

X_train, X_test, y_train, y_test = train_test_split(X,y,train_size = .8, random_state=12)

object_columns= X_train.select_dtypes(include = 'object')
X_train_encoded = pd.get_dummies(X_train, columns= object_columns.columns, drop_first =
X_test_encoded = pd.get_dummies(X_test, columns = object_columns.columns, drop_first = T

## Account for missing columns in X_test that are in X_train
missing_columns = set(X_train_encoded.columns) - set(X_test_encoded.columns)

# Add missing columns to X_test and fill them with zeros
for col in missing_columns:
    X_test_encoded[col] = 0

int_columns = X_train_encoded.select_dtypes(include = 'int')
scaler = StandardScaler()
X_train_encoded[int_columns.columns] = scaler.fit_transform(X_train_encoded[int_columns.
X_test_encoded[int_columns.columns] = scaler.transform(X_test_encoded[int_columns.column

```

```

In [43]: correl = df.corr()
plt.figure(figsize = (12,12))
sns.heatmap(correl, cmap = 'mako',linewidth = .2, linecolor= 'white', annot=True)
plt.title('Correlation Matrix of Numerical Features')
plt.show()

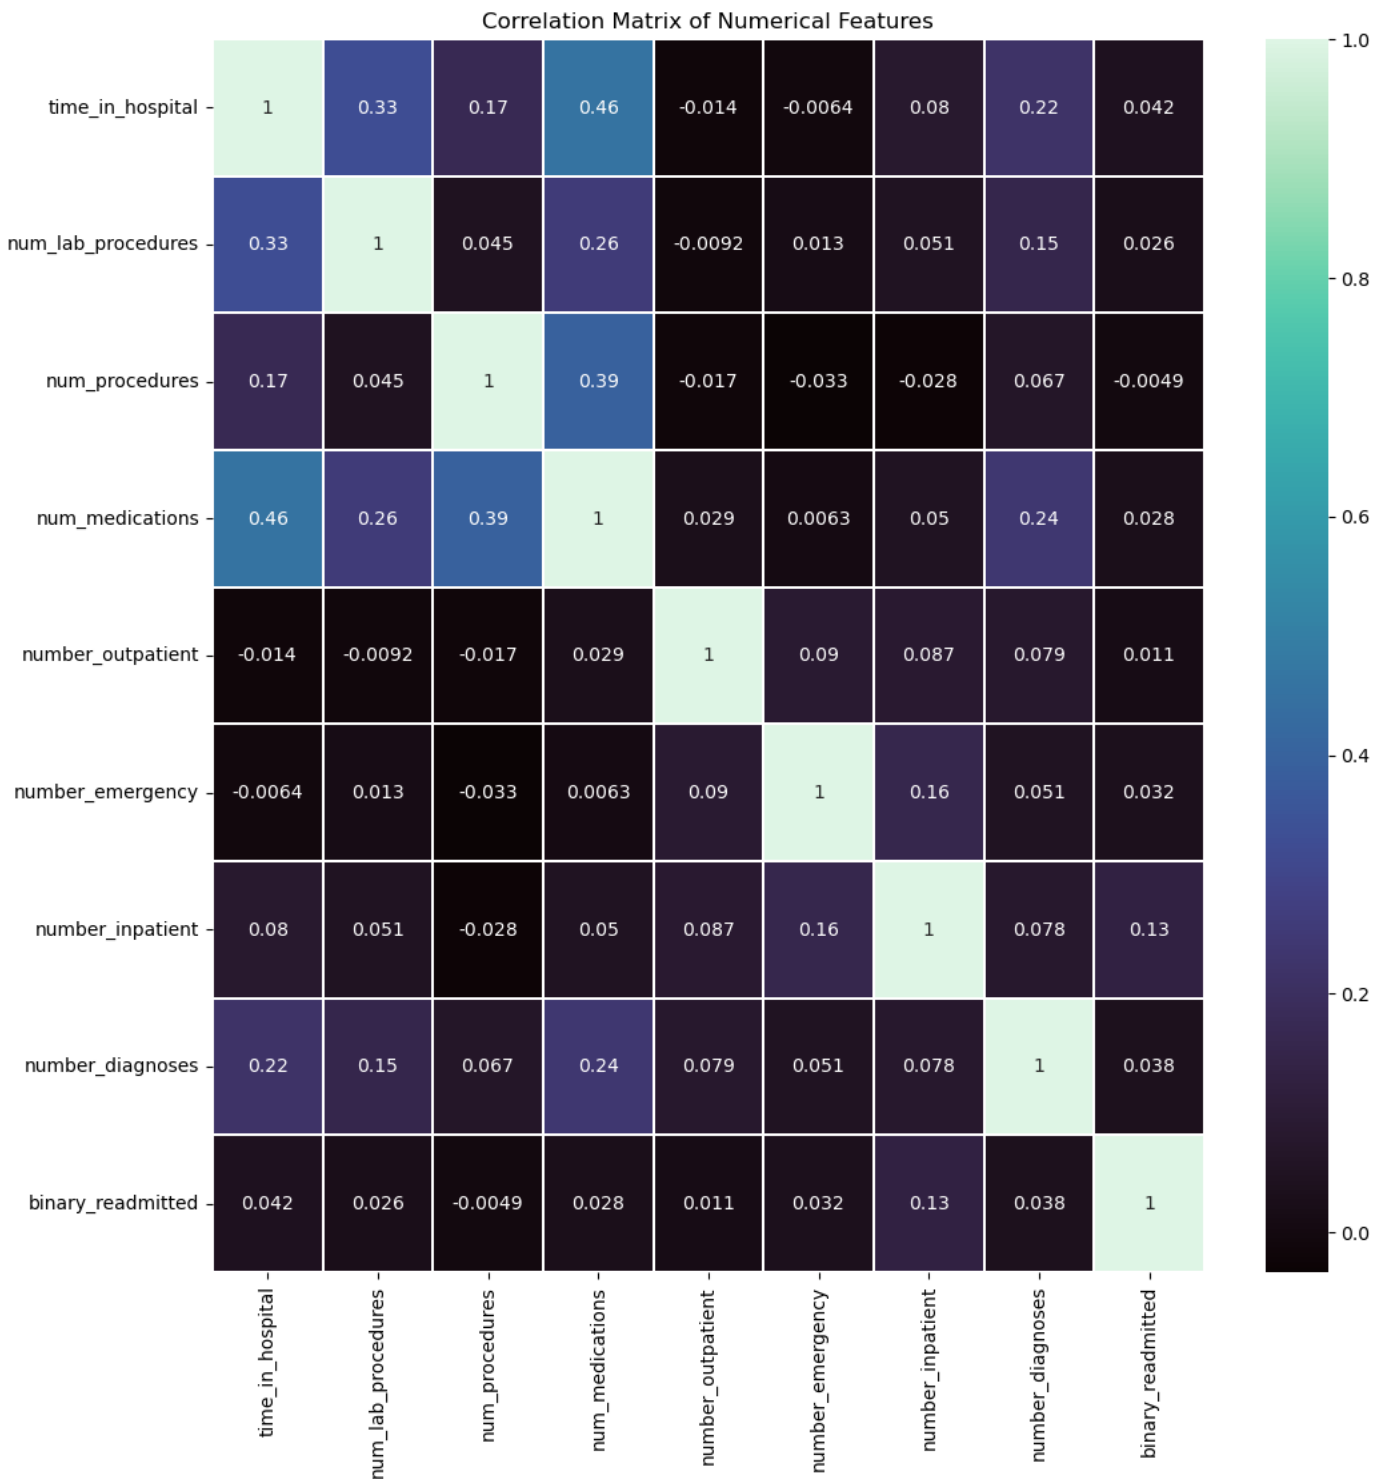
```

C:\Users\Joey\AppData\Local\Temp\ipykernel_4540\1708422235.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```

    correl = df.corr()

```

Lasso / Logistic Regression + Grid Search

```
In [44]: from sklearn.linear_model import LassoCV, LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc

# Define range of alpha values
alphas = np.logspace(-2, 2, 11)

# Create the LassoCV model with CV alpha selection
lasso_model = LassoCV(alphas=alphas, cv=5)

# Fit the LassoCV model to train
```

```

lasso_model.fit(X_train_encoded, y_train)

# Get the best alpha value
best_alpha = lasso_model.alpha_

# Apply the same one-hot encoding to the test data
object_columns_test = X_test.select_dtypes(include='object')
X_test_encoded = pd.get_dummies(X_test, columns=object_columns_test.columns, drop_first=

# Ensure the test data has the same columns as the training data
missing_columns = set(X_train_encoded.columns) - set(X_test_encoded.columns)
for column in missing_columns:
    X_test_encoded[column] = 0 # Add missing columns with all zeros

# Reorder columns in X_test_encoded to match X_train_encoded
X_test_encoded = X_test_encoded[X_train_encoded.columns]

# Create the Lasso logistic regression model with the best alpha value
lasso_logistic_best = LogisticRegression(penalty='l1', solver='saga', C=1 / best_alpha,

# Fit the model to the training data
lasso_logistic_best.fit(X_train_encoded, y_train)

```

Out[44]:

```

▼ LogisticRegression
LogisticRegression(C=100.0, max_iter=10000, penalty='l1', random_state=12,
                    solver='saga')

```

Predictions, Coefficients, Classification Report, ROC, and Confusion Matrix

In [45]:

```

# Make predictions on the test set
y_pred_prob = lasso_logistic_best.predict_proba(X_test_encoded)[: , 1]
y_pred = (y_pred_prob >= 0.65).astype(int) # Change threshold ***

# Calculate accuracy on the test set
test_accuracy = accuracy_score(y_test, y_pred)

# Display features with non-zero coefficients
coefficients = lasso_logistic_best.coef_
feature_names = X_train_encoded.columns
coef_df = pd.DataFrame({'Feature': feature_names, 'Coefficient': coefficients[0]})
non_zero_coef_df = coef_df[coef_df['Coefficient'] != 0]

print("Features with Non-Zero Coefficients:")
print(non_zero_coef_df)

print(f"Best Alpha (Lambda): {best_alpha:.4f}")
print(f"Test Set Accuracy: {test_accuracy:.4f}")

# threshold
user_defined_threshold = 0.65

# Calculate classification metrics with threshold
y_pred_custom_threshold = (y_pred_prob >= user_defined_threshold).astype(int)

classification_report_custom_threshold = classification_report(y_test, y_pred_custom_thr
confusion_matrix_custom_threshold = confusion_matrix(y_test, y_pred_custom_threshold)

print(f"User-Defined Threshold: {user_defined_threshold:.2f}")
print("Classification Report with Custom Threshold:")
print(classification_report_custom_threshold)

```

```

custom_label = ['Not Readmitted', 'Readmitted']
# Plot confusion matrix with threshold
plt.figure(figsize=(8, 5))
sns.heatmap(confusion_matrix_custom_threshold, annot=True, fmt="d", cmap="Blues", cbar=F
            xticklabels=custom_label, yticklabels=custom_label)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Lasso Prediction Matrix")
plt.show()

```

Features with Non-Zero Coefficients:

	Feature	Coefficient
0	time_in_hospital	0.022256
1	num_lab_procedures	0.037550
2	num_procedures	-0.024361
3	num_medications	0.030043
4	number_outpatient	0.008690
..
209	diag_2_new_Respiratory	-0.050334
210	diag_3_new_Heart Disease	-0.031457
211	diag_3_new_Injury	-0.004772
212	diag_3_new_Other	-0.034090
213	diag_3_new_Respiratory	0.005635

[214 rows x 2 columns]

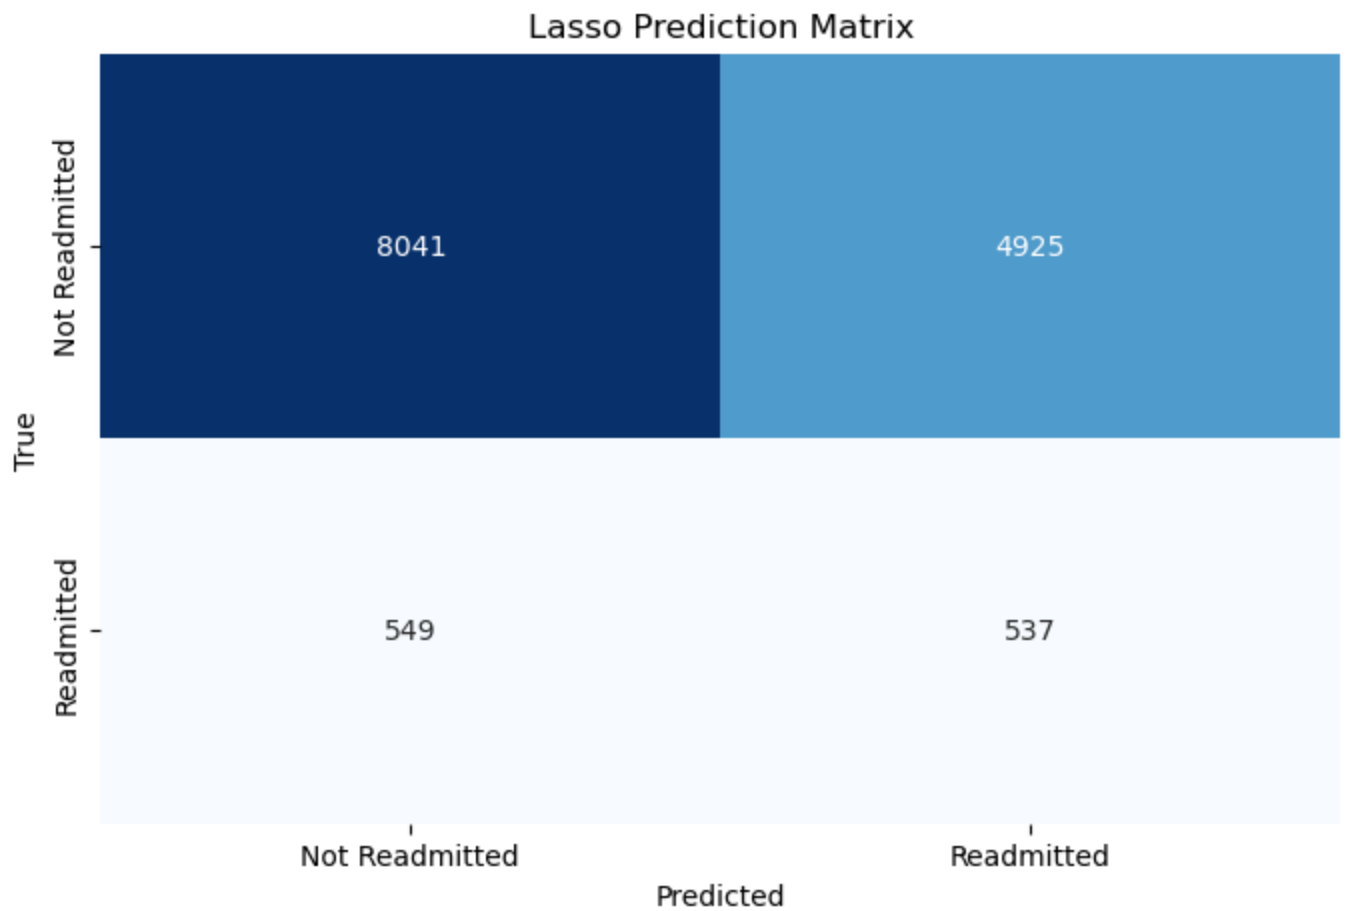
Best Alpha (Lambda): 0.0100

Test Set Accuracy: 0.6104

User-Defined Threshold: 0.65

Classification Report with Custom Threshold:

	precision	recall	f1-score	support
0	0.94	0.62	0.75	12966
1	0.10	0.49	0.16	1086
accuracy			0.61	14052
macro avg	0.52	0.56	0.46	14052
weighted avg	0.87	0.61	0.70	14052



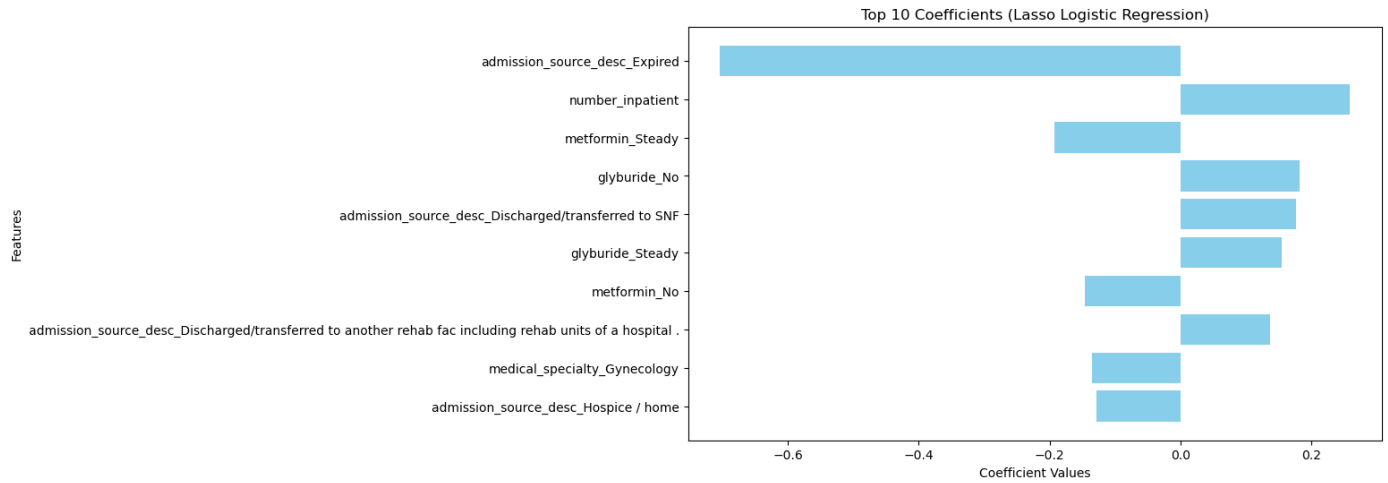
```
In [46]: # Get the coefficients and corresponding feature names
coefficients = lasso_logistic_best.coef_
feature_names = X_train_encoded.columns.tolist()
coef_df = pd.DataFrame({'Feature': feature_names, 'Coefficient': coefficients[0]})

# Sort the coefficients by absolute value in descending order
sorted_coef_df = coef_df.reindex(coef_df['Coefficient'].abs().sort_values(ascending=False))

# Select the top 10 coefficients
top_10_coef = sorted_coef_df.head(10)
print(top_10_coef)

# Create a bar plot for the top 10 coefficients
plt.figure(figsize=(10, 6))
plt.barh(top_10_coef['Feature'], top_10_coef['Coefficient'], color='skyblue')
plt.xlabel('Coefficient Values')
plt.ylabel('Features')
plt.title('Top 10 Coefficients (Lasso Logistic Regression)')
plt.gca().invert_yaxis()
plt.show()
```

	Feature	Coefficient
176	admission_source_desc_Expired	-0.704329
6	number_inpatient	0.258546
115	metformin_Steady	-0.192531
132	glyburide_No	0.180965
163	admission_source_desc_Discharged/transferred t...	0.175976
133	glyburide_Steady	0.154605
114	metformin_No	-0.146662
167	admission_source_desc_Discharged/transferred t...	0.136092
51	medical_specialty_Gynecology	-0.134850
179	admission_source_desc_Hospice / home	-0.128575



In [47]: `from sklearn.metrics import classification_report, roc_curve, roc_auc_score, confusion_m`

```
# Generate classification report
classification_rep = classification_report(y_test, y_pred)

# Calculate ROC curve and AUC score
y_pred_proba = lasso_logistic_best.predict_proba(X_test_encoded[:, 1])
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

# create confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# threshold here:
user_defined_threshold = 0.65

# Plot classification report
print("Classification Report:")
print(classification_rep)

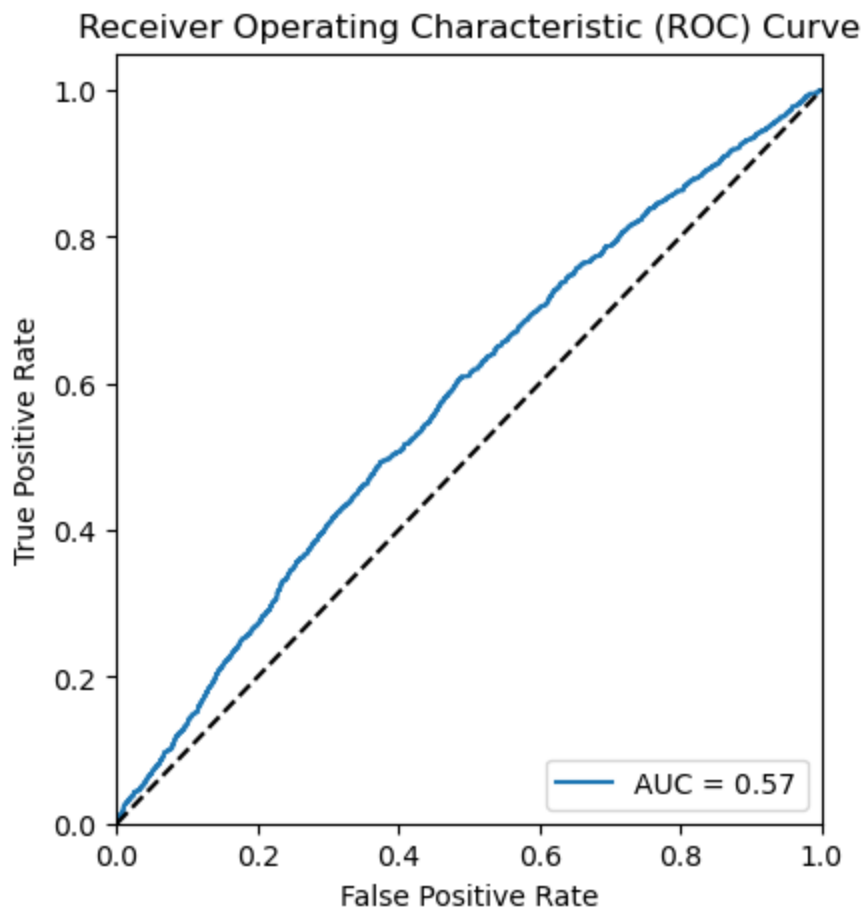
# Plot ROC curve with threshold
plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')

# apply the created threshold
y_pred_threshold = (y_pred_proba >= user_defined_threshold).astype(int)
conf_matrix_threshold = confusion_matrix(y_test, y_pred_threshold)
```

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.62	0.75	12966
1	0.10	0.49	0.16	1086
accuracy			0.61	14052
macro avg	0.52	0.56	0.46	14052
weighted avg	0.87	0.61	0.70	14052



Undersampling - Logistic Regression

```
In [48]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from imblearn.under_sampling import RandomUnderSampler
from sklearn.linear_model import LassoCV, LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc

# Initialize RandomUnderSampler
undersampler = RandomUnderSampler(sampling_strategy='auto', random_state=12)

# Fit and apply undersampling to the training data
X_train_resampled, y_train_resampled = undersampler.fit_resample(X_train_encoded, y_train_encoded)

In [49]: # Create the LassoCV model with CV alpha selection
alphas = np.logspace(-2, 2, 11)
lasso_model = LassoCV(alphas=alphas, cv=5)

# Fit the LassoCV model to the resampled training data
lasso_model.fit(X_train_resampled, y_train_resampled)

# Get the best alpha value
best_alpha = lasso_model.alpha_

# Apply the same one-hot encoding to the test data
X_test_encoded = pd.get_dummies(X_test, drop_first=True).astype(int)

# Ensure the test data has the same columns as the training data
missing_columns = set(X_train_encoded.columns) - set(X_test_encoded.columns)
```

```

for column in missing_columns:
    X_test_encoded[column] = 0 # Add missing columns with all zeros

# Reorder columns in X_test_encoded to match X_train_encoded
X_test_encoded = X_test_encoded[X_train_encoded.columns]

# Create the Lasso logistic regression model with the best alpha value
lasso_logistic_best = LogisticRegression(penalty='l1', solver='saga', C=1 / best_alpha,

# Fit the model to the resampled training data
lasso_logistic_best.fit(X_train_resampled, y_train_resampled)

# Make predictions on the test set
y_pred_prob = lasso_logistic_best.predict_proba(X_test_encoded)[: , 1]
y_pred = (y_pred_prob >= 0.982).astype(int) # Change threshold ***

# Calculate accuracy on the test set
test_accuracy = accuracy_score(y_test, y_pred)

# Display features with non-zero coefficients
coefficients = lasso_logistic_best.coef_
feature_names = X_train_encoded.columns
coef_df = pd.DataFrame({'Feature': feature_names, 'Coefficient': coefficients[0]})
non_zero_coef_df = coef_df[coef_df['Coefficient'] != 0]

print("Features with Non-Zero Coefficients:")
print(non_zero_coef_df)

print(f"Best Alpha (Lambda): {best_alpha:.4f}")
print(f"Test Set Accuracy: {test_accuracy:.4f}")

# threshold
user_defined_threshold = 0.982

# Calculate classification metrics with threshold
y_pred_custom_threshold = (y_pred_prob >= user_defined_threshold).astype(int)

classification_report_custom_threshold = classification_report(y_test, y_pred_custom_thr
confusion_matrix_custom_threshold = confusion_matrix(y_test, y_pred_custom_threshold)

print(f"User-Defined Threshold: {user_defined_threshold:.2f}")
print("Classification Report with Custom Threshold:")
print(classification_report_custom_threshold)

custom_labels = ['Not Readmitted', 'Readmitted']
# Plot confusion matrix with threshold
plt.figure(figsize=(8, 5))
sns.heatmap(confusion_matrix_custom_threshold, annot=True, fmt="d", cmap="Blues", cbar=F
            xticklabels=custom_labels, yticklabels=custom_labels)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Logistic Regression With Undersampling - Prediction Matrix")
plt.show()

```

Features with Non-Zero Coefficients:

	Feature	Coefficient
0	time_in_hospital	0.023936
1	num_lab_procedures	0.061299
2	num_procedures	-0.015450
3	num_medications	-0.013577
4	number_outpatient	0.012571
..
209	diag_2_new_Respiratory	-0.055289
210	diag_3_new_Heart Disease	-0.038178
211	diag_3_new_Injury	-0.012873
212	diag_3_new_Other	-0.045747

213 diag_3_new_Respiratory 0.018719

[214 rows x 2 columns]

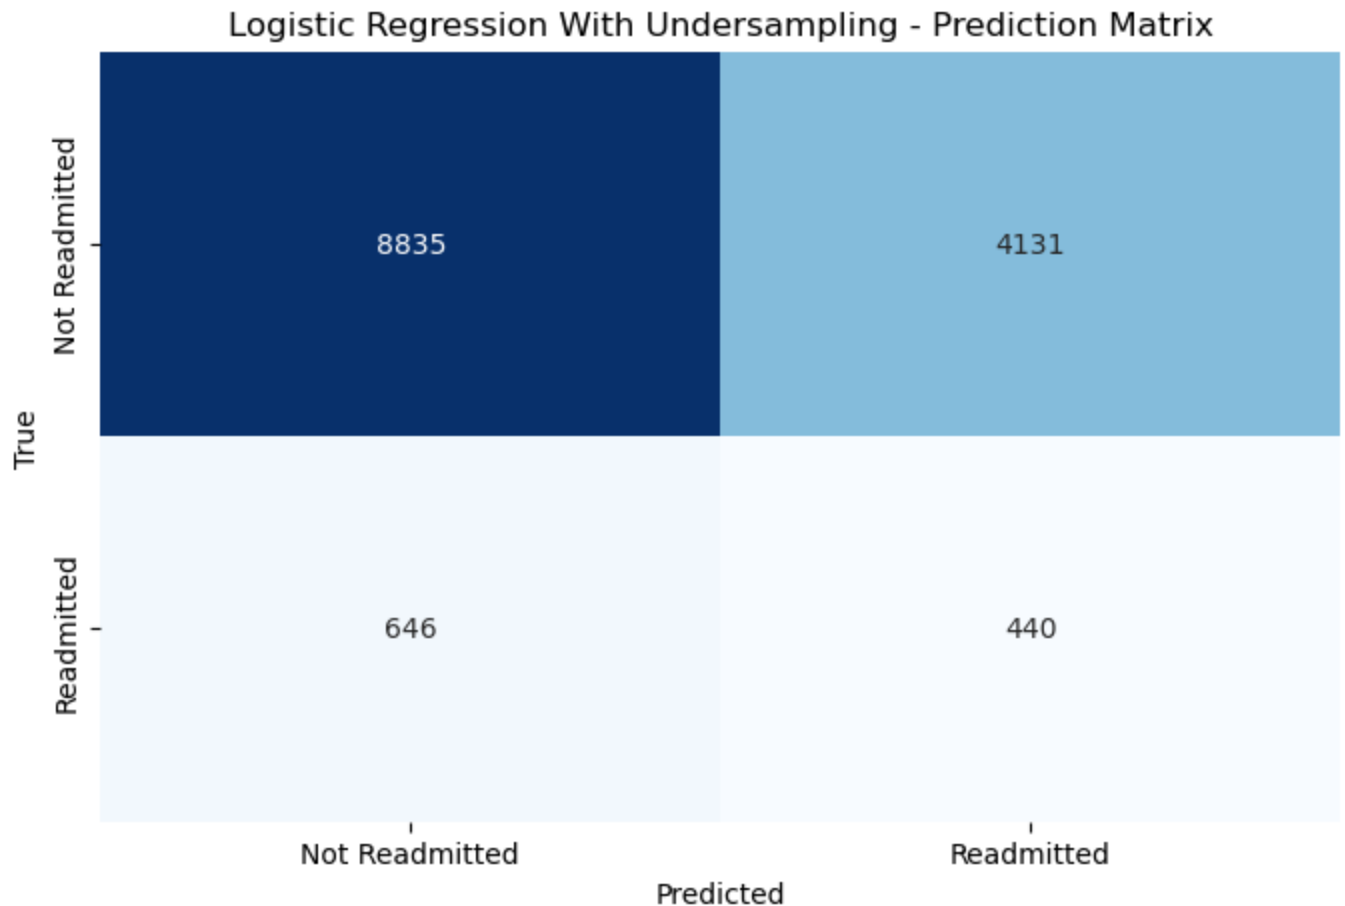
Best Alpha (Lambda): 0.0100

Test Set Accuracy: 0.6600

User-Defined Threshold: 0.98

Classification Report with Custom Threshold:

	precision	recall	f1-score	support
0	0.93	0.68	0.79	12966
1	0.10	0.41	0.16	1086
accuracy			0.66	14052
macro avg	0.51	0.54	0.47	14052
weighted avg	0.87	0.66	0.74	14052



```
In [50]: from sklearn.metrics import classification_report, roc_curve, roc_auc_score, confusion_m
```

```
# Generate classification report
classification_rep = classification_report(y_test, y_pred)

# Calculate ROC curve and AUC score
y_pred_proba = lasso_logistic_best.predict_proba(X_test_encoded)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

# create confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# threshold here:
user_defined_threshold = 0.982

# Plot classification report
print("Classification Report:")
```



```

print(classification_rep)

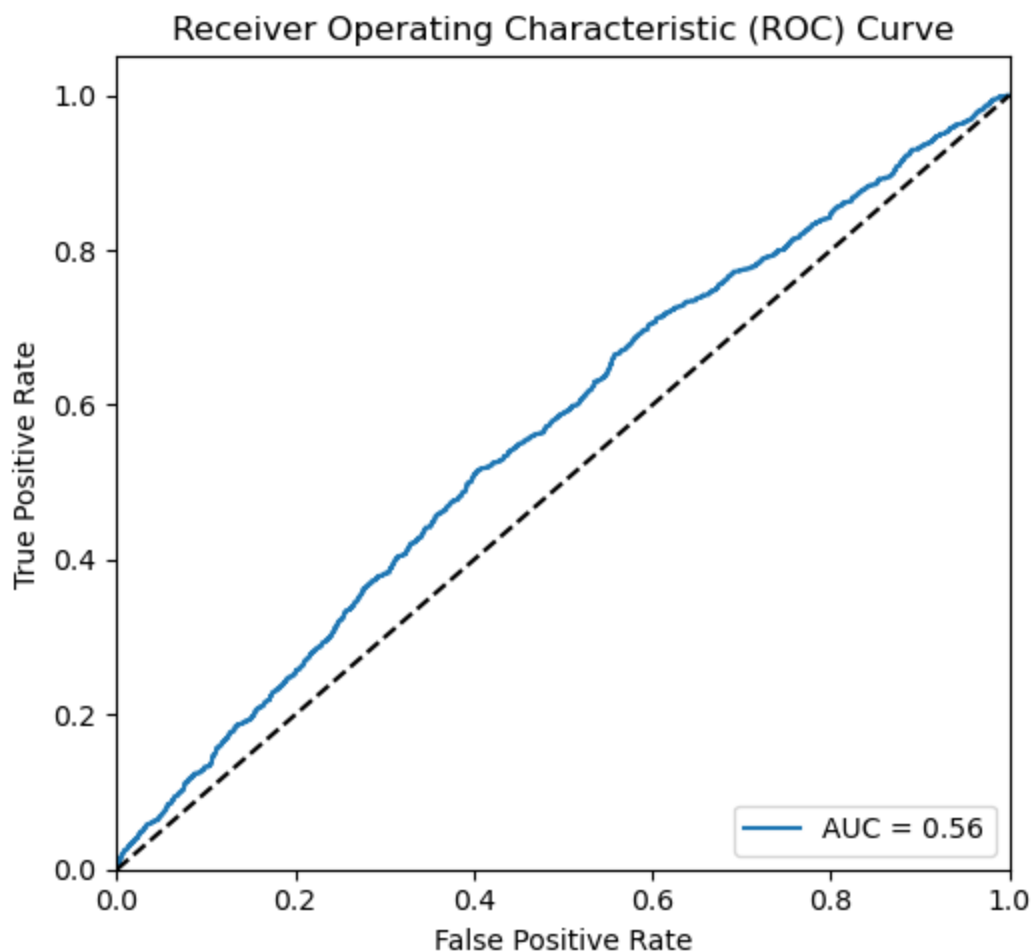
# Plot ROC curve with threshold
plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()

```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.68	0.79	12966
1	0.10	0.41	0.16	1086
accuracy			0.66	14052
macro avg	0.51	0.54	0.47	14052
weighted avg	0.87	0.66	0.74	14052



Random Forest

```

In [51]: from sklearn.ensemble import RandomForestClassifier
         from sklearn.model_selection import GridSearchCV

         # Random Forest classifier
         rf_classifier = RandomForestClassifier(random_state=12)

         # GridSearchCV
         param_grid = {
             'n_estimators': [100, 200, 300, 400, 500],
             'max_depth': [None, 10, 15, 20]
         }

         # GridSearchCV with 5-fold cross-validation
         grid_search = GridSearchCV(estimator=rf_classifier, param_grid=param_grid, cv=5, n_jobs=

         # Fit the grid search to the data
         grid_search.fit(X_train_encoded, y_train)

         # Get the best parameters and the best estimator
         best_params = grid_search.best_params_
         best_rf_classifier = grid_search.best_estimator_

         # Print the best parameters
         print("Best Parameters:")
         print(best_params)

```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
 Best Parameters:
 {'max_depth': None, 'n_estimators': 200}

```

In [52]: # Random Forest classifier with best parameters
         best_rf_classifier.fit(X_train_encoded, y_train)

         # Make probability predictions on the test set
         y_prob_rf = best_rf_classifier.predict_proba(X_test_encoded)

         # threshold
         threshold = 0.313

         # Apply the threshold
         y_pred_rf = (y_prob_rf[:, 1] >= threshold).astype(int)

         from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

         # accuracy
         accuracy = accuracy_score(y_test, y_pred_rf)
         print(f"Test Set Accuracy: {accuracy:.4f}")

         # classification report
         classification_rep_rf = classification_report(y_test, y_pred_rf)
         print("Classification Report:")
         print(classification_rep_rf)

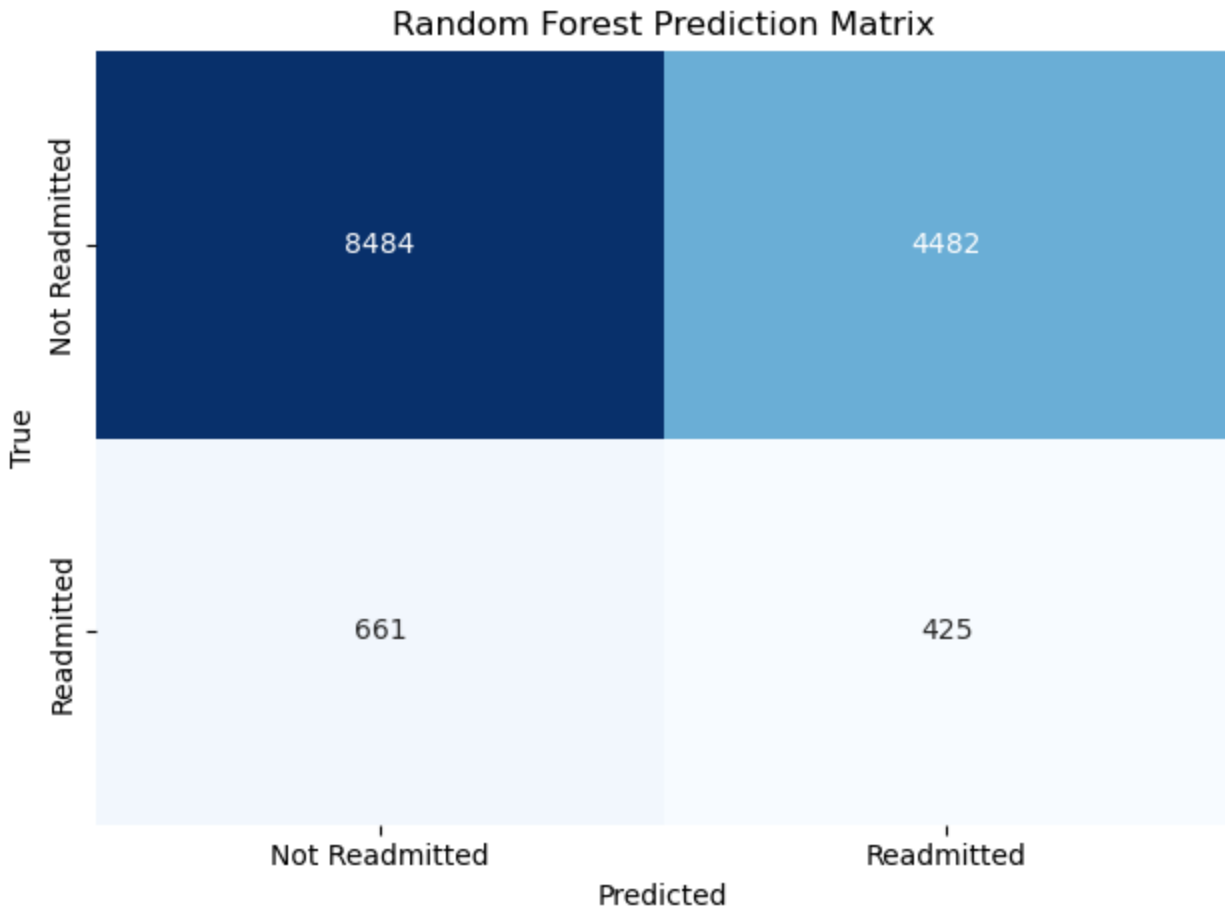
         # confusion matrix
         conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

         sns.heatmap(conf_matrix_rf, annot=True, fmt="d", cmap="Blues", cbar=False,
                     xticklabels=custom_label, yticklabels=custom_label)
         plt.xlabel("Predicted")
         plt.ylabel("True")
         plt.title(f"Random Forest Prediction Matrix")
         plt.tight_layout()
         plt.show()

```

Test Set Accuracy: 0.6340
 Classification Report:

	precision	recall	f1-score	support
0	0.93	0.65	0.77	12966
1	0.09	0.39	0.14	1086
accuracy			0.63	14052
macro avg	0.51	0.52	0.45	14052
weighted avg	0.86	0.63	0.72	14052



```
In [53]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

# Random Forest classifier
rf_classifier = RandomForestClassifier(random_state=12, class_weight='balanced')

# GridSearchCV
param_grid = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_depth': [None, 10, 15, 20]
}

# GridSearchCV with 5-fold cross-validation
grid_search = GridSearchCV(estimator=rf_classifier, param_grid=param_grid, cv=5, n_jobs=

# Fit the grid search to the data
grid_search.fit(X_train_encoded, y_train)

# Get the best parameters and the best estimator
best_params = grid_search.best_params_
best_rf_classifier = grid_search.best_estimator_

# Print the best parameters
print("Best Parameters:")
print(best_params)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
Best Parameters:
{ 'max_depth': None, 'n_estimators': 100 }

```
In [54]: import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc

# Random Forest classifier with best parameters
best_rf_classifier.fit(X_train_encoded, y_train)

# Make probability predictions on the test set
y_prob_rf = best_rf_classifier.predict_proba(X_test_encoded)

# threshold
threshold = 0.12

# Apply the threshold
y_pred_rf = (y_prob_rf[:, 1] >= threshold).astype(int)

# Calculate ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob_rf[:, 1])

# Calculate AUC (Area Under the Curve)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

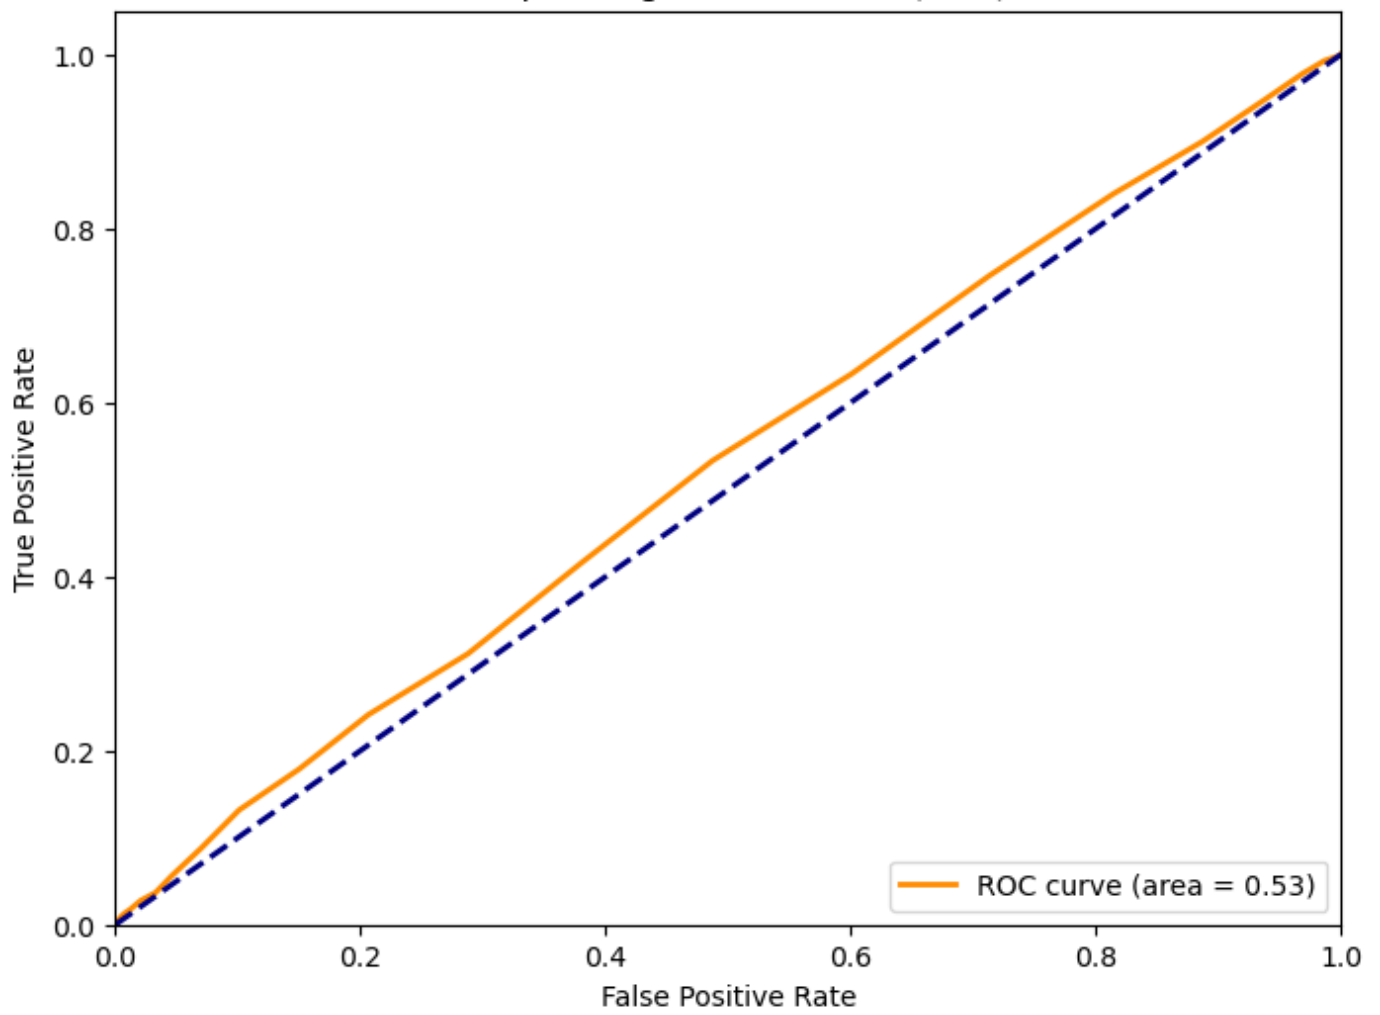
# accuracy
accuracy = accuracy_score(y_test, y_pred_rf)
print(f"Test Set Accuracy: {accuracy:.4f}")

# classification report
classification_rep_rf = classification_report(y_test, y_pred_rf)
print("Classification Report:")
print(classification_rep_rf)

# confusion matrix
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

sns.heatmap(conf_matrix_rf, annot=True, fmt="d", cmap="Blues", cbar=False,
            xticklabels=custom_label, yticklabels=custom_label)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title(f"Random Forest Prediction Matrix")
plt.tight_layout()
plt.show()
```

Receiver Operating Characteristic (ROC) Curve



Test Set Accuracy: 0.6050

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.62	0.74	12966
1	0.08	0.41	0.14	1086
accuracy			0.61	14052
macro avg	0.51	0.52	0.44	14052
weighted avg	0.86	0.61	0.70	14052

Random Forest Prediction Matrix

